

# Day 4: Advanced SQL for Data Engineering - Complete Guide

## What You'll Learn Today

- **Advanced Window Functions** for complex analytics
  - **Complex JOINS and Subqueries** for multi-table analysis
  - **Common Table Expressions (CTEs)** for readable, modular queries
  - **Query Performance Optimization** for production systems
  - **Real-world Advanced Analytics** with multiple Kaggle datasets
- 

## Learning Objectives

By the end of Day 4, you will:

1. Master advanced window functions for sophisticated analytics
  2. Build complex multi-table queries with optimal performance
  3. Use CTEs for hierarchical and recursive data processing
  4. Optimize query performance for large datasets
  5. Create production-ready analytical queries
- 

## Real Kaggle Datasets for Day 4

**Primary Dataset:** Sample Superstore Dataset - Practice Your Data Analysis Skills as a Superstore Data Analyst

- **Kaggle Link:** [kaggle.com/datasets/bravehart101/sample-supermarket-dataset](https://kaggle.com/datasets/bravehart101/sample-supermarket-dataset)
- **Size:** 9,426 records
- **Use Case:** Advanced analytics, complex JOINS, performance optimization

**Secondary Dataset:** E-Commerce Transactions Dataset

- **Kaggle Link:** [kaggle.com/datasets/smayanj/e-commerce-transactions-dataset](https://kaggle.com/datasets/smayanj/e-commerce-transactions-dataset)
- **Size:** 50,000+ records
- **Use Case:** Large dataset performance testing, complex analytics

**Download Instructions:**

```
bash
```

```
# Download both datasets
```

```
kaggle datasets download -d bravehart101/sample-supermarket-dataset --unzip
```

```
kaggle datasets download -d smayanj/e-commerce-transactions-dataset --unzip
```

---

## **Advanced Window Functions**

### **1. Ranking and Dense Ranking**

Window Functions allow calculations across a set of table rows that are related to the current row. Unlike traditional aggregate functions, which collapse the result set into a single value per group, window functions return a value for every row in the result set.

sql

-- Advanced ranking analysis

```
WITH customer_rankings AS (  
    SELECT  
        customer_name,  
        customer_id,  
        segment,  
        region,  
        SUM(sales) as total_sales,  
        SUM(profit) as total_profit,  
        COUNT(DISTINCT order_id) as total_orders,  
  
        -- Different ranking functions  
        ROW_NUMBER() OVER (ORDER BY SUM(sales) DESC) as sales_rank,  
        RANK() OVER (ORDER BY SUM(sales) DESC) as sales_rank_with_ties,  
        DENSE_RANK() OVER (ORDER BY SUM(sales) DESC) as dense_sales_rank,  
  
        -- Ranking within segments  
        ROW_NUMBER() OVER (PARTITION BY segment ORDER BY SUM(sales) DESC) as rank_in_segment,  
  
        -- Percentile ranking  
        PERCENT_RANK() OVER (ORDER BY SUM(sales)) as sales_percentile,  
        NTILE(10) OVER (ORDER BY SUM(sales)) as sales_decile  
    FROM superstore  
    GROUP BY customer_name, customer_id, segment, region  
)  
SELECT  
    customer_name,  
    segment,  
    region,  
    total_sales,  
    total_profit,  
    sales_rank,  
    rank_in_segment,  
    sales_decile,  
    CASE  
        WHEN sales_decile >= 9 THEN 'Top 20%'  
        WHEN sales_decile >= 7 THEN 'High Value'  
        WHEN sales_decile >= 4 THEN 'Medium Value'  
        ELSE 'Low Value'  
    END as customer_tier  
FROM customer_rankings
```

```
WHERE sales_rank <= 100  
ORDER BY sales_rank;
```

## 2. Advanced Lag/Lead Analysis

sql

-- Customer behavior analysis with lag/lead

WITH customer\_orders AS (

SELECT

customer\_id,  
customer\_name,  
order\_date,  
sales,  
profit,  
category,

-- Previous and next order analysis

LAG(order\_date, 1) OVER (PARTITION BY customer\_id ORDER BY order\_date) as prev\_order\_date,  
LEAD(order\_date, 1) OVER (PARTITION BY customer\_id ORDER BY order\_date) as next\_order\_date,  
LAG(sales, 1) OVER (PARTITION BY customer\_id ORDER BY order\_date) as prev\_order\_sales,

-- First and last values

FIRST\_VALUE(order\_date) OVER (PARTITION BY customer\_id ORDER BY order\_date) as first\_order\_date,  
LAST\_VALUE(order\_date) OVER (  
PARTITION BY customer\_id  
ORDER BY order\_date  
ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING  
) as last\_order\_date,

-- Order sequence numbering

ROW\_NUMBER() OVER (PARTITION BY customer\_id ORDER BY order\_date) as order\_sequence,

FROM superstore

),

customer\_insights AS (

SELECT

\*,

-- Calculate days between orders

COALESCE(order\_date - prev\_order\_date, 0) as days\_since\_last\_order,  
COALESCE(next\_order\_date - order\_date, 0) as days\_to\_next\_order,

-- Customer lifetime in days

last\_order\_date - first\_order\_date as customer\_lifetime\_days,

-- Sales trend analysis

CASE

WHEN prev\_order\_sales IS NULL THEN 'First Order'  
WHEN sales > prev\_order\_sales THEN 'Increasing'  
WHEN sales < prev\_order\_sales THEN 'Decreasing'  
ELSE 'Stable'

```

        END as sales_trend
    FROM customer_orders
)
SELECT
    customer_name,
    order_sequence,
    order_date,
    sales,
    days_since_last_order,
    sales_trend,
    customer_lifetime_days,

    -- Customer lifecycle stage
    CASE
        WHEN order_sequence = 1 THEN 'New Customer'
        WHEN days_since_last_order <= 30 THEN 'Active'
        WHEN days_since_last_order <= 90 THEN 'At Risk'
        ELSE 'Churned'
    END as customer_status
FROM customer_insights
WHERE customer_id IN (
    SELECT customer_id
    FROM customer_insights
    GROUP BY customer_id
    HAVING COUNT(*) >= 5
)
ORDER BY customer_name, order_sequence;

```

### 3. Moving Averages and Rolling Calculations



sql

*-- Advanced time series analysis*

```
WITH daily_sales AS (  
    SELECT  
        order_date,  
        SUM(sales) as daily_sales,  
        SUM(profit) as daily_profit,  
        COUNT(DISTINCT order_id) as daily_orders,  
        COUNT(DISTINCT customer_id) as daily_customers  
    FROM superstore  
    GROUP BY order_date  
)
```

sales\_analytics AS (  
 SELECT

```
        order_date,  
        daily_sales,  
        daily_profit,  
        daily_orders,  
        daily_customers,
```

*-- Moving averages*

```
    AVG(daily_sales) OVER (  
        ORDER BY order_date  
        ROWS BETWEEN 6 PRECEDING AND CURRENT ROW  
    ) as sales_7day_ma,
```

```
    AVG(daily_sales) OVER (  
        ORDER BY order_date  
        ROWS BETWEEN 29 PRECEDING AND CURRENT ROW  
    ) as sales_30day_ma,
```

*-- Rolling sums*

```
    SUM(daily_sales) OVER (  
        ORDER BY order_date  
        ROWS BETWEEN 6 PRECEDING AND CURRENT ROW  
    ) as sales_7day_rolling,
```

*-- Standard deviation for volatility*

```
    STDDEV(daily_sales) OVER (  
        ORDER BY order_date  
        ROWS BETWEEN 29 PRECEDING AND CURRENT ROW  
    ) as sales_30day_stddev,
```

*-- Min/Max in rolling window*

```

        MIN(daily_sales) OVER (
            ORDER BY order_date
            ROWS BETWEEN 29 PRECEDING AND CURRENT ROW
        ) as sales_30day_min,

        MAX(daily_sales) OVER (
            ORDER BY order_date
            ROWS BETWEEN 29 PRECEDING AND CURRENT ROW
        ) as sales_30day_max
    FROM daily_sales
)
SELECT
    order_date,
    daily_sales,
    sales_7day_ma,
    sales_30day_ma,
    sales_30day_stddev,

    -- Trend analysis
    CASE
        WHEN daily_sales > sales_30day_ma + (2 * sales_30day_stddev) THEN 'Exceptional'
        WHEN daily_sales > sales_30day_ma + sales_30day_stddev THEN 'High'
        WHEN daily_sales < sales_30day_ma - (2 * sales_30day_stddev) THEN 'Exceptional'
        WHEN daily_sales < sales_30day_ma - sales_30day_stddev THEN 'Low'
        ELSE 'Normal'
    END as sales_performance,

    -- Volatility measure
    CASE
        WHEN sales_30day_stddev / NULLIF(sales_30day_ma, 0) > 0.5 THEN 'High Volatility'
        WHEN sales_30day_stddev / NULLIF(sales_30day_ma, 0) > 0.2 THEN 'Medium Volatility'
        ELSE 'Low Volatility'
    END as volatility_level
FROM sales_analytics
WHERE order_date >= (SELECT MIN(order_date) + INTERVAL '30 days' FROM daily_sales)
ORDER BY order_date;

```

---

## Complex JOINS and Subqueries

### 1. Advanced Multi-Table Analysis

sql

-- Create dimension tables for complex analysis

-- Customer dimension

CREATE TABLE dim\_customers AS

SELECT DISTINCT

customer\_id,  
customer\_name,  
segment,  
city,  
state,  
region,  
country,

-- Customer metrics

(SELECT MIN(order\_date) FROM superstore s2 WHERE s2.customer\_id = s1.customer\_id) as min\_order\_date,

(SELECT MAX(order\_date) FROM superstore s2 WHERE s2.customer\_id = s1.customer\_id) as max\_order\_date,

(SELECT COUNT(\*) FROM superstore s2 WHERE s2.customer\_id = s1.customer\_id) as total\_orders,

(SELECT SUM(sales) FROM superstore s2 WHERE s2.customer\_id = s1.customer\_id) as total\_sales

FROM superstore s1;

-- Product dimension

CREATE TABLE dim\_products AS

SELECT DISTINCT

product\_id,  
product\_name,  
category,  
sub\_category,

-- Product metrics

(SELECT AVG(sales) FROM superstore s2 WHERE s2.product\_id = s1.product\_id) as avg\_sales,

(SELECT SUM(quantity) FROM superstore s2 WHERE s2.product\_id = s1.product\_id) as total\_quantity,

(SELECT COUNT(DISTINCT customer\_id) FROM superstore s2 WHERE s2.product\_id = s1.product\_id) as unique\_customers

FROM superstore s1;

-- Time dimension

CREATE TABLE dim\_time AS

WITH date\_range AS (

SELECT generate\_series(  
    (SELECT MIN(order\_date) FROM superstore),  
    (SELECT MAX(order\_date) FROM superstore),  
    '1 day'::interval  
)::date as date\_value

)

SELECT

date\_value,  
EXTRACT(YEAR FROM date\_value) as year,

```

EXTRACT(MONTH FROM date_value) as month,
EXTRACT(DAY FROM date_value) as day,
EXTRACT(QUARTER FROM date_value) as quarter,
EXTRACT(DOW FROM date_value) as day_of_week,
TO_CHAR(date_value, 'Day') as day_name,
TO_CHAR(date_value, 'Month') as month_name,
CASE
    WHEN EXTRACT(DOW FROM date_value) IN (0, 6) THEN 'Weekend'
    ELSE 'Weekday'
END as day_type
FROM date_range;

-- Complex multi-table analysis
WITH customer_product_affinity AS (
    SELECT
        c.customer_id,
        c.customer_name,
        c.segment,
        p.category,
        p.sub_category,
        COUNT(*) as purchase_frequency,
        SUM(s.sales) as category_spend,
        AVG(s.sales) as avg_order_value,

        -- Rank categories by customer preference
        ROW_NUMBER() OVER (
            PARTITION BY c.customer_id
            ORDER BY COUNT(*) DESC, SUM(s.sales) DESC
        ) as category_preference_rank
    FROM superstore s
    JOIN dim_customers c ON s.customer_id = c.customer_id
    JOIN dim_products p ON s.product_id = p.product_id
    JOIN dim_time t ON s.order_date = t.date_value
    WHERE t.year >= 2020
    GROUP BY c.customer_id, c.customer_name, c.segment, p.category, p.sub_category
),
customer_segments AS (
    SELECT
        customer_id,
        customer_name,
        segment,

        -- Primary category (most frequent purchases)
        MAX(CASE WHEN category_preference_rank = 1 THEN category END) as primary_category

```

```

-- Category diversity (number of different categories purchased)
COUNT(DISTINCT category) as category_diversity,

-- Total spend and frequency
SUM(category_spend) as total_spend,
SUM(purchase_frequency) as total_frequency,

-- Calculate category concentration (how focused customer is)
MAX(category_spend) / SUM(category_spend) as category_concentration
FROM customer_product_affinity
GROUP BY customer_id, customer_name, segment
)
SELECT
    cs.customer_name,
    cs.segment,
    cs.primary_category,
    cs.category_diversity,
    cs.total_spend,
    cs.total_frequency,
    ROUND(cs.category_concentration * 100, 2) as concentration_pct,

-- Customer behavior classification
CASE
    WHEN cs.category_concentration > 0.8 THEN 'Specialist'
    WHEN cs.category_concentration > 0.5 THEN 'Focused'
    ELSE 'Diversified'
END as shopping_behavior,

-- Value classification
CASE
    WHEN cs.total_spend > 15000 THEN 'High Value'
    WHEN cs.total_spend > 5000 THEN 'Medium Value'
    ELSE 'Low Value'
END as value_segment
FROM customer_segments cs
WHERE cs.total_frequency >= 5
ORDER BY cs.total_spend DESC;

```

## 2. Correlated Subqueries for Advanced Analytics

sql



-- Advanced customer analysis with correlated subqueries

SELECT

```
c.customer_name,  
c.segment,  
c.region,  
c.lifetime_value,
```

-- Compare to segment average

```
(  
    SELECT AVG(lifetime_value)  
    FROM dim_customers c2  
    WHERE c2.segment = c.segment  
) as segment_avg_value,
```

-- Customer rank within segment

```
(  
    SELECT COUNT(*) + 1  
    FROM dim_customers c2  
    WHERE c2.segment = c.segment  
    AND c2.lifetime_value > c.lifetime_value  
) as rank_in_segment,
```

-- Most purchased category

```
(  
    SELECT p.category  
    FROM superstore s  
    JOIN dim_products p ON s.product_id = p.product_id  
    WHERE s.customer_id = c.customer_id  
    GROUP BY p.category  
    ORDER BY SUM(s.sales) DESC  
    LIMIT 1  
) as favorite_category,
```

-- Average days between orders

```
(  
    SELECT AVG(days_between)  
    FROM (  
        SELECT  
            order_date - LAG(order_date) OVER (ORDER BY order_date) as days_between  
        FROM superstore s2  
        WHERE s2.customer_id = c.customer_id  
    ) t  
    WHERE days_between IS NOT NULL
```

```
) as avg_days_between_orders,  
  
-- Has made recent purchase (last 90 days)  
EXISTS (  
    SELECT 1  
    FROM superstore s3  
    WHERE s3.customer_id = c.customer_id  
    AND s3.order_date >= CURRENT_DATE - INTERVAL '90 days'  
) as recent_customer  
FROM dim_customers c  
WHERE c.total_orders >= 3  
ORDER BY c.lifetime_value DESC;
```

---



## Common Table Expressions (CTEs) Advanced Patterns

### 1. Hierarchical Data Processing

A Common Table Expression (CTE) is a temporary result set that you can reference within a SELECT, INSERT, UPDATE, or DELETE statement. CTEs make complex queries easier to read, write, and maintain by breaking them down into logical, manageable parts.

sql

-- Hierarchical product category analysis

WITH RECURSIVE category\_hierarchy AS (

-- Base case: main categories

SELECT

category as path,  
category,  
sub\_category,  
0 as level,  
category as root\_category

FROM dim\_products

WHERE category IS NOT NULL

UNION ALL

-- Recursive case: sub-categories

SELECT

ch.path || ' > ' || p.sub\_category as path,  
p.category,  
p.sub\_category,  
ch.level + 1,  
ch.root\_category

FROM category\_hierarchy ch

JOIN dim\_products p ON ch.category = p.category

WHERE p.sub\_category IS NOT NULL

AND ch.level < 2

),

category\_performance AS (

SELECT

ch.path,  
ch.level,  
ch.root\_category,  
COUNT(DISTINCT s.product\_id) as product\_count,  
COUNT(DISTINCT s.customer\_id) as customer\_count,  
SUM(s.sales) as total\_sales,  
SUM(s.profit) as total\_profit,  
AVG(s.sales) as avg\_order\_value,

-- Profit margin

CASE

WHEN SUM(s.sales) > 0 THEN SUM(s.profit) / SUM(s.sales) \* 100  
ELSE 0

END as profit\_margin\_pct

FROM category\_hierarchy ch

```

LEFT JOIN dim_products p ON (ch.category = p.category AND ch.sub_category = p.sub_
LEFT JOIN superstore s ON p.product_id = s.product_id
GROUP BY ch.path, ch.level, ch.root_category
)
SELECT
    REPEAT(' ', level) || path as hierarchy_display,
    product_count,
    customer_count,
    total_sales,
    total_profit,
    ROUND(profit_margin_pct, 2) as profit_margin_pct,

    -- Performance ranking within level
    RANK() OVER (PARTITION BY level ORDER BY total_sales DESC) as sales_rank_in_level
FROM category_performance
WHERE total_sales > 0
ORDER BY root_category, level, total_sales DESC;

```

## 2. Complex Time Series Analysis with CTEs

sql

*-- Multi-layered time series analysis*

WITH monthly\_base AS (

*-- Base monthly aggregations*

SELECT

DATE\_TRUNC('month', order\_date) as month,  
category,  
segment,  
SUM(sales) as monthly\_sales,  
SUM(profit) as monthly\_profit,  
COUNT(DISTINCT customer\_id) as monthly\_customers,  
COUNT(DISTINCT order\_id) as monthly\_orders

FROM superstore s

JOIN dim\_products p ON s.product\_id = p.product\_id

JOIN dim\_customers c ON s.customer\_id = c.customer\_id

GROUP BY DATE\_TRUNC('month', order\_date), category, segment

),

monthly\_trends AS (

*-- Add trend calculations*

SELECT

\*,  
LAG(monthly\_sales, 1) OVER (  
PARTITION BY category, segment  
ORDER BY month  
) as prev\_month\_sales,

LAG(monthly\_sales, 12) OVER (  
PARTITION BY category, segment  
ORDER BY month  
) as same\_month\_last\_year,

AVG(monthly\_sales) OVER (  
PARTITION BY category, segment  
ORDER BY month  
ROWS BETWEEN 11 PRECEDING AND CURRENT ROW  
) as trailing\_12\_month\_avg

FROM monthly\_base

),

monthly\_insights AS (

*-- Calculate growth rates and trends*

SELECT

\*,  
CASE  
WHEN prev\_month\_sales > 0 THEN

```

        ((monthly_sales - prev_month_sales) / prev_month_sales) * 100
    ELSE NULL
END as mom_growth_pct,

CASE
    WHEN same_month_last_year > 0 THEN
        ((monthly_sales - same_month_last_year) / same_month_last_year) * 100
    ELSE NULL
END as yoy_growth_pct,

CASE
    WHEN trailing_12_month_avg > 0 THEN
        ((monthly_sales - trailing_12_month_avg) / trailing_12_month_avg) * 100
    ELSE NULL
END as vs_12mo_avg_pct
FROM monthly_trends
),
performance_classification AS (
    -- Classify performance
    SELECT
        *,
        CASE
            WHEN mom_growth_pct > 20 THEN 'Accelerating'
            WHEN mom_growth_pct > 5 THEN 'Growing'
            WHEN mom_growth_pct > -5 THEN 'Stable'
            WHEN mom_growth_pct > -20 THEN 'Declining'
            ELSE 'Steep Decline'
        END as mom_trend,

        CASE
            WHEN yoy_growth_pct > 15 THEN 'Strong Growth'
            WHEN yoy_growth_pct > 5 THEN 'Moderate Growth'
            WHEN yoy_growth_pct > -5 THEN 'Flat'
            ELSE 'Declining'
        END as yoy_trend
    FROM monthly_insights
)
SELECT
    month,
    category,
    segment,
    monthly_sales,
    monthly_profit,
    monthly_customers,

```



```

ROUND(mom_growth_pct, 2) as mom_growth_pct,
ROUND(yoy_growth_pct, 2) as yoy_growth_pct,
mom_trend,
yoy_trend,

-- Overall performance score
CASE
    WHEN mom_trend IN ('Accelerating', 'Growing') AND yoy_trend IN ('Strong Growth', '
    WHEN mom_trend IN ('Growing', 'Stable') AND yoy_trend IN ('Moderate Growth', '
    WHEN mom_trend = 'Stable' AND yoy_trend = 'Flat' THEN 'Steady'
    WHEN mom_trend IN ('Declining', 'Steep Decline') OR yoy_trend = 'Declining' TH
    ELSE 'Mixed'
END as overall_performance
FROM performance_classification
WHERE month >= DATE_TRUNC('month', CURRENT_DATE) - INTERVAL '24 months'
ORDER BY category, segment, month DESC;

```

### 3. Advanced Cohort Analysis with CTEs

sql

*-- Customer cohort analysis using CTEs*

```
WITH customer_cohorts AS (  
    -- Define customer cohorts by first purchase month  
    SELECT  
        customer_id,  
        DATE_TRUNC('month', MIN(order_date)) as cohort_month  
    FROM superstore  
    GROUP BY customer_id  
,  
customer_activities AS (  
    -- Track customer activities by month  
    SELECT  
        cc.cohort_month,  
        cc.customer_id,  
        DATE_TRUNC('month', s.order_date) as activity_month,  
        SUM(s.sales) as monthly_sales,  
        COUNT(DISTINCT s.order_id) as monthly_orders  
    FROM customer_cohorts cc  
    JOIN superstore s ON cc.customer_id = s.customer_id  
    GROUP BY cc.cohort_month, cc.customer_id, DATE_TRUNC('month', s.order_date)  
,  
cohort_metrics AS (  
    -- Calculate cohort metrics  
    SELECT  
        cohort_month,  
        activity_month,  
        EXTRACT(MONTH FROM AGE(activity_month, cohort_month)) as months_since_first_purchase,  
        COUNT(DISTINCT customer_id) as active_customers,  
        SUM(monthly_sales) as cohort_revenue,  
        AVG(monthly_sales) as avg_customer_spend  
    FROM customer_activities  
    GROUP BY cohort_month, activity_month  
,  
cohort_sizes AS (  
    -- Calculate initial cohort sizes  
    SELECT  
        cohort_month,  
        COUNT(DISTINCT customer_id) as cohort_size  
    FROM customer_cohorts  
    GROUP BY cohort_month  
,  
cohort_retention AS (  
    -- Calculate retention rates
```

SELECT

```
cm.cohort_month,  
cm.months_since_first_purchase,  
cs.cohort_size,  
cm.active_customers,  
cm.cohort_revenue,  
cm.avg_customer_spend,
```

*-- Retention rate*

```
ROUND(  
    (cm.active_customers::DECIMAL / cs.cohort_size) * 100, 2  
) as retention_rate,
```

*-- Revenue per original customer*

```
ROUND(  
    cm.cohort_revenue / cs.cohort_size, 2  
) as revenue_per_original_customer
```

FROM cohort\_metrics cm

JOIN cohort\_sizes cs ON cm.cohort\_month = cs.cohort\_month

)

SELECT

```
cohort_month,  
months_since_first_purchase,  
cohort_size,  
active_customers,  
retention_rate,  
cohort_revenue,  
revenue_per_original_customer,
```

*-- Cohort performance classification*

CASE

WHEN months\_since\_first\_purchase = 0 THEN 'Acquisition'

WHEN months\_since\_first\_purchase <= 3 AND retention\_rate >= 30 THEN 'Strong Ea

WHEN months\_since\_first\_purchase <= 3 AND retention\_rate >= 15 THEN 'Moderate |

WHEN months\_since\_first\_purchase <= 3 THEN 'Weak Early Retention'

WHEN months\_since\_first\_purchase <= 12 AND retention\_rate >= 15 THEN 'Strong L

WHEN months\_since\_first\_purchase <= 12 AND retention\_rate >= 8 THEN 'Moderate |

ELSE 'Weak Long-term'

END as retention\_classification

FROM cohort\_retention

WHERE months\_since\_first\_purchase <= 24

ORDER BY cohort\_month, months\_since\_first\_purchase;

---

# Query Performance Optimization

## 1. Indexing Strategies for Data Engineering

sql

```
-- Create strategic indexes for performance
-- Covering index for customer analysis
CREATE INDEX CONCURRENTLY idx_superstore_customer_analysis
ON superstore (customer_id, order_date)
INCLUDE (sales, profit, quantity);

-- Partial index for recent high-value transactions
CREATE INDEX CONCURRENTLY idx_superstore_recent_high_value
ON superstore (order_date, sales)
WHERE order_date >= '2020-01-01' AND sales > 1000;

-- Composite index for category analysis
CREATE INDEX CONCURRENTLY idx_superstore_category_segment
ON superstore (category, segment, order_date);

-- Function-based index for date analysis
CREATE INDEX CONCURRENTLY idx_superstore_year_month
ON superstore (EXTRACT(YEAR FROM order_date), EXTRACT(MONTH FROM order_date));

-- Check index usage
SELECT
    schemaname,
    tablename,
    indexname,
    idx_scan as index_scans,
    idx_tup_read as tuples_read,
    idx_tup_fetch as tuples_fetched
FROM pg_stat_user_indexes
WHERE tablename = 'superstore'
ORDER BY idx_scan DESC;
```

## 2. Query Optimization Techniques

sql

*-- Example: Optimized vs Unoptimized queries*

*-- BEFORE: Inefficient query*

*-- This query has multiple performance issues*

```
SELECT DISTINCT
    c.customer_name,
    (SELECT SUM(s2.sales) FROM superstore s2 WHERE s2.customer_id = c.customer_id) as total_sales,
    (SELECT COUNT(*) FROM superstore s3 WHERE s3.customer_id = c.customer_id) as order_count
FROM (SELECT DISTINCT customer_id, customer_name FROM superstore) c
WHERE (SELECT SUM(s4.sales) FROM superstore s4 WHERE s4.customer_id = c.customer_id) > 5000
ORDER BY (SELECT SUM(s5.sales) FROM superstore s5 WHERE s5.customer_id = c.customer_id) DESC;
```

*-- AFTER: Optimized query*

*-- Much more efficient with single table scan and proper aggregation*

```
SELECT
    customer_name,
    SUM(sales) as total_sales,
    COUNT(*) as order_count
FROM superstore
GROUP BY customer_id, customer_name
HAVING SUM(sales) > 5000
ORDER BY total_sales DESC;
```

*-- Performance comparison query*

```
EXPLAIN (ANALYZE, BUFFERS, FORMAT JSON)
```

```
SELECT
    customer_name,
    SUM(sales) as total_sales,
    COUNT(*) as order_count,
    AVG(sales) as avg_order_value,
    MAX(order_date) as last_order_date
FROM superstore
WHERE order_date >= '2020-01-01'
GROUP BY customer_id, customer_name
HAVING SUM(sales) > 1000
ORDER BY total_sales DESC
LIMIT 100;
```

### 3. Materialized Views for Performance

sql

*-- Create materialized views for heavy analytical queries*

```
CREATE MATERIALIZED VIEW mv_customer_monthly_summary AS
WITH monthly_customer_metrics AS (
    SELECT
        customer_id,
        customer_name,
        segment,
        region,
        DATE_TRUNC('month', order_date) as month,
        SUM(sales) as monthly_sales,
        SUM(profit) as monthly_profit,
        COUNT(DISTINCT order_id) as monthly_orders,
        AVG(sales) as avg_order_value
    FROM superstore
    GROUP BY customer_id, customer_name, segment, region, DATE_TRUNC('month', order_da
)
SELECT
    *,
    LAG(monthly_sales, 1) OVER (
        PARTITION BY customer_id
        ORDER BY month
    ) as prev_month_sales,

-- Growth calculations
CASE
    WHEN LAG(monthly_sales, 1) OVER (PARTITION BY customer_id ORDER BY month) > 0
        ((monthly_sales - LAG(monthly_sales, 1) OVER (PARTITION BY customer_id ORDER BY month))
        / LAG(monthly_sales, 1) OVER (PARTITION BY customer_id ORDER BY month)) *
    ELSE NULL
END as month_over_month_growth,

-- Running totals
SUM(monthly_sales) OVER (
    PARTITION BY customer_id
    ORDER BY month
    ROWS UNBOUNDED PRECEDING
) as cumulative_sales
FROM monthly_customer_metrics;

-- Create indexes on materialized view
CREATE INDEX idx_mv_customer_monthly_customer_month
ON mv_customer_monthly_summary (customer_id, month);
```



```
CREATE INDEX idx_mv_customer_monthly_segment_month
ON mv_customer_monthly_summary (segment, month);

-- Refresh strategy (run this in your ETL pipeline)
REFRESH MATERIALIZED VIEW CONCURRENTLY mv_customer_monthly_summary;
```

---



## Advanced Analytics Patterns

### 1. Market Basket Analysis

sql

*-- Advanced market basket analysis*

WITH order\_products AS (

SELECT

order\_id,  
product\_name,  
category,  
sub\_category,  
sales,  
profit

FROM superstore

),

product\_pairs AS (

*-- Find products bought together*

SELECT

a.product\_name as product\_a,  
b.product\_name as product\_b,  
a.category as category\_a,  
b.category as category\_b,  
COUNT(\*) as frequency,  
AVG(a.sales + b.sales) as avg\_combined\_value,  
SUM(a.profit + b.profit) as total\_combined\_profit

FROM order\_products a

JOIN order\_products b ON a.order\_id = b.order\_id

WHERE a.product\_name < b.product\_name *-- Avoid duplicates and self-joins*

GROUP BY a.product\_name, b.product\_name, a.category, b.category

HAVING COUNT(\*) >= 5 *-- Minimum support threshold*

),

product\_statistics AS (

*-- Calculate individual product statistics*

SELECT

product\_name,  
COUNT(DISTINCT order\_id) as total\_orders

FROM order\_products

GROUP BY product\_name

),

association\_metrics AS (

*-- Calculate association rule metrics*

SELECT

pp.\*,  
psa.total\_orders as orders\_a,  
psb.total\_orders as orders\_b,  
(SELECT COUNT(DISTINCT order\_id) FROM order\_products) as total\_unique\_orders,

```

-- Support:  $P(A \cap B)$ 
pp.frequency::DECIMAL / (SELECT COUNT(DISTINCT order_id) FROM order_products) ;

-- Confidence:  $P(B|A) = P(A \cap B) / P(A)$ 
pp.frequency::DECIMAL / psa.total_orders as confidence_a_to_b,

-- Confidence:  $P(A|B) = P(A \cap B) / P(B)$ 
pp.frequency::DECIMAL / psb.total_orders as confidence_b_to_a,

-- Lift:  $P(A \cap B) / (P(A) * P(B))$ 
(pp.frequency::DECIMAL / (SELECT COUNT(DISTINCT order_id) FROM order_products)
((psa.total_orders::DECIMAL / (SELECT COUNT(DISTINCT order_id) FROM order_products)
(psb.total_orders::DECIMAL / (SELECT COUNT(DISTINCT order_id) FROM order_products)
FROM product_pairs pp
JOIN product_statistics psa ON pp.product_a = psa.product_name
JOIN product_statistics psb ON pp.product_b = psb.product_name
)
SELECT
    product_a,
    product_b,
    category_a,
    category_b,
    frequency,
    ROUND(support * 100, 3) as support_pct,
    ROUND(confidence_a_to_b * 100, 2) as confidence_a_to_b_pct,
    ROUND(confidence_b_to_a * 100, 2) as confidence_b_to_a_pct,
    ROUND(lift, 3) as lift,
    ROUND(avg_combined_value, 2) as avg_combined_value,

-- Business interpretation
CASE
    WHEN lift > 2 THEN 'Strong Association'
    WHEN lift > 1.5 THEN 'Moderate Association'
    WHEN lift > 1 THEN 'Weak Association'
    ELSE 'Negative Association'
END as association_strength,

-- Cross-category insights
CASE
    WHEN category_a = category_b THEN 'Same Category'
    ELSE 'Cross Category'
END as category_relationship
FROM association_metrics
WHERE lift > 1 -- Only show positive associations

```

```
ORDER BY lift DESC, frequency DESC  
LIMIT 50;
```

## 2. Customer Lifetime Value Prediction

sql

-- Advanced CLV analysis with predictive elements

```
WITH customer_purchase_history AS (  
    SELECT  
        customer_id,  
        customer_name,  
        segment,  
        region,  
        MIN(order_date) as first_purchase_date,  
        MAX(order_date) as last_purchase_date,  
        COUNT(DISTINCT order_id) as total_orders,  
        COUNT(DISTINCT DATE_TRUNC('month', order_date)) as active_months,  
        SUM(sales) as total_revenue,  
        AVG(sales) as avg_order_value,  
        SUM(profit) as total_profit,  
  
        -- Calculate customer lifespan in days  
        MAX(order_date) - MIN(order_date) as customer_lifespan_days,  
  
        -- Calculate average days between orders  
        CASE  
            WHEN COUNT(DISTINCT order_id) > 1 THEN  
                (MAX(order_date) - MIN(order_date)) / (COUNT(DISTINCT order_id) - 1)  
            ELSE NULL  
        END as avg_days_between_orders,  
  
        -- Recency (days since last purchase)  
        CURRENT_DATE - MAX(order_date) as days_since_last_purchase  
    FROM superstore  
    GROUP BY customer_id, customer_name, segment, region  
) ,  
customer_clv_metrics AS (  
    SELECT  
        *,  
        -- Purchase frequency (orders per month)  
        CASE  
            WHEN customer_lifespan_days > 0 THEN  
                (total_orders::DECIMAL / (customer_lifespan_days::DECIMAL / 30.44))  
            ELSE total_orders  
        END as purchase_frequency_monthly,  
  
        -- Customer lifetime value calculation  
        -- CLV = Average Order Value × Purchase Frequency × Customer Lifespan  
        CASE
```

```

        WHEN customer_lifespan_days > 0 AND avg_days_between_orders > 0 THEN
            avg_order_value * (total_orders::DECIMAL / (customer_lifespan_days::DECIMAL / 30.44))
        ELSE total_revenue
    END as historical_clv,

    -- Predicted future value (simple model)
    CASE
        WHEN days_since_last_purchase <= 90 AND avg_days_between_orders > 0 THEN
            avg_order_value * (365.0 / avg_days_between_orders) * 2 -- Predict next 2 orders
        ELSE 0
    END as predicted_future_value
FROM customer_purchase_history
WHERE total_orders >= 2 -- Only customers with multiple purchases
),
customer_segmentation AS (
    SELECT
        *,
        -- RFM-based segmentation
        NTILE(5) OVER (ORDER BY days_since_last_purchase DESC) as recency_score,
        NTILE(5) OVER (ORDER BY purchase_frequency_monthly) as frequency_score,
        NTILE(5) OVER (ORDER BY avg_order_value) as monetary_score,

        -- CLV-based segmentation
        NTILE(10) OVER (ORDER BY historical_clv) as clv_decile,

        -- Customer status
        CASE
            WHEN days_since_last_purchase <= 30 THEN 'Active'
            WHEN days_since_last_purchase <= 90 THEN 'At Risk'
            WHEN days_since_last_purchase <= 180 THEN 'Dormant'
            ELSE 'Lost'
        END as customer_status
    FROM customer_clv_metrics
)
SELECT
    customer_name,
    segment,
    region,
    total_orders,
    ROUND(avg_order_value, 2) as avg_order_value,
    ROUND(purchase_frequency_monthly, 3) as monthly_frequency,
    days_since_last_purchase,
    ROUND(historical_clv, 2) as historical_clv,

```



```

ROUND(predicted_future_value, 2) as predicted_future_value,
ROUND(historical_clv + predicted_future_value, 2) as total_clv,

-- RFM scores combined
CONCAT(recency_score, frequency_score, monetary_score) as rfm_score,

-- CLV segment
CASE
    WHEN clv_decile >= 9 THEN 'Champions'
    WHEN clv_decile >= 7 THEN 'Loyal Customers'
    WHEN clv_decile >= 5 THEN 'Potential Loyalists'
    WHEN clv_decile >= 3 THEN 'New Customers'
    ELSE 'At Risk'
END as clv_segment,

customer_status,

-- Business recommendations
CASE
    WHEN clv_decile >= 8 AND customer_status = 'Active' THEN 'VIP Treatment'
    WHEN clv_decile >= 6 AND customer_status = 'At Risk' THEN 'Retention Campaign'
    WHEN clv_decile >= 4 AND customer_status = 'Active' THEN 'Upsell Opportunity'
    WHEN customer_status = 'Lost' AND historical_clv > 5000 THEN 'Win Back Campaign'
    ELSE 'Standard Treatment'
END as recommended_action
FROM customer_segmentation
ORDER BY total_clv DESC;

```

### 3. Time Series Forecasting Foundations

sql

*-- Advanced time series analysis for forecasting*

```
WITH daily_metrics AS (  
    SELECT  
        order_date,  
        category,  
        SUM(sales) as daily_sales,  
        COUNT(DISTINCT order_id) as daily_orders,  
        COUNT(DISTINCT customer_id) as daily_customers,  
        AVG(sales) as avg_order_value  
    FROM superstore  
    WHERE order_date >= '2020-01-01'  
    GROUP BY order_date, category  
) ,  
time_series_features AS (  
    SELECT  
        *,  
        -- Day of week effects  
        EXTRACT(DOW FROM order_date) as day_of_week,  
        EXTRACT(DAY FROM order_date) as day_of_month,  
        EXTRACT(MONTH FROM order_date) as month,  
        EXTRACT(QUARTER FROM order_date) as quarter,  
  
        -- Lag features  
        LAG(daily_sales, 1) OVER (PARTITION BY category ORDER BY order_date) as sales_  
        LAG(daily_sales, 7) OVER (PARTITION BY category ORDER BY order_date) as sales_  
        LAG(daily_sales, 30) OVER (PARTITION BY category ORDER BY order_date) as sales_  
  
        -- Moving averages  
        AVG(daily_sales) OVER (  
            PARTITION BY category  
            ORDER BY order_date  
            ROWS BETWEEN 6 PRECEDING AND CURRENT ROW  
        ) as ma_7day,  
  
        AVG(daily_sales) OVER (  
            PARTITION BY category  
            ORDER BY order_date  
            ROWS BETWEEN 29 PRECEDING AND CURRENT ROW  
        ) as ma_30day,  
  
        -- Trend calculation (simple linear)  
        ROW_NUMBER() OVER (PARTITION BY category ORDER BY order_date) as time_index  
    FROM daily_metrics
```

```

),
seasonal_analysis AS (
    SELECT
        *,
        -- Seasonal decomposition components
        AVG(daily_sales) OVER (PARTITION BY category, month) as seasonal_monthly,
        AVG(daily_sales) OVER (PARTITION BY category, day_of_week) as seasonal_weekly,
        AVG(daily_sales) OVER (PARTITION BY category) as overall_mean,

        -- Detrended and deseasonalized values
        daily_sales - ma_30day as detrended_sales,
        daily_sales / NULLIF(seasonal_monthly, 0) as deseasonalized_sales
    FROM time_series_features
),
forecast_base AS (
    SELECT
        *,
        -- Simple forecast components
        -- Trend component (using linear regression on time_index)
        -- Note: This is a simplified approach; real forecasting would use more sophis

        -- Combine trend + seasonal for basic forecast
        ma_30day + (seasonal_monthly - overall_mean) as basic_forecast,

        -- Calculate forecast confidence based on historical volatility
        STDDEV(daily_sales) OVER (
            PARTITION BY category
            ORDER BY order_date
            ROWS BETWEEN 29 PRECEDING AND CURRENT ROW
        ) as volatility_30day
    FROM seasonal_analysis
)
SELECT
    order_date,
    category,
    daily_sales,
    ma_7day,
    ma_30day,
    ROUND(basic_forecast, 2) as basic_forecast,
    ROUND(volatility_30day, 2) as volatility,

    -- Forecast bounds (simple confidence intervals)
    ROUND(basic_forecast - (1.96 * volatility_30day), 2) as forecast_lower_bound,
    ROUND(basic_forecast + (1.96 * volatility_30day), 2) as forecast_upper_bound,

```

```

-- Forecast accuracy metrics (for historical validation)
CASE
    WHEN basic_forecast > 0 THEN
        ABS(daily_sales - basic_forecast) / basic_forecast * 100
    ELSE NULL
END as forecast_error_pct,

-- Trend classification
CASE
    WHEN ma_7day > ma_30day * 1.1 THEN 'Strong Upward'
    WHEN ma_7day > ma_30day * 1.05 THEN 'Upward'
    WHEN ma_7day < ma_30day * 0.9 THEN 'Strong Downward'
    WHEN ma_7day < ma_30day * 0.95 THEN 'Downward'
    ELSE 'Stable'
END as trend_direction
FROM forecast_base
WHERE order_date >= '2021-01-01' -- Focus on recent data for better forecast quality
ORDER BY category, order_date;

```

---

## Production-Ready Query Templates

### 1. Automated Data Quality Monitoring

sql

-- Comprehensive data quality monitoring system

CREATE OR REPLACE FUNCTION run\_data\_quality\_checks()

RETURNS TABLE(

check\_date DATE,  
table\_name TEXT,  
check\_type TEXT,  
check\_description TEXT,  
expected\_result TEXT,  
actual\_result TEXT,  
status TEXT,  
severity TEXT

) AS \$

BEGIN

RETURN QUERY

-- Completeness checks

SELECT

CURRENT\_DATE as check\_date,  
'superstore'::TEXT as table\_name,  
'COMPLETENESS'::TEXT as check\_type,  
'Order ID completeness'::TEXT as check\_description,  
'100%'::TEXT as expected\_result,  
ROUND((COUNT(order\_id)::DECIMAL / COUNT(\*)) \* 100, 2)::TEXT || '%' as actual\_r  
CASE WHEN COUNT(order\_id) = COUNT(\*) THEN 'PASS' ELSE 'FAIL' END as status,  
'HIGH'::TEXT as severity

FROM superstore

UNION ALL

-- Validity checks

SELECT

CURRENT\_DATE,  
'superstore'::TEXT,  
'VALIDITY'::TEXT,  
'Sales amount validity (positive values)',  
'100%'::TEXT,  
ROUND((COUNT(CASE WHEN sales > 0 THEN 1 END)::DECIMAL / COUNT(\*)) \* 100, 2)::T  
CASE WHEN COUNT(CASE WHEN sales <= 0 THEN 1 END) = 0 THEN 'PASS' ELSE 'FAIL' EI  
'HIGH'::TEXT

FROM superstore

UNION ALL

-- Consistency checks

SELECT

```
CURRENT_DATE,  
'superstore'::TEXT,  
'CONSISTENCY'::TEXT,  
'Order date before ship date',  
'100%'::TEXT,  
ROUND((COUNT(CASE WHEN order_date <= ship_date THEN 1 END)::DECIMAL / COUNT(*)  
CASE WHEN COUNT(CASE WHEN order_date > ship_date THEN 1 END) = 0 THEN 'PASS' E  
'MEDIUM'::TEXT
```

FROM superstore

UNION ALL

-- Uniqueness checks

SELECT

```
CURRENT_DATE,  
'superstore'::TEXT,  
'UNIQUENESS'::TEXT,  
'Row ID uniqueness',  
'100%'::TEXT,  
CASE WHEN COUNT(*) = COUNT(DISTINCT row_id) THEN '100%'  
      ELSE ROUND((COUNT(DISTINCT row_id)::DECIMAL / COUNT(*)) * 100, 2)::TEXT |  
END,  
CASE WHEN COUNT(*) = COUNT(DISTINCT row_id) THEN 'PASS' ELSE 'FAIL' END,  
'HIGH'::TEXT
```

FROM superstore

UNION ALL

-- Timeliness checks

SELECT

```
CURRENT_DATE,  
'superstore'::TEXT,  
'TIMELINESS'::TEXT,  
'Recent data availability (last 7 days)',  
'YES'::TEXT,  
CASE WHEN MAX(order_date) >= CURRENT_DATE - INTERVAL '7 days' THEN 'YES' ELSE  
CASE WHEN MAX(order_date) >= CURRENT_DATE - INTERVAL '7 days' THEN 'PASS' ELSE  
'MEDIUM'::TEXT
```

FROM superstore;

END;

\$ LANGUAGE plpgsql;



```
-- Execute data quality checks  
SELECT * FROM run_data_quality_checks();
```

## 2. Performance Monitoring Queries

sql

-- Query performance monitoring

CREATE VIEW query\_performance\_monitor AS

WITH query\_stats AS (

SELECT

query,

calls,

total\_time,

mean\_time,

stddev\_time,

rows,

100.0 \* shared\_blks\_hit / nullif(shared\_blks\_hit + shared\_blks\_read, 0) AS hit,

100.0 \* shared\_blks\_dirtied / nullif(shared\_blks\_hit + shared\_blks\_read, 0) AS

FROM pg\_stat\_statements

WHERE query NOT LIKE '%pg\_stat%'

AND query NOT LIKE '%information\_schema%'

AND calls > 5

),

slow\_queries AS (

SELECT \*,

CASE

WHEN mean\_time > 1000 THEN 'SLOW'

WHEN mean\_time > 500 THEN 'MODERATE'

ELSE 'FAST'

END as performance\_category

FROM query\_stats

)

SELECT

LEFT(query, 100) as query\_preview,

calls,

ROUND(total\_time::numeric, 2) as total\_time\_ms,

ROUND(mean\_time::numeric, 2) as avg\_time\_ms,

ROUND(stddev\_time::numeric, 2) as stddev\_time\_ms,

rows,

ROUND(hit\_percent::numeric, 2) as cache\_hit\_percent,

performance\_category,

-- Recommendations

CASE

WHEN hit\_percent < 90 THEN 'Consider adding indexes'

WHEN mean\_time > 1000 AND calls > 100 THEN 'Priority optimization target'

WHEN stddev\_time > mean\_time \* 2 THEN 'Inconsistent performance - investigate'

ELSE 'Performance acceptable'

END as recommendation

```
FROM slow_queries
ORDER BY total_time DESC;
```

---

## Essential Resources for Day 4

### Advanced SQL Learning

#### 1. Advanced SQL Techniques for Data Scientists

- Explore 7 essential advanced SQL techniques, including CTEs, window functions, and more, to streamline your data analysis process
- Source: SQLPad.io

#### 2. Window Functions and CTEs Guide

- Advanced SQL features such as Window Functions and Common Table Expressions (CTEs) provide powerful tools for performing complex data analysis and manipulation
- Source: CSInfo360

#### 3. 15 Advanced SQL Concepts

- Common Table Expressions (CTEs) and recursive queries can also be used to calculate running totals
- Source: Airbyte

### Video Courses

#### 1. SQL for Data Analysis: Advanced Querying - Udemy

- Learn advanced data analysis with SQL, and master topics like subqueries, CTEs, window functions, and more
- Focus: Practical application of advanced SQL techniques

#### 2. Advanced SQL Techniques for Data Engineering

- Window Functions, Common Table Expressions (CTEs), and other complex methods are used in Advanced SQL approaches for Data Engineering
- Source: Oracle PL/SQL Tutorial

### Practice Datasets

#### 1. **Primary:** Sample Superstore Dataset

- **Link:** [kaggle.com/datasets/bravehart101/sample-supermarket-dataset](https://www.kaggle.com/datasets/bravehart101/sample-supermarket-dataset)
- **Use:** Complex JOINS, performance optimization

#### 2. **Secondary:** E-Commerce Transactions

- **Link:** [kaggle.com/datasets/smayanj/e-commerce-transactions-dataset](https://kaggle.com/datasets/smayanj/e-commerce-transactions-dataset)
- **Use:** Large dataset performance testing

## Tools and Performance

### 1. PostgreSQL Documentation

- **Query Optimization:** Official PostgreSQL performance tuning guide
- **Window Functions:** Comprehensive function reference

### 2. pgAdmin and Performance Tools

- **Query Analysis:** EXPLAIN ANALYZE for query optimization
  - **Index Management:** Index usage statistics and recommendations
- 

## Day 4 Practical Tasks

### Task 1: Advanced Window Functions (60 minutes)

- ☐ Implement customer ranking analysis with multiple ranking functions
- ☐ Create moving averages and rolling calculations for time series
- ☐ Build lag/lead analysis for customer behavior tracking
- ☐ Practice FIRST\_VALUE/LAST\_VALUE for cohort analysis

### Task 2: Complex CTEs and Joins (75 minutes)

- ☐ Build hierarchical product category analysis
- ☐ Create recursive CTEs for organizational data
- ☐ Implement multi-layered time series analysis
- ☐ Design complex multi-table analytical queries

### Task 3: Performance Optimization (45 minutes)

- ☐ Create strategic indexes for query performance
- ☐ Analyze query execution plans with EXPLAIN
- ☐ Build materialized views for heavy analytics
- ☐ Implement query performance monitoring

### Task 4: Advanced Analytics (90 minutes)

- ☐ Build market basket analysis with association rules
- ☐ Create customer lifetime value analysis
- ☐ Implement cohort retention analysis
- ☐ Design time series forecasting foundations

## Task 5: Production Readiness (30 minutes)

- ☐ Set up automated data quality monitoring
  - ☐ Create performance monitoring views
  - ☐ Document query optimization strategies
  - ☐ Build reusable analytical query templates
- 



## Day 4 Deliverables

### 1. Advanced SQL Mastery

- Complex window functions for sophisticated analytics
- Multi-layered CTEs for readable, maintainable queries
- Optimized query performance with proper indexing
- Production-ready analytical query templates

### 2. Real Business Analytics

- Customer segmentation and lifetime value analysis
- Market basket analysis with association rules
- Time series forecasting foundations
- Cohort retention and behavior analysis

### 3. Performance Optimization

- Strategic indexing for large datasets
- Materialized views for heavy analytical workloads
- Query performance monitoring and alerting
- Execution plan analysis and optimization

### 4. Skills Assessment

Rate yourself after today (1-10):

- ☐ Advanced window functions: \_\_\_\_/10
- ☐ Complex CTEs and recursive queries: \_\_\_\_/10
- ☐ Multi-table JOINS and subqueries: \_\_\_\_/10
- ☐ Query performance optimization: \_\_\_\_/10
- ☐ Production SQL patterns: \_\_\_\_/10

## 5. GitHub Repository Update

bash

*# Commit your Day 4 work*

git add .

git commit -m "Day 4: Advanced SQL techniques and performance optimization"

git push origin main

## 6. Learning Journal Entry

Create `day-04/learning-notes.md`:

markdown

## # Day 4: Advanced SQL – Learning Notes

### ## Key Concepts Mastered

- Advanced window functions for complex analytics
- Common Table Expressions (CTEs) for modular queries
- Complex JOINS and correlated subqueries
- Query performance optimization strategies
- Production-ready analytical patterns

### ## Advanced Techniques Learned

- Market basket analysis with association rules
- Customer lifetime value calculation
- Cohort retention analysis
- Time series forecasting foundations
- Hierarchical data processing with recursive CTEs

### ## Performance Optimization Skills

- Strategic indexing for analytical workloads
- Materialized views for heavy computations
- Query execution plan analysis
- Performance monitoring and alerting

### ## Real Business Applications

- Customer segmentation and RFM analysis
- Product recommendation systems
- Sales forecasting and trend analysis
- Data quality monitoring automation

### ## Production Patterns Implemented

- Automated data quality checking functions
- Performance monitoring views
- Reusable analytical query templates
- Error handling and optimization strategies

### ## Tomorrow's Preparation

- Review data modeling concepts
- Prepare for dimensional modeling
- Study star schema design patterns




### What to expect:



- Dimensional modeling and star schema design
- Fact and dimension table architecture
- Data warehouse design patterns
- Normalization vs denormalization strategies
- Building scalable data models for analytics

### Preparation:

- Review relational database concepts
- Understand business requirements analysis
- Prepare for hands-on data warehouse design

---

 Congratulations on completing Day 4! You now have advanced SQL skills that rival experienced data engineers. Tomorrow, we'll learn how to design robust data models that scale.

**Progress:** 8% (4/50 days) | **Next:** Day 5 - Data Modeling | **Skills:** Python  + SQL  + Advanced SQL 