Solution ■ Mastering Workflow Orchestration for Data Engineers

What You'll Learn Today (Concept-First Approach)

Primary Focus: Understanding workflow orchestration and why it's the heart of data engineering **Secondary Focus**: Hands-on implementation through Airflow Web UI and visual tools **Dataset for Context**: Customer Analytics Dataset from Kaggle for automated processing

Solution Learning Philosophy for Day 10

"Understand the conductor before learning the orchestra"

We'll start with workflow concepts, explore Airflow's visual interface, understand DAG design patterns, and build production-ready automated pipelines.

💢 The Orchestration Revolution: Why Airflow Matters

The Problem: Chaos in Data Pipeline Management

Scenario: You're managing a customer analytics pipeline with multiple dependent tasks...

```
Without Orchestration (Manual Chaos):

Monday 6 AM: Run data extraction script

Monday 6:15 AM: Check if extraction completed

Monday 6:30 AM: Run data cleaning (if extraction worked)

Monday 6:45 AM: Check cleaning status

Monday 7:00 AM: Run analytics (if cleaning worked)

Monday 7:15 AM: Generate reports (if analytics worked)

Monday 7:30 AM: Send email alerts
```

X Problems:

- Manual intervention required at each step
- No automatic retry on failures
- Hard to track what succeeded/failed
- Impossible to scale across teams
- No historical execution tracking

With Airflow Orchestration:

- Define workflow once
- Automatic execution based on schedule
- ✓ Smart dependency management
- ✓ Automatic retries on failures
- ✓ Rich monitoring and alerting
- ✓ Historical tracking and debugging
- ✓ Scalable across entire organization

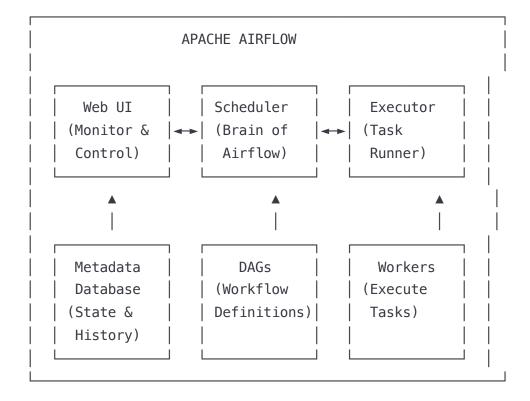
The Airflow Solution: Workflows as Code

Think of Airflow like this:

- Traditional Way: Manually conducting an orchestra, cuing each musician
- Airflow Way: Written musical score that orchestras can follow automatically

Understanding Airflow Architecture (Visual Approach)

The Airflow Mental Model



Key Airflow Components

- 1. DAGs (Directed Acyclic Graphs)
 - **Directed**: Tasks flow in specific directions

- **Acyclic**: No circular dependencies (no infinite loops)
- **Graph**: Visual representation of workflow

2. Tasks

- Individual units of work (extract data, run analysis, send email)
- Can be Python functions, bash commands, SQL queries, etc.

3. Operators

Templates for creating tasks (PythonOperator, BashOperator, SqlOperator)

4. Scheduler

Monitors DAGs and triggers tasks when dependencies are met

5. Executor

• Runs tasks on workers (local machine, cluster, cloud)

6. Web UI

Visual interface for monitoring, debugging, and controlling workflows

Airflow Installation and Setup (UI-First Approach)

Quick Start with Docker (Visual Learning)

Step 1: Download Airflow with Docker

Create project structure:

```
airflow-customer-analytics/

— docker-compose.yml

— dags/

— customer_analytics_dag.py

— data/

— customer_data.csv

— logs/
— plugins/
— scripts/
```

Step 2: Docker Compose for Airflow

Create (docker-compose.yml):

```
x-airflow-common:
  &airflow-common
  image: apache/airflow:2.7.1
 environment: &airflow-common-env
    AIRFLOW CORE EXECUTOR: LocalExecutor
   AIRFLOW__DATABASE__SQL_ALCHEMY_CONN: postgresql+psycopg2://airflow:airflow@postgres
   AIRFLOW__CORE__FERNET_KEY: ''
   AIRFLOW__CORE__DAGS_ARE_PAUSED_AT_CREATION: 'true'
   AIRFLOW__CORE__LOAD_EXAMPLES: 'false'
   AIRFLOW__API__AUTH_BACKENDS: 'airflow.api.auth.backend.basic_auth'
   AIRFLOW__WEBSERVER__EXPOSE_CONFIG: 'true'
 volumes:
   - ./dags:/opt/airflow/dags
   - ./logs:/opt/airflow/logs
   - ./plugins:/opt/airflow/plugins
   - ./data:/opt/airflow/data
 user: "${AIRFLOW UID:-50000}:0"
  depends_on: &airflow-common-depends-on
    postgres:
      condition: service_healthy
services:
  postgres:
    image: postgres:13
    environment:
      POSTGRES USER: airflow
     POSTGRES PASSWORD: airflow
     POSTGRES DB: airflow
    volumes:
      - postgres_db_volume:/var/lib/postgresql/data
    healthcheck:
      test: ["CMD", "pg isready", "-U", "airflow"]
      interval: 5s
      retries: 5
  airflow-webserver:
    <<: *airflow-common
    command: webserver
    ports:
      - 8080:8080
```

version: '3.8'

healthcheck:

```
test: ["CMD", "curl", "--fail", "http://localhost:8080/health"]
    interval: 10s
   timeout: 10s
    retries: 5
airflow-scheduler:
 <<: *airflow-common
 command: scheduler
 healthcheck:
   test: ["CMD-SHELL", 'airflow jobs check --job-type SchedulerJob --hostname "$${H
   interval: 10s
   timeout: 10s
    retries: 5
airflow-init:
 <<: *airflow-common
 entrypoint: /bin/bash
 command:
   − −c
    - |
      function ver() {
        printf "%04d%04d%04d%04d" $${1//./ }
     }
      airflow version=$$(AIRFLOW LOGGING LOGGING LEVEL=INFO && airflow version)
      airflow_version_comparable=$$(ver $${airflow_version})
      min airflow version=2.2.0
      min airflow version comparable=$$(ver $${min airflow version})
      if (( airflow_version_comparable < min_airflow_version_comparable )); then
        echo -e "\033[1;31mERROR!!!: Too old Airflow version $${airflow_version}!\e[|
        exit 1
      fi
      if [[ -z "${AIRFLOW_UID}" ]]; then
        echo -e "\033[1;33mWARNING!!!: AIRFLOW UID not set!\e[0m"
        echo "Setting AIRFLOW UID to 50000"
        export AIRFLOW UID=50000
      fi
      one meg=1048576
      mem_available=$$(($$(getconf _PHYS_PAGES) * $$(getconf PAGE_SIZE) / one_meg))
      cpus_available=$$(grep -cE 'cpu[0-9]+' /proc/stat)
      disk_available=$$(df / | tail -1 | awk '{print $$4}')
      warning resources="false"
      if (( mem available < 4000 )); then
        echo -e "\033[1;33mWARNING!!!: Not enough memory available for Docker.\e[0m"
       warning resources="true"
```

```
fi
        if (( cpus_available < 2 )); then</pre>
          echo -e "\033[1;33mWARNING!!!: Not enough CPUS available for Docker.\e[0m"
         warning resources="true"
        fi
        if (( disk available < one meg )); then</pre>
          echo -e "\033[1;33mWARNING!!!: Not enough Disk space available for Docker.\e
         warning resources="true"
        fi
        if [[ $${warning resources} == "true" ]]; then
          echo -e "\033[1;33mWARNING!!!: You have not enough resources to run Airflow
          echo "Please follow the instructions to increase amount of resources availab"
                   https://airflow.apache.org/docs/apache-airflow/stable/howto/docker-
         echo "
        fi
        mkdir -p /sources/logs /sources/dags /sources/plugins
        chown -R "${AIRFLOW_UID}:0" /sources/{logs,dags,plugins}
        exec /entrypoint airflow version
   environment:
     <<: *airflow-common-env
     _AIRFLOW_DB_UPGRADE: 'true'
     _AIRFLOW_WWW_USER_CREATE: 'true'
     AIRFLOW WWW USER USERNAME: ${ AIRFLOW WWW USER USERNAME:-airflow}
     AIRFLOW WWW USER PASSWORD: ${ AIRFLOW WWW USER PASSWORD:-airflow}
volumes:
  postgres_db_volume:
```

Step 3: Launch Airflow

```
# Set Airflow user ID
echo -e "AIRFLOW_UID=$(id -u)" > .env
# Initialize Airflow
docker-compose up airflow-init
# Start Airflow services
docker-compose up -d
# Access Airflow Web UI
# http://localhost:8080
# Username: airflow
# Password: airflow
```

First Look at Airflow Web UI

Main Interface Sections:

- 1. DAGs View: List of all workflows
 - Status indicators (running, success, failed)
 - Last run information
 - DAG schedule and next run time
- 2. **Graph View**: Visual representation of workflow
 - Task dependencies as connected nodes
 - Task status color coding
 - Interactive task details
- 3. Tree View: Historical runs over time
 - Timeline of DAG executions
 - Task success/failure patterns
 - Easy identification of recurring issues
- 4. Gantt Chart: Task duration analysis
 - Performance bottlenecks identification
 - Resource utilization patterns
 - Optimization opportunities

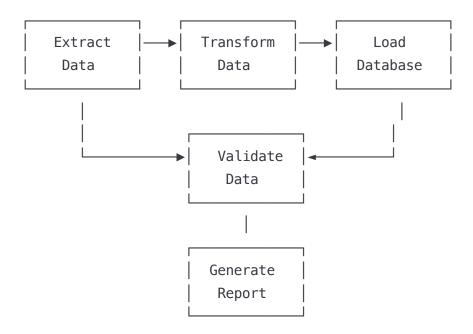
Understanding DAGs: The Workflow Blueprint



DAG Fundamentals (Visual Learning)

What is a DAG?

Simple DAG Example (Customer Analytics Pipeline):



Key DAG Properties:

- **Directed**: Arrows show task execution order
- Acyclic: No circular dependencies
- Connected: Related tasks form a coherent workflow

Creating Your First DAG (Hands-On)

Step 1: Download Sample Data

- 1. Visit: kaggle.com/datasets/imakash3011/customer-personality-analysis
- 2. Download (marketing_campaign.csv)
- Place in (data/customer_data.csv)

Step 2: Basic DAG Structure

Create (dags/customer_analytics_dag.py):

0.0001

Customer Analytics DAG

```
This DAG demonstrates a complete customer analytics workflow:
1. Extract customer data from CSV
2. Validate data quality
3. Transform data (RFM analysis)
4. Load to database
5. Generate insights report
6. Send email notification
Schedule: Daily at 6 AM
Owner: Data Engineering Team
.....
from datetime import datetime, timedelta
from airflow import DAG
from airflow.operators.python operator import PythonOperator
from airflow.operators.bash_operator import BashOperator
from airflow.operators.email_operator import EmailOperator
from airflow.operators.dummy_operator import DummyOperator
# Default arguments for all tasks
default args = {
    'owner': 'data-engineering-team',
    'depends on past': False,
    'start date': datetime(2024, 1, 1),
    'email_on_failure': True,
    'email on retry': False,
    'retries': 2,
    'retry_delay': timedelta(minutes=5),
    'catchup': False # Don't run historical DAGs
}
# Define the DAG
daq = DAG(
    'customer_analytics_pipeline',
    default_args=default_args,
    description='Daily customer analytics processing',
    schedule_interval='0 6 * * * *', # Daily at 6 AM
    max_active_runs=1, # Only one DAG run at a time
    tags=['customer', 'analytics', 'daily']
```

```
# Task functions (we'll implement these)
def extract customer data(**context):
    """Extract customer data from CSV file"""
    print("② Starting data extraction...")
    # Implementation here
    print("✓ Data extraction completed!")
def validate_data_quality(**context):
    """Validate data quality and completeness"""
    print(" Starting data validation...")
    # Implementation here
    print("✓ Data validation completed!")
def transform customer data(**context):
    """Perform RFM analysis and customer segmentation"""
    print("
    Starting data transformation...")
    # Implementation here
    print("▼ Data transformation completed!")
def load to database(**context):
    """Load processed data to PostgreSQL"""
    print(" Starting data loading...")
    # Implementation here
    print("▼ Data loading completed!")
def generate_insights_report(**context):
    """Generate customer insights report"""
    print("
    Generating insights report...")
    # Implementation here
    print("▼ Report generation completed!")
# Define tasks
start task = DummyOperator(
    task_id='start_pipeline',
    dag=dag
)
extract_task = PythonOperator(
    task_id='extract_customer_data',
    python callable=extract customer data,
    dag=dag
)
```

)

```
validate_task = PythonOperator(
            task id='validate data quality',
           python callable=validate data quality,
           dag=dag
)
transform_task = PythonOperator(
           task_id='transform_customer_data',
           python_callable=transform_customer_data,
           dag=dag
)
load task = PythonOperator(
           task id='load to database',
           python_callable=load_to_database,
           dag=dag
)
report_task = PythonOperator(
           task_id='generate_insights_report',
           python callable=generate insights report,
           dag=dag
)
# Health check task
health check = BashOperator(
           task_id='health_check',
           bash_command='echo "Pipeline health check passed ✓"',
           dag=dag
)
# Email notification on success
success email = EmailOperator(
           task id='send success email',
           to=['data-team@company.com'],
            subject='Customer Analytics Pipeline - Success',
           html content="""
           <h3>Customer Analytics Pipeline Completed Successfully</h3>
           The daily customer analytics pipeline has completed successfully.
           <strong>Execution Date:</strong> {{ ds }}
           <strong>Total Runtime:</strong> {{ macros.timedelta(dag_run.end_date - dag_run.end_date - dag_run.end_dat
           dag=dag
```

```
end_task = DummyOperator(
    task_id='end_pipeline',
    dag=dag
)

# Define task dependencies
start_task >> extract_task >> validate_task >> transform_task >> load_task >> report_task
```

STATE OF THE PROOF OF THE PARK OF THE PAR

Step 1: Refresh DAGs

- 1. Open Airflow UI (http://localhost:8080)
- 2. Look for "customer_analytics_pipeline" in DAGs list
- 3. Toggle the DAG "ON" (unpause it)

Step 2: Explore Graph View

- 1. Click on the DAG name
- 2. Go to "Graph View" tab
- 3. Observe the visual workflow representation
- 4. Click on individual tasks to see details

Step 3: Trigger Manual Run

- 1. Click "Trigger DAG" button (play icon)
- 2. Watch tasks execute in real-time
- 3. Observe color changes:

• White: Not started

• Yellow: Running

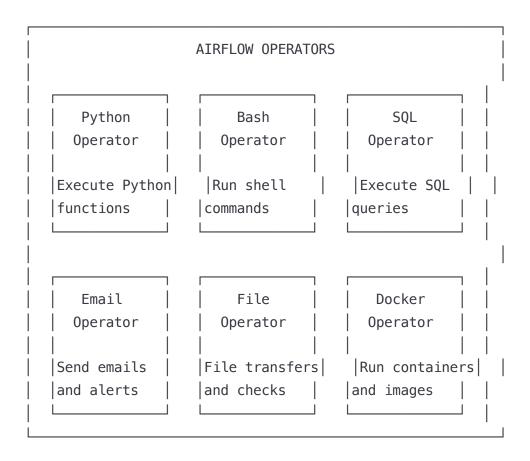
• Green: Success

• **Red**: Failed

X Airflow Operators: Building Blocks of Workflows

Understanding Operators (Concept First)

Operators are task templates that define what to do:



© Common Operators for Data Engineering

1. PythonOperator - Most versatile for data processing

```
def process_customer_data(**context):
    import pandas as pd
    # Your data processing logic
    df = pd.read_csv('/opt/airflow/data/customer_data.csv')
    # Process the data
    return df.shape[0] # Return number of records processed

python_task = PythonOperator(
    task_id='process_data',
    python_callable=process_customer_data,
    dag=dag
)
```

2. BashOperator - For shell commands and scripts

```
python

bash_task = BashOperator(
    task_id='run_data_quality_check',
    bash_command='python /opt/airflow/scripts/data_quality.py',
    dag=dag
)
```

3. SqlOperator - For database operations

```
python
from airflow.providers.postgres.operators.postgres import PostgresOperator
sql_task = PostgresOperator(
    task_id='create_customer_table',
    postgres_conn_id='postgres_default',
    sql="""
    CREATE TABLE IF NOT EXISTS customer_segments (
        customer_id VARCHAR(50),
        segment VARCHAR(50),
        rfm_score VARCHAR(10),
        created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
    );
    """,
    dag=dag
)
```

4. EmailOperator - For notifications

```
python

email_task = EmailOperator(
    task_id='send_report',
    to=['team@company.com'],
    subject='Daily Customer Report',
    html_content='<h1>Report attached</h1>',
    dag=dag
)
```

Mands-On Operator Implementation

Let's implement the customer analytics DAG with real functionality:

```
.....
Complete Customer Analytics DAG Implementation
.....
import pandas as pd
import numpy as np
from datetime import datetime, timedelta
from sqlalchemy import create_engine
from airflow import DAG
from airflow.operators.python_operator import PythonOperator
from airflow.operators.bash_operator import BashOperator
from airflow.providers.postgres.operators.postgres import PostgresOperator
# DAG configuration
default args = {
    'owner': 'data-team',
    'depends_on_past': False,
    'start_date': datetime(2024, 1, 1),
    'email_on_failure': True,
    'retries': 2,
    'retry_delay': timedelta(minutes=5)
}
dag = DAG(
    'customer_analytics_complete',
    default_args=default_args,
    description='Complete customer analytics pipeline',
    schedule_interval='0 6 * * * *',
    catchup=False,
    tags=['customer', 'rfm', 'analytics']
)
def extract_and_validate_data(**context):
    """Extract customer data and perform basic validation"""
    print("② Starting data extraction and validation...")
    # Load data
    df = pd.read_csv('/opt/airflow/data/customer_data.csv')
    # Basic validation
    validation_results = {
        'total records': len(df),
```

'missing values': df.isnull().sum().sum(),

```
'duplicate_records': df.duplicated().sum(),
        'date range': f"{df['Dt Customer'].min()} to {df['Dt Customer'].max()}"
   }
    print(f"

Validation Results: {validation results}")
   # Store validation results for downstream tasks
   context['task_instance'].xcom_push(key='validation_results', value=validation_resu
   # Save cleaned data
   df clean = df.dropna().drop duplicates()
   df_clean.to_csv('/opt/airflow/data/customer_data_clean.csv', index=False)
    print("✓ Data extraction and validation completed!")
    return validation results
def perform rfm analysis(**context):
    """Perform RFM (Recency, Frequency, Monetary) analysis"""
   print("
    Starting RFM analysis...")
   # Load cleaned data
   df = pd.read csv('/opt/airflow/data/customer data clean.csv')
   # Convert date column
   df['Dt Customer'] = pd.to datetime(df['Dt Customer'])
   # Calculate current date for recency
    current date = df['Dt Customer'].max() + pd.Timedelta(days=1)
   # Calculate RFM metrics
    rfm = df.groupby('ID').agg({
        'Dt_Customer': lambda x: (current_date - x.max()).days, # Recency
        'NumTotalPurchases': 'sum', # Frequency
        'MntTotal': 'sum' # Monetary
   }).reset index()
    rfm.columns = ['CustomerID', 'Recency', 'Frequency', 'Monetary']
   # Create RFM scores (quintiles)
    rfm['R_Score'] = pd.qcut(rfm['Recency'], 5, labels=[5,4,3,2,1])
    rfm['F_Score'] = pd.qcut(rfm['Frequency'].rank(method='first'), 5, labels=[1,2,3,4
    rfm['M Score'] = pd.qcut(rfm['Monetary'], 5, labels=[1,2,3,4,5])
   # Combine scores
```

```
rfm['RFM_Score'] = rfm['R_Score'].astype(str) + rfm['F_Score'].astype(str) + rfm['I
   # Customer segmentation
   def segment customers(score):
        if score in ['555', '554', '544', '545', '454', '455', '445']:
            return 'Champions'
        elif score in ['543', '444', '435', '355', '354', '345', '344', '335']:
            return 'Loyal Customers'
        elif score in ['553', '551', '552', '541', '542', '533', '532', '531']:
            return 'Potential Loyalists'
       elif score in ['512', '511', '422', '421', '412', '411', '311']:
            return 'New Customers'
       elif score in ['155', '154', '144', '214', '215', '115', '114']:
            return 'At Risk'
       else:
            return 'Others'
    rfm['Segment'] = rfm['RFM_Score'].apply(segment_customers)
   # Save results
    rfm.to_csv('/opt/airflow/data/customer_rfm_analysis.csv', index=False)
   # Generate summary statistics
    segment summary = rfm.groupby('Segment').agg({
        'CustomerID': 'count',
        'Recency': 'mean',
        'Frequency': 'mean',
        'Monetary': ['mean', 'sum']
   }).round(2)
    segment_summary.to_csv('/opt/airflow/data/segment_summary.csv')
   print(f" ■ RFM Analysis completed. Segments identified: {rfm['Segment'].value coun
   # Store summary for reporting
    context['task_instance'].xcom_push(key='segment_summary', value=segment_summary.to
    return rfm.shape[0]
def generate_business_insights(**context):
    """Generate actionable business insights from RFM analysis"""
   print("I Generating business insights...")
   # Load RFM results
```

```
rfm = pd.read_csv('/opt/airflow/data/customer_rfm_analysis.csv')
    # Calculate kev business metrics
    insights = {
        'total customers': len(rfm),
        'total revenue': rfm['Monetary'].sum(),
        'avg customer value': rfm['Monetary'].mean(),
        'champions_count': len(rfm[rfm['Segment'] == 'Champions']),
        'at risk count': len(rfm[rfm['Segment'] == 'At Risk']),
        'champions_revenue_share': rfm[rfm['Segment'] == 'Champions']['Monetary'].sum(
   }
   # Generate recommendations
    recommendations = []
    if insights['champions revenue share'] > 30:
        recommendations.append("High champion revenue share - focus on retention progra
    if insights['at_risk_count'] > insights['total_customers'] * 0.2:
        recommendations.append("High at-risk customers - implement win-back campaigns"
   # Create insights report
    report = {
        'execution date': context['ds'],
        'insights': insights,
        'recommendations': recommendations
   }
   # Save report
    import json
   with open('/opt/airflow/data/business_insights_report.json', 'w') as f:
        json.dump(report, f, indent=2, default=str)
   print("▼ Business insights generated successfully!")
    context['task_instance'].xcom_push(key='business_insights', value=insights)
    return insights
# Create PostgreSQL table
create_table_task = PostgresOperator(
   task id='create customer segments table',
    postgres conn id='postgres default',
    sql="""
```

```
CREATE TABLE IF NOT EXISTS customer_segments (
        customer_id VARCHAR(50) PRIMARY KEY,
        recency INTEGER,
        frequency INTEGER,
        monetary DECIMAL(10,2),
        rfm score VARCHAR(10),
        segment VARCHAR(50),
        created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
   );
   CREATE TABLE IF NOT EXISTS segment summary (
        segment VARCHAR(50) PRIMARY KEY,
        customer count INTEGER,
        avg recency DECIMAL(10,2),
        avg frequency DECIMAL(10,2),
        avg_monetary DECIMAL(10,2),
        total revenue DECIMAL(15,2),
        updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
    );
   dag=dag
def load data to postgres(**context):
   """Load RFM analysis results to PostgreSQL"""
   print("H Loading data to PostgreSQL...")
   # Database connection
   engine = create_engine('postgresql://airflow:airflow@postgres:5432/airflow')
   # Load RFM data
    rfm = pd.read_csv('/opt/airflow/data/customer_rfm_analysis.csv')
   # Load to database
    rfm.to_sql('customer_segments', engine, if_exists='replace', index=False, method='
   # Load segment summary
    segment_summary = pd.read_csv('/opt/airflow/data/segment_summary.csv')
    segment_summary.to_sql('segment_summary', engine, if_exists='replace', index=False
    print(f" Loaded {len(rfm)} customer records to PostgreSQL!")
    return len(rfm)
```

)

```
# Define all tasks
extract_validate_task = PythonOperator(
    task id='extract and validate data',
   python callable=extract and validate data,
   dag=dag
)
rfm_analysis_task = PythonOperator(
    task_id='perform_rfm_analysis',
   python_callable=perform_rfm_analysis,
   dag=dag
)
insights task = PythonOperator(
   task id='generate business insights',
   python_callable=generate_business_insights,
   dag=dag
)
load_postgres_task = PythonOperator(
   task_id='load_data_to_postgres',
   python callable=load data to postgres,
   dag=dag
)
# Data quality check
quality check task = BashOperator(
   task_id='data_quality_check',
   bash_command="""
   echo "Running data quality checks..."
    if [ -f /opt/airflow/data/customer_rfm_analysis.csv ]; then
        echo "☑ RFM analysis file exists"
        record count=$(wc -l < /opt/airflow/data/customer rfm analysis.csv)
        echo "■ Record count: $record count"
        if [ $record count -qt 1 ]; then
            echo " Data quality check passed"
        else
            echo "X Data quality check failed - insufficient records"
            exit 1
        fi
   else
        echo "X Data quality check failed - RFM file not found"
       exit 1
    fi
```

```
dag=dag
)

# Define task dependencies
extract_validate_task >> rfm_analysis_task >> insights_task >> create_table_task >> log
```

Task Dependencies and Scheduling Concepts

Understanding Dependencies (Visual Approach)

Dependency Types:

Solution Implementing Dependencies in Code

Multiple Dependency Syntax Options:

```
# Method 1: Using >> and << operators
task_a >> task_b >> task_c # Linear

# Method 2: Using set_downstream/set_upstream
task_a.set_downstream([task_b, task_c]) # Fan-out
[task_b, task_c].set_downstream(task_d) # Fan-in

# Method 3: Using lists
task_a >> [task_b, task_c] >> task_d # Diamond pattern

# Method 4: Complex dependencies
task_a >> task_b >> task_d
task_a >> task_c >> task_d
```

Scheduling Concepts (Cron and Beyond)

Understanding Schedule Intervals:

```
# Common scheduling patterns
dag = DAG(
    'my_dag',
    schedule_interval='@daily',  # Every day at midnight
    # schedule_interval='@hourly',  # Every hour
    # schedule_interval='0 6 * * *', # Every day at 6 AM
    # schedule_interval='0 6 * * 1', # Every Monday at 6 AM
    # schedule_interval=timedelta(hours=2), # Every 2 hours
    # schedule_interval=None,  # Manual trigger only
)
```

Cron Expression Breakdown:

```
0 6 * * *
| | | | | |
| | | | Day of week (0-7, both 0 and 7 are Sunday)
| | | Month (1-12)
| | Day of month (1-31)
| Hour (0-23)
| Minute (0-59)
```

Visual Cron Examples:

```
'0 6 * * * * ' = Daily at 6:00 AM
'0 */2 * * * ' = Every 2 hours
'0 6 * * 1' = Every Monday at 6:00 AM
'0 6 1 * * ' = First day of every month at 6:00 AM
'0 6 1 1 * ' = January 1st at 6:00 AM (yearly)
```

Airflow Web UI Mastery

UI Navigation and Monitoring

DAGs View Overview:

DAG Name	Status	Last Run	Next Run	Actions
customer_analytics data_validation	×	05:30:00	05:30:00	

Status Indicators:

- Green (☑): All tasks successful
- Red (X): At least one task failed
- Yellow (): Currently running
- Purple (): Upstream failed
- **Gray**: Not yet started

© Graph View Deep Dive

Interactive Graph Features:

1. Task Status Colors:

• White: Not started

• Light Green: Queued

• Yellow: Running

• Dark Green: Success

• **Red**: Failed

• **Orange**: Retry

• Purple: Skipped

2. Task Details Panel:

- Click any task to see:
 - Task Instance Details
 - Logs
 - XCom (cross-communication) data
 - Task Duration
 - Retry History

3. Zoom and Pan:

- Mouse wheel to zoom
- Click and drag to pan
- Fit to screen button

Tree View for Historical Analysis

Tree View Benefits:

Task Timeline View:

	Today	Yesterday	2 Days Ago
extract_data	\checkmark	\checkmark	×
validate_data	\checkmark	\checkmark	I
transform_data	$\overline{\checkmark}$	$\overline{\checkmark}$	I
load_data	$\overline{\checkmark}$	\checkmark	I
generate_report		\checkmark	I

Pattern Recognition:

- Consistent failures: Infrastructure issues
- Intermittent failures: Data quality issues
- Cascading failures: Dependency problems
- Performance degradation: Resource constraints

Gantt Chart for Performance Analysis

Performance Insights:

XCom: Task Communication

Understanding XCom (Cross-Communication)

Concept: Tasks sharing data with each other

```
python
def upstream task(**context):
    """Task that produces data"""
    result = {"processed_records": 1000, "quality_score": 0.95}
    # Push data to XCom
    context['task instance'].xcom push(key='processing results', value=result)
    return result
def downstream task(**context):
    """Task that consumes data"""
    # Pull data from XCom
    results = context['task_instance'].xcom_pull(
        task_ids='upstream_task',
        key='processing results'
    )
    print(f"Processing {results['processed records']} records")
    print(f"Quality score: {results['quality_score']}")
```

STATE OF STATE OF ST

1. Keep Data Small:

```
python

#  Good: Small metadata

xcom_push(key='record_count', value=1000)

#  Bad: Large datasets

# xcom_push(key='full_dataset', value=huge_dataframe)
```

2. Use for Coordination, Not Data Transfer:

```
python

#  Good: Status and metrics

return {
    'status': 'success',
    'records_processed': len(df),
    'file_path': '/tmp/processed_data.csv'
}

#  X Bad: Actual data transfer
# return df # Don't pass DataFrames via XCom
```

3. Error Handling:

```
def consume_xcom_data(**context):
    try:
        data = context['task_instance'].xcom_pull(task_ids='producer_task')
        if data is None:
            raise ValueError("No data received from upstream task")
        # Process data
    except Exception as e:
        print(f"Error processing XCom data: {e}")
        raise
```

- **§** Error Handling and Retries
- **Retry Strategies**

Configuring Retries:

```
python
default_args = {
    'retries': 3,
                                            # Number of retries
    'retry_delay': timedelta(minutes=5),
                                          # Wait between retries
    'retry_exponential_backoff': True,  # Exponential backoff
    'max_retry_delay': timedelta(hours=1), # Maximum retry delay
}
# Task-specific retry configuration
risky_task = PythonOperator(
    task_id='risky_operation',
    python_callable=some_function,
    retries=5, # Override default
    retry_delay=timedelta(minutes=10),
    dag=dag
)
```

Email Alerts and Notifications

Email Configuration:

```
python
default_args = {
    'email': ['data-team@company.com', 'devops@company.com'],
    'email on failure': True,
    'email on retry': False,
    'email_on_success': False, # Usually too noisy
}
# Custom email content
def task_fail_slack_alert(context):
    """Send custom Slack alert on task failure"""
    slack_msg = f"""
    :red_circle: Task Failed
   *Task*: {context.get('task instance').task id}
    *DAG*: {context.get('task instance').dag id}
    *Execution Time*: {context.get('ds')}
   *Log*: {context.get('task_instance').log_url}
    # Send to Slack (implementation depends on your setup)
task_with_alerts = PythonOperator(
    task_id='important_task',
    python callable=important function,
    on_failure_callback=task_fail_slack_alert,
    dag=dag
)
```

Failure Handling Strategies

1. Task-Level Handling:

```
def robust_data_processing(**context):
    try:
        # Main processing logic
        df = pd.read_csv('/path/to/data.csv')
        processed_df = transform_data(df)
        processed_df.to_csv('/path/to/output.csv')

except FileNotFoundError:
    # Handle missing input gracefully
    print("Input file not found, skipping processing")
    return "skipped"

except Exception as e:
    # Log error and re-raise for Airflow to handle
```

print(f"Unexpected error in data processing: {e}")

2. Skip Tasks on Upstream Failure:

raise

```
from airflow.operators.dummy_operator import DummyOperator
from airflow.utils.trigger_rule import TriggerRule

cleanup_task = PythonOperator(
    task_id='cleanup',
    python_callable=cleanup_function,
    trigger_rule=TriggerRule.ALL_DONE, # Run regardless of upstream success/failure
    dag=dag
)
```

3. Branching for Error Recovery:

```
from airflow.operators.python_operator import BranchPythonOperator
def check data quality(**context):
    """Branch based on data quality results"""
    quality_score = check_quality()
    if quality_score > 0.9:
        return 'high_quality_processing'
    else:
        return 'data_cleaning_task'
branching_task = BranchPythonOperator(
    task_id='quality_check_branch',
    python_callable=check_data_quality,
    dag=dag
)
high_quality_task = PythonOperator(
    task_id='high_quality_processing',
    python_callable=standard_processing,
    dag=dag
)
cleaning_task = PythonOperator(
    task_id='data_cleaning_task',
    python_callable=intensive_cleaning,
    dag=dag
)
branching_task >> [high_quality_task, cleaning_task]
```

Airflow Connections and Variables

Managing Connections (UI Approach)

Setting Up Database Connection:

- 1. Access Admin Menu:
 - Go to Admin → Connections in Airflow UI
- 2. Add New Connection:
 - Conn ld: postgres_customer_db

- Conn Type: (Postgres)
- Host: (postgres)
- Schema: (customer_analytics)
- Login: (airflow)
- Password: (airflow)
- **Port**: (5432)

3. Test Connection:

Click "Test" button to verify connectivity

Using Connections in DAGs:

```
python
from airflow.providers.postgres.hooks.postgres import PostgresHook

def load_data_with_connection(**context):
    """Load data using Airflow connection"""
    # Get connection
    postgres_hook = PostgresHook(postgres_conn_id='postgres_customer_db')

# Execute SQL
    postgres_hook.run("""
        INSERT INTO customer_segments (customer_id, segment, rfm_score)
        VALUES (%s, %s, %s)
    """, parameters=('CUST001', 'Champion', '555'))

print("V Data loaded using Airflow connection!")
```

Nanaging Variables (Configuration)

Setting Variables via UI:

- 1. Go to Admin → Variables
- 2. Add Key-Value Pairs:
 - (data_source_path): (/opt/airflow/data/customer_data.csv)
 - (min_quality_threshold): (0.95)
 - (notification_email): (data-team@company.com)

Using Variables in DAGs:

```
python
```

```
def configurable_task(**context):
    """Task using Airflow Variables for configuration"""
    # Get variables
    data_path = Variable.get('data_source_path')
    quality_threshold = float(Variable.get('min_quality_threshold'))
    notification_email = Variable.get('notification_email')

    print(f"Processing data from: {data_path}")
    print(f"Quality threshold: {quality_threshold}")

# Use in processing logic
    df = pd.read_csv(data_path)
    quality_score = calculate_quality(df)

if quality_score < quality_threshold:
    send_alert(notification_email, f"Quality below threshold: {quality_score}")</pre>
```

Solution Advanced DAG Patterns

💢 Dynamic DAG Generation

Problem: Creating similar DAGs for different datasets or environments

Solution: Generate DAGs programmatically

```
Dynamic DAG Generation Example
Generate customer analytics DAGs for different regions
.....
from airflow import DAG
from airflow.operators.python_operator import PythonOperator
from datetime import datetime, timedelta
# Configuration for different regions
REGIONS = ['north_america', 'europe', 'asia_pacific']
def create_region_dag(region_name):
    """Factory function to create DAG for specific region"""
    default_args = {
        'owner': 'data-team',
        'depends_on_past': False,
        'start_date': datetime(2024, 1, 1),
        'retries': 2,
        'retry_delay': timedelta(minutes=5)
    }
    dag = DAG(
        f'customer analytics {region name}',
        default args=default args,
        description=f'Customer analytics for {region_name}',
        schedule_interval='0 6 * * * *',
        catchup=False,
        tags=['customer', 'analytics', region_name]
    )
    def process_region_data(**context):
        print(f"Processing customer data for {region name}")
        # Region-specific processing logic
        data_path = f'/opt/airflow/data/{region_name}_customers.csv'
        # Process data...
    process_task = PythonOperator(
        task_id=f'process_{region_name}_data',
        python_callable=process_region_data,
        dag=dag
```

)

return dag

```
# Generate DAGs for each region
for region in REGIONS:
   globals()[f'customer_analytics_{region}'] = create_region_dag(region)
```

SubDAGs for Reusable Workflows

Concept: Encapsulate common workflow patterns

```
from airflow.operators.subdag_operator import SubDagOperator
from airflow.models import DAG
def create data quality subdag(parent dag id, child dag id, schedule interval, default
    """Reusable data quality checking workflow"""
    subdag = DAG(
        f'{parent_dag_id}.{child_dag_id}',
        default_args=default_args,
        schedule_interval=schedule_interval,
    )
    # Data quality tasks
    completeness check = PythonOperator(
        task id='check completeness',
        python_callable=check_data_completeness,
        dag=subdag
    )
    accuracy_check = PythonOperator(
        task_id='check_accuracy',
        python_callable=check_data_accuracy,
        dag=subdag
    )
    consistency_check = PythonOperator(
        task_id='check_consistency',
        python_callable=check_data_consistency,
        dag=subdag
    )
    # Set dependencies
    [completeness_check, accuracy_check, consistency_check]
    return subdag
# Use SubDAG in main DAG
main_dag = DAG(...)
quality_check_subdag = SubDagOperator(
    task_id='data_quality_checks',
    subdag=create_data_quality_subdag(
        parent dag id='customer analytics main',
```

```
child_dag_id='data_quality_checks',
    schedule_interval=main_dag.schedule_interval,
    default_args=main_dag.default_args
),
    dag=main_dag
)
```

♦ TaskGroups for Better Organization

Modern Alternative to SubDAGs:

```
from airflow.utils.task_group import TaskGroup
with DAG('customer analytics with groups', ...) as dag:
    # Data ingestion group
    with TaskGroup('data_ingestion') as ingestion_group:
        extract_customers = PythonOperator(
            task_id='extract_customers',
            python_callable=extract_customer_data
        )
        extract_transactions = PythonOperator(
            task_id='extract_transactions',
            python callable=extract transaction data
        )
        validate_data = PythonOperator(
            task_id='validate_data',
            python_callable=validate_extracted_data
        )
        [extract_customers, extract_transactions] >> validate_data
    # Processing group
    with TaskGroup('data_processing') as processing_group:
        clean_data = PythonOperator(
            task_id='clean_data',
            python_callable=clean_customer_data
        )
        rfm_analysis = PythonOperator(
            task_id='rfm_analysis',
            python_callable=perform_rfm_analysis
        )
        segment_customers = PythonOperator(
            task_id='segment_customers',
            python_callable=segment_customers_func
        )
        clean_data >> rfm_analysis >> segment_customers
```

```
with TaskGroup('reporting') as reporting_group:
    generate_dashboard = PythonOperator(
        task_id='generate_dashboard',
        python_callable=create_dashboard
)

send_email_report = EmailOperator(
        task_id='send_email_report',
        to=['stakeholders@company.com'],
        subject='Daily Customer Analytics Report',
        html_content='Dashboard updated successfully!'
)

generate_dashboard >> send_email_report

# Define group dependencies
ingestion_group >> processing_group >> reporting_group
```

Monitoring and Logging

Built-in Monitoring Features

DAG Performance Metrics:

- **Duration Trends**: Track pipeline execution time
- Success Rate: Percentage of successful runs
- **Task Distribution**: Which tasks take longest
- Resource Usage: CPU and memory consumption

Accessing Metrics in UI:

- 1. Browse → DAG Runs: Historical execution data
- 2. Browse → Task Instances: Individual task performance
- 3. **Browse** → **Jobs**: Scheduler and executor status

Structured Logging

Best Practices for Logging:

```
python
import logging
from airflow.utils.log.logging_mixin import LoggingMixin
class CustomerAnalyticsProcessor(LoggingMixin):
    """Class with proper logging for Airflow"""
   def process_customers(self, **context):
        """Process customer data with structured logging"""
       # Use Airflow's logging
        self.log.info("Starting customer data processing")
        try:
            # Log processing steps
            self.log.info("Loading customer data from CSV")
            df = pd.read csv('/opt/airflow/data/customer data.csv')
            self.log.info(f"Loaded {len(df)} customer records")
            # Process data
            processed_df = self.transform_data(df)
            self.log.info(f"Processed {len(processed df)} customer records")
            self.log.info("Customer data processing completed successfully")
            return len(processed_df)
        except Exception as e:
            self.log.error(f"Error processing customer data: {str(e)}")
            self.log.error(f"Error type: {type(e).__name__}")
            raise
def airflow task with logging(**context):
   """Airflow task function with proper logging"""
   processor = CustomerAnalyticsProcessor()
```

Custom Alerting

Slack Integration Example:

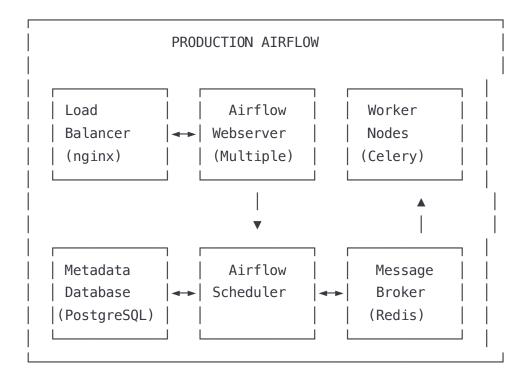
return processor.process_customers(**context)

```
from airflow.providers.slack.operators.slack_webhook import SlackWebhookOperator
def send success alert(**context):
    """Send custom success notification"""
    # Get task instance details
    ti = context['task instance']
    dag_run = context['dag_run']
    # Prepare message
    message = f'''''
    *Pipeline Success*
    *DAG*: {ti.dag_id}
    *Execution Date*: {context['ds']}
    *Duration*: {dag_run.end_date - dag_run.start_date}
    *Tasks Completed*: All tasks successful
    Dashboard: http://localhost:8080/dags/{ti.dag_id}/grid
    .....
    slack alert = SlackWebhookOperator(
        task_id='slack_success_alert',
        http_conn_id='slack_webhook',
        message=message,
        dag=dag
    )
    return slack_alert.execute(context=context)
# Add to DAG
success_alert_task = PythonOperator(
    task id='send success alert',
    python_callable=send_success_alert,
    trigger_rule='all_success',
    dag=dag
)
```

Production Deployment Considerations

Airflow Architecture for Production

Components in Production:



Configuration Management

Environment-Specific Settings:

```
python
import os
from airflow import DAG
# Environment-aware configuration
ENVIRONMENT = os.getenv('AIRFLOW_ENV', 'development')
if ENVIRONMENT == 'production':
    DEFAULT ARGS = {
        'retries': 3,
        'retry_delay': timedelta(minutes=10),
        'email_on_failure': True,
        'email': ['data-team@company.com', 'ops@company.com']
    }
    SCHEDULE_INTERVAL = '0 6 * * * * # Daily at 6 AM
elif ENVIRONMENT == 'staging':
    DEFAULT_ARGS = {
        'retries': 2,
        'retry_delay': timedelta(minutes=5),
        'email_on_failure': True,
        'email': ['data-team@company.com']
    }
    SCHEDULE_INTERVAL = '0 8 * * * * # Daily at 8 AM
else: # development
    DEFAULT ARGS = {
        'retries': 1,
        'retry_delay': timedelta(minutes=1),
        'email_on_failure': False
    }
    SCHEDULE_INTERVAL = None # Manual trigger only
dag = DAG(
    'customer_analytics_env_aware',
    default_args=DEFAULT_ARGS,
    schedule_interval=SCHEDULE_INTERVAL,
    tags=[ENVIRONMENT, 'customer', 'analytics']
)
```

§ Security Best Practices

1. Connection Security:

- Store sensitive credentials in Airflow Connections
- Use environment variables for secrets
- Enable encryption for metadata database

2. DAG Security:

```
# Secure file access
def secure_file_processing(**context):
    """Process files with security checks"""

file_path = context['params'].get('file_path')

# Validate file path
    if not file_path.startswith('/opt/airflow/data/'):
        raise ValueError("Unauthorized file path access")

# Check file permissions
    if not os.access(file_path, os.R_OK):
        raise PermissionError("Insufficient file permissions")

# Process file securely
    process_file(file_path)
```

3. Resource Limits:

```
from airflow.operators.python_operator import PythonOperator
from airflow.configuration import conf

# Set resource limits
task_with_limits = PythonOperator(
    task_id='resource_limited_task',
    python_callable=memory_intensive_function,
    pool='memory_intensive_pool', # Use resource pool
    dag=dag
)
```



Official Documentation

- Apache Airflow Documentation: https://airflow.apache.org/docs/
- Airflow Best Practices: https://airflow.apache.org/docs/apache-airflow/stable/best-practices.html
- Airflow Operators: https://airflow.apache.org/docs/apache-airflow/stable/operators-and-hooks