

Day 14: Week 2 Integration Project - Building Your First End-to-End Data Platform

What You'll Build Today (Systems-First Approach)

Primary Focus: Integrating all Week 2 technologies into a cohesive data platform **Secondary Focus:** Production deployment patterns and monitoring strategies

Dataset Integration: Multi-source data processing combining all previous datasets

Learning Philosophy for Day 14

"Architecture is about understanding the whole before optimizing the parts"

We'll start with system design principles, understand component interactions, design data flow architecture, and build a production-ready integrated platform.

The Integration Challenge: From Tools to Platform

The Problem: Tool Fragmentation in Data Engineering

Scenario: You've learned individual tools but need to build a real business solution...

Without Integration (Tool Chaos):

- ✗ Monday 6 AM: Manually start PostgreSQL
- ✗ Monday 6:15 AM: Check if Docker containers are running
- ✗ Monday 6:30 AM: Manually trigger Airflow DAG
- ✗ Monday 6:45 AM: Monitor Spark job logs separately
- ✗ Monday 7:00 AM: Check data quality in different tools
- ✗ Monday 7:15 AM: Manually verify database loads
- ✗ Monday 7:30 AM: Pray everything worked together

Problems:

- No unified monitoring across tools
- Manual coordination between systems
- Failure in one component breaks everything
- Impossible to track data lineage
- No automated recovery mechanisms
- Difficult to scale or reproduce

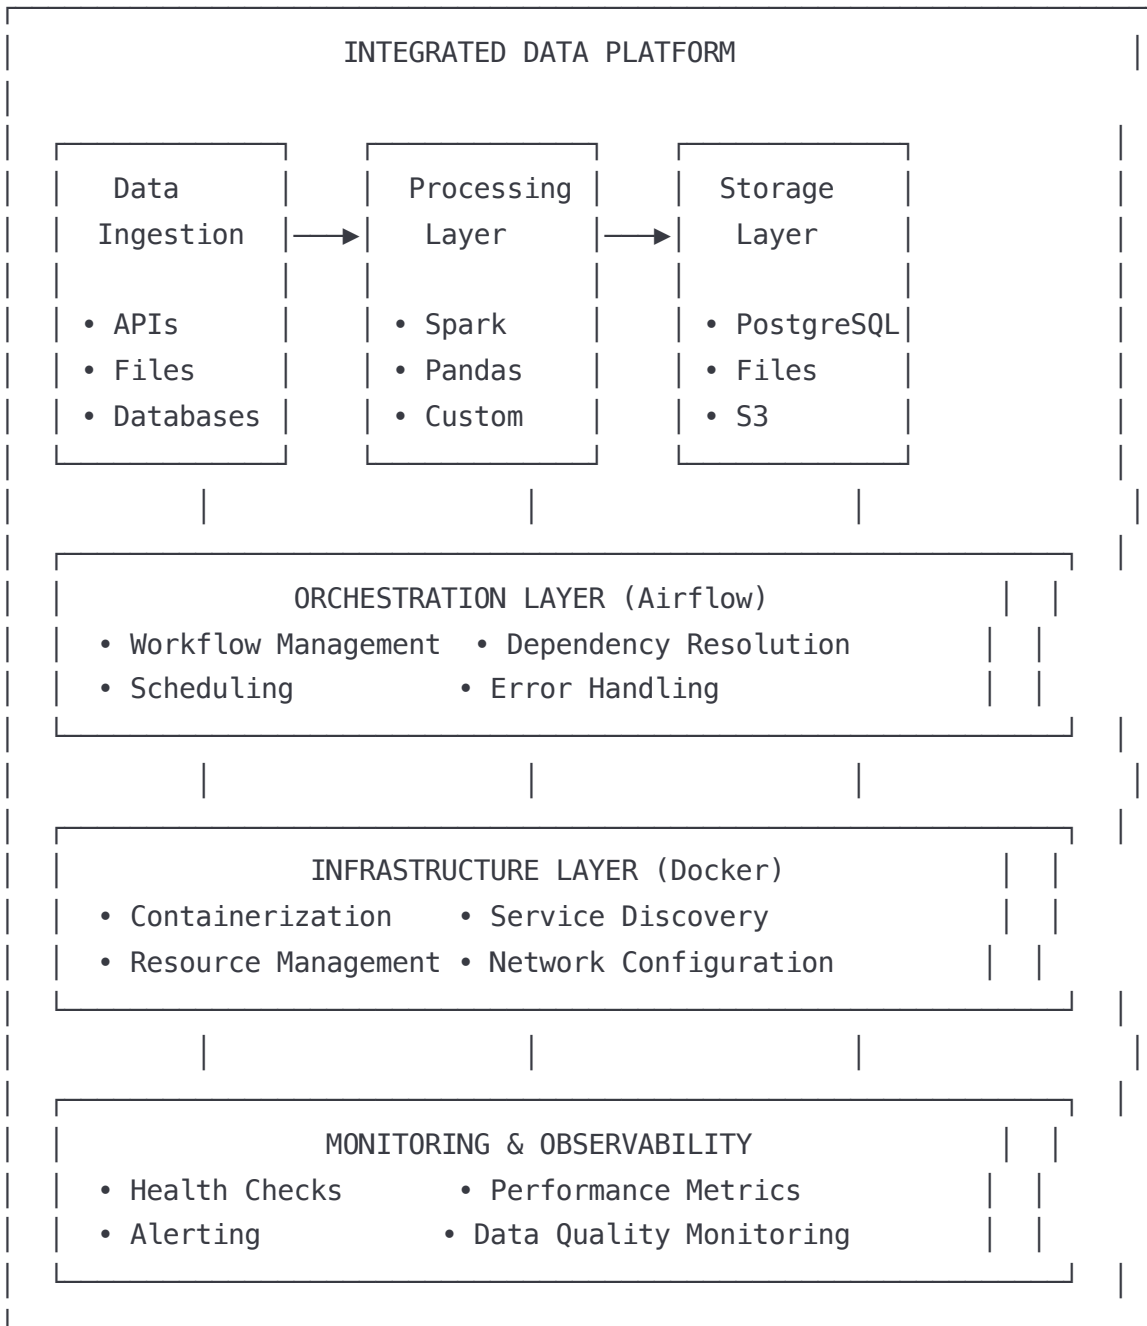
💡 The Integrated Platform Solution

With Platform Integration:

- ✅ Automated infrastructure deployment
- ✅ Orchestrated cross-tool workflows
- ✅ Unified monitoring and alerting
- ✅ Automatic failure recovery
- ✅ End-to-end data lineage tracking
- ✅ Scalable, reproducible architecture
- ✅ Single command deployment and operation

🏗️ Platform Architecture Design (Visual Approach)

🎨 The Integrated Data Platform Mental Model



Platform Components and Interactions

1. Infrastructure Layer (Foundation)

- Docker Compose for service orchestration
- Network configuration and service discovery
- Volume management for data persistence
- Resource allocation and scaling

2. Data Storage Layer

- PostgreSQL for structured analytics
- File system for raw and processed data
- Staging areas for data transformations
- Backup and recovery mechanisms

3. Processing Layer

- Apache Spark for large-scale processing
- Pandas for data manipulation
- Custom Python modules for business logic
- Data quality validation frameworks

4. Orchestration Layer

- Airflow DAGs for workflow management
- Cross-service dependency management
- Scheduling and trigger mechanisms
- Error handling and recovery

5. Monitoring Layer

- Health checks across all services
- Performance metrics collection
- Data quality monitoring
- Alerting and notification systems

Platform Setup and Architecture (Infrastructure-First)

Complete Platform Docker Compose

Project Structure:

```
week2-integration-project/
├─ docker-compose.yml
├─ airflow/
│   ├── dags/
│   │   └─ integrated_data_platform_dag.py
│   ├── plugins/
│   ├── logs/
│   └─ requirements.txt
├─ spark/
│   ├── apps/
│   │   └─ customer_analytics_spark.py
│   └─ conf/
├─ data/
│   ├── raw/
│   │   ├── ecommerce_transactions.csv
│   │   ├── superstore_sales.csv
│   │   ├── amazon_products.csv
│   │   └─ retail_analytics.csv
│   ├── staging/
│   └─ processed/
├─ sql/
│   ├── schema/
│   │   └─ create_tables.sql
│   └─ queries/
├─ scripts/
│   ├── data_quality/
│   │   └─ quality_checks.py
│   ├── ingestion/
│   │   └─ multi_source_ingestion.py
│   ├── monitoring/
│   │   └─ health_checks.py
├─ config/
│   ├── spark-defaults.conf
│   └─ postgres.conf
└─ monitoring/
    ├── grafana/
    └─ prometheus/
```

Complete docker-compose.yml:

yaml

version: '3.8'

Shared environment variables

x-airflow-common:

&airflow-common

image: apache/airflow:2.7.1-python3.10

environment: &airflow-common-env

AIRFLOW__CORE__EXECUTOR: LocalExecutor

AIRFLOW__DATABASE__SQL_ALCHEMY_CONN: postgresql+psycopg2://airflow:airflow@postgre

AIRFLOW__CORE__FERNET_KEY: ''

AIRFLOW__CORE__DAGS_ARE_PAUSED_AT_CREATION: 'true'

AIRFLOW__CORE__LOAD_EXAMPLES: 'false'

AIRFLOW__API__AUTH_BACKENDS: 'airflow.api.auth.backend.basic_auth'

AIRFLOW__WEBSERVER__EXPOSE_CONFIG: 'true'

_PIP_ADDITIONAL_REQUIREMENTS: >-

apache-airflow-providers-postgres

apache-airflow-providers-docker

pyspark==3.4.1

pandas

requests

great-expectations

volumes:

- ./airflow/dags:/opt/airflow/dags
- ./airflow/logs:/opt/airflow/logs
- ./airflow/plugins:/opt/airflow/plugins
- ./data:/opt/airflow/data
- ./scripts:/opt/airflow/scripts
- ./sql:/opt/airflow/sql

user: "\${AIRFLOW_UID:-50000}:0"

depends_on: &airflow-common-depends-on

postgres-airflow:

condition: service_healthy

services:

PostgreSQL for Airflow metadata

postgres-airflow:

image: postgres:13

environment:

POSTGRES_USER: airflow

POSTGRES_PASSWORD: airflow

POSTGRES_DB: airflow

volumes:

- postgres-airflow-db-volume:/var/lib/postgresql/data

```
healthcheck:
  test: ["CMD", "pg_isready", "-U", "airflow"]
  interval: 5s
  retries: 5
ports:
  - "5433:5432"
```

PostgreSQL for business data

```
postgres-data:
  image: postgres:13
  environment:
    POSTGRES_USER: datauser
    POSTGRES_PASSWORD: datapass
    POSTGRES_DB: analytics
  volumes:
    - postgres_data_db_volume:/var/lib/postgresql/data
    - ./sql/schema:/docker-entrypoint-initdb.d
  healthcheck:
    test: ["CMD", "pg_isready", "-U", "datauser"]
    interval: 5s
    retries: 5
  ports:
    - "5434:5432"
```

Redis for task queuing

```
redis:
  image: redis:7.0-alpine
  healthcheck:
    test: ["CMD", "redis-cli", "ping"]
    interval: 5s
    timeout: 3s
    retries: 5
  ports:
    - "6379:6379"
```

Airflow Webserver

```
airflow-webserver:
  <<: *airflow-common
  command: webserver
  ports:
    - "8080:8080"
  healthcheck:
    test: ["CMD", "curl", "--fail", "http://localhost:8080/health"]
    interval: 10s
```



```
    timeout: 10s
    retries: 5
depends_on:
  <<: *airflow-common-depends-on
```

Airflow Scheduler

```
airflow-scheduler:
  <<: *airflow-common
  command: scheduler
  healthcheck:
    test: ["CMD-SHELL", 'airflow jobs check --job-type SchedulerJob --hostname "${HOSTNAME}']
    interval: 10s
    timeout: 10s
    retries: 5
  depends_on:
    <<: *airflow-common-depends-on
```

Spark Master

```
spark-master:
  image: bitnami/spark:3.4.1
  environment:
    - SPARK_MODE=master
    - SPARK_RPC_AUTHENTICATION_ENABLED=no
    - SPARK_RPC_ENCRYPTION_ENABLED=no
    - SPARK_LOCAL_STORAGE_ENCRYPTION_ENABLED=no
    - SPARK_SSL_ENABLED=no
  ports:
    - "8081:8080"
    - "7077:7077"
  volumes:
    - ./spark/apps:/opt/spark-apps
    - ./data:/opt/spark-data
    - ./spark/conf:/opt/spark/conf
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost:8080"]
    interval: 30s
    timeout: 10s
    retries: 3
```

Spark Worker

```
spark-worker:
  image: bitnami/spark:3.4.1
  environment:
    - SPARK_MODE=worker
```

- SPARK_MASTER_URL=spark://spark-master:7077
- SPARK_WORKER_MEMORY=2G
- SPARK_WORKER_CORES=2
- SPARK_RPC_AUTHENTICATION_ENABLED=no
- SPARK_RPC_ENCRYPTION_ENABLED=no
- SPARK_LOCAL_STORAGE_ENCRYPTION_ENABLED=no
- SPARK_SSL_ENABLED=no

volumes:

- ./spark/apps:/opt/spark-apps
- ./data:/opt/spark-data

depends_on:

- spark-master

Jupyter Notebook for development

jupyter:

image: jupyter/pyspark-notebook:spark-3.4.1

ports:

- "8888:8888"

environment:

- JUPYTER_ENABLE_LAB=yes

volumes:

- ./notebooks:/home/jovyan/work
- ./data:/home/jovyan/data

command: start-notebook.sh --NotebookApp.token='' --NotebookApp.password=''

Airflow Init Service

airflow-init:

<<: *airflow-common

entrypoint: /bin/bash

command:

- -c
- |


```
function ver() {
    printf "%04d%04d%04d%04d" ${1//./ }
}
airflow_version=${$(AIRFLOW__LOGGING__LOGGING_LEVEL=INFO && airflow version)}
airflow_version_comparable=${(ver ${airflow_version})}
min_airflow_version=2.2.0
min_airflow_version_comparable=${(ver ${min_airflow_version})}
if (( airflow_version_comparable < min_airflow_version_comparable )); then
    echo "ERROR!!!: Too old Airflow version ${airflow_version}!"
    exit 1
fi
if [[ -z "${AIRFLOW_UID}" ]]; then
```

```

    echo "WARNING!!!: AIRFLOW_UID not set!"
    export AIRFLOW_UID=50000
fi
one_meg=1048576
mem_available=$((($(getconf _PHYS_PAGES) * $(getconf PAGE_SIZE) / one_meg))
cpus_available=$(grep -cE 'cpu[0-9]+' /proc/stat)
if (( mem_available < 4000 )) ; then
    echo "WARNING!!!: Not enough memory available for Docker."
fi
if (( cpus_available < 2 )); then
    echo "WARNING!!!: Not enough CPUS available for Docker."
fi
mkdir -p /sources/logs /sources/dags /sources/plugins
chown -R "${AIRFLOW_UID}:0" /sources/{logs,dags,plugins}
exec /entrypoint airflow version
environment:
    <<: *airflow-common-env
    _AIRFLOW_DB_UPGRADE: 'true'
    _AIRFLOW_WWW_USER_CREATE: 'true'
    _AIRFLOW_WWW_USER_USERNAME: ${_AIRFLOW_WWW_USER_USERNAME:-airflow}
    _AIRFLOW_WWW_USER_PASSWORD: ${_AIRFLOW_WWW_USER_PASSWORD:-airflow}

volumes:
    postgres_airflow_db_volume:
    postgres_data_db_volume:

networks:
    default:
        driver: bridge

```

Platform Initialization and Startup

Step 1: Environment Setup

```
bash
```

```
# Create project directory structure
```

```
mkdir -p week2-integration-project/{airflow/{dags,plugins,logs},spark/{apps,conf},data,
```

```
# Set environment variables
```

```
echo "AIRFLOW_UID=$(id -u)" > .env
```

```
echo "COMPOSE_PROJECT_NAME=week2-platform" >> .env
```

```
# Download datasets
```

```
cd data/raw
```

```
# Download all datasets from Kaggle (instructions below)
```

Step 2: Database Schema Setup

sql

```
-- sql/schema/create_tables.sql
-- Create analytics database schema

-- Customer dimension table
CREATE TABLE IF NOT EXISTS dim_customers (
    customer_id VARCHAR(50) PRIMARY KEY,
    customer_name VARCHAR(255),
    segment VARCHAR(100),
    country VARCHAR(100),
    city VARCHAR(100),
    registration_date DATE,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

```
-- Product dimension table
CREATE TABLE IF NOT EXISTS dim_products (
    product_id VARCHAR(50) PRIMARY KEY,
    product_name VARCHAR(255),
    category VARCHAR(100),
    sub_category VARCHAR(100),
    brand VARCHAR(100),
    price DECIMAL(10,2),
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

```
-- Time dimension table
CREATE TABLE IF NOT EXISTS dim_time (
    date_id DATE PRIMARY KEY,
    year INTEGER,
    quarter INTEGER,
    month INTEGER,
    week INTEGER,
    day_of_week INTEGER,
    is_weekend BOOLEAN,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

```
-- Sales fact table
CREATE TABLE IF NOT EXISTS fact_sales (
    sale_id SERIAL PRIMARY KEY,
    customer_id VARCHAR(50) REFERENCES dim_customers(customer_id),
    product_id VARCHAR(50) REFERENCES dim_products(product_id),
    sale_date DATE REFERENCES dim_time(date_id),
```

```
quantity INTEGER,  
unit_price DECIMAL(10,2),  
discount DECIMAL(5,4),  
total_amount DECIMAL(12,2),  
profit DECIMAL(12,2),  
created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
);
```

-- Customer analytics table

```
CREATE TABLE IF NOT EXISTS customer_analytics (  
    customer_id VARCHAR(50) PRIMARY KEY,  
    total_orders INTEGER,  
    total_spent DECIMAL(12,2),  
    average_order_value DECIMAL(10,2),  
    last_order_date DATE,  
    days_since_last_order INTEGER,  
    rfm_recency INTEGER,  
    rfm_frequency INTEGER,  
    rfm_monetary INTEGER,  
    rfm_score VARCHAR(10),  
    customer_segment VARCHAR(50),  
    lifetime_value DECIMAL(12,2),  
    churn_probability DECIMAL(5,4),  
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
);
```

-- Data quality monitoring table

```
CREATE TABLE IF NOT EXISTS data_quality_metrics (  
    id SERIAL PRIMARY KEY,  
    table_name VARCHAR(100),  
    metric_name VARCHAR(100),  
    metric_value DECIMAL(15,4),  
    threshold_value DECIMAL(15,4),  
    status VARCHAR(20),  
    execution_date DATE,  
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
);
```

-- Pipeline execution log

```
CREATE TABLE IF NOT EXISTS pipeline_execution_log (  
    id SERIAL PRIMARY KEY,  
    pipeline_name VARCHAR(100),  
    execution_date DATE,  
    start_time TIMESTAMP,
```

```

    end_time TIMESTAMP,
    status VARCHAR(20),
    records_processed INTEGER,
    error_message TEXT,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Create indexes for performance
CREATE INDEX IF NOT EXISTS idx_fact_sales_customer_date ON fact_sales(customer_id, sale_date);
CREATE INDEX IF NOT EXISTS idx_fact_sales_product_date ON fact_sales(product_id, sale_date);
CREATE INDEX IF NOT EXISTS idx_customer_analytics_segment ON customer_analytics(customer_id, segment);
CREATE INDEX IF NOT EXISTS idx_data_quality_table_date ON data_quality_metrics(table_name, date);

```

Step 3: Launch the Platform

```

bash

# Initialize Airflow
docker-compose up airflow-init

# Start all services
docker-compose up -d

# Verify all services are healthy
docker-compose ps

# Access platform components:
# Airflow Web UI: http://localhost:8080 (airflow/airflow)
# Spark Master UI: http://localhost:8081
# Jupyter Lab: http://localhost:8888
# PostgreSQL Data: localhost:5434 (datauser/datapass/analytics)
# PostgreSQL Airflow: localhost:5433 (airflow/airflow/airflow)

```

Data Acquisition and Setup

Multi-Source Dataset Download

Required Datasets for Integration:

1. E-Commerce Transactions

- Source: `kaggle.com/datasets/smeyanj/e-commerce-transactions-dataset`
- File: `ecommerce_transactions.csv`
- Purpose: Transaction fact data

2. Superstore Sales

- Source: kaggle.com/datasets/vivek468/superstore-dataset-final
- File: [superstore_sales.csv](#)
- Purpose: Historical sales analysis

3. Amazon Products

- Source: kaggle.com/datasets/jithinanievarghese/amazon-product-dataset
- File: [amazon_products.csv](#)
- Purpose: Product dimension data

4. Retail Analytics

- Source: kaggle.com/datasets/manjeetsingh/retaildataset
- File: [retail_analytics.csv](#)
- Purpose: Customer behavior analysis

Dataset Download Script:

python

```
# scripts/ingestion/download_datasets.py
```

```
"""
```

Dataset Download and Preparation Script

Automates the download and initial preparation of all required datasets

```
"""
```

```
import os
```

```
import pandas as pd
```

```
import requests
```

```
from pathlib import Path
```

```
def setup_data_directories():
```

```
    """Create required data directory structure"""
```

```
    directories = [
```

```
        'data/raw',
```

```
        'data/staging',
```

```
        'data/processed',
```

```
        'data/archive'
```

```
    ]
```

```
    for directory in directories:
```

```
        Path(directory).mkdir(parents=True, exist_ok=True)
```

```
        print(f"✅ Created directory: {directory}")
```

```
def validate_datasets():
```

```
    """Validate downloaded datasets"""
```

```
    required_files = [
```

```
        'data/raw/ecommerce_transactions.csv',
```

```
        'data/raw/superstore_sales.csv',
```

```
        'data/raw/amazon_products.csv',
```

```
        'data/raw/retail_analytics.csv'
```

```
    ]
```

```
    validation_results = {}
```

```
    for file_path in required_files:
```

```
        if os.path.exists(file_path):
```

```
            try:
```

```
                df = pd.read_csv(file_path)
```

```
                validation_results[file_path] = {
```

```
                    'exists': True,
```

```
                    'rows': len(df),
```

```
                    'columns': len(df.columns),
```

```

        'size_mb': round(os.path.getsize(file_path) / (1024*1024), 2)
    }
    print(f"✅ {file_path}: {len(df)} rows, {len(df.columns)} columns")
except Exception as e:
    validation_results[file_path] = {
        'exists': True,
        'error': str(e)
    }
    print(f"❌ {file_path}: Error reading file - {e}")
else:
    validation_results[file_path] = {'exists': False}
    print(f"❌ {file_path}: File not found")

return validation_results

def create_sample_data():
    """Create sample datasets if originals not available"""
    print("🔄 Creating sample datasets for development...")

    # Sample e-commerce transactions
    ecommerce_data = {
        'transaction_id': [f'TXN{i:06d}' for i in range(1, 1001)],
        'customer_id': [f'CUST{(i % 200) + 1:04d}' for i in range(1, 1001)],
        'product_id': [f'PROD{(i % 50) + 1:04d}' for i in range(1, 1001)],
        'transaction_date': pd.date_range('2023-01-01', periods=1000, freq='H'),
        'quantity': np.random.randint(1, 5, 1000),
        'unit_price': np.random.uniform(10, 200, 1000).round(2),
        'total_amount': lambda x: x['quantity'] * x['unit_price']
    }

    # Create and save sample datasets
    # (Additional sample data generation code here)

    print(f"✅ Sample datasets created successfully!")

if __name__ == "__main__":
    setup_data_directories()
    validation_results = validate_datasets()

    # If datasets missing, create samples
    missing_files = [k for k, v in validation_results.items() if not v.get('exists')]
    if missing_files:
        print(f"⚠️ Missing {len(missing_files)} required datasets")
        print("Please download from Kaggle or run with --create-samples")

```

```
else:  
    print("🎉 All datasets validated successfully!")
```

Integrated DAG Implementation

Complete Platform DAG

Primary DAG File: `airflow/dags/integrated_data_platform_dag.py`

python

.....

Integrated Data Platform DAG

=====

This DAG orchestrates the complete data platform workflow:

1. Multi-source data ingestion
2. Cross-dataset data quality validation
3. Distributed processing with Spark
4. Data warehouse loading
5. Analytics computation
6. Cross-platform monitoring

Schedule: Daily at 6 AM

Dependencies: Docker services (Spark, PostgreSQL, Redis)

.....

```
from datetime import datetime, timedelta
import pandas as pd
import logging
from pathlib import Path

from airflow import DAG
from airflow.operators.python_operator import PythonOperator
from airflow.operators.bash_operator import BashOperator
from airflow.operators.docker_operator import DockerOperator
from airflow.providers.postgres.operators.postgres import PostgresOperator
from airflow.providers.postgres.hooks.postgres import PostgresHook
from airflow.utils.task_group import TaskGroup
from airflow.models import Variable

# DAG Configuration
default_args = {
    'owner': 'data-platform-team',
    'depends_on_past': False,
    'start_date': datetime(2024, 1, 1),
    'email_on_failure': True,
    'email_on_retry': False,
    'retries': 2,
    'retry_delay': timedelta(minutes=5),
    'catchup': False
}

dag = DAG(
    'integrated_data_platform',
```

```

default_args=default_args,
description='Complete integrated data platform workflow',
schedule_interval='0 6 * * *',
max_active_runs=1,
tags=['platform', 'integration', 'multi-source', 'enterprise']
)

```

Configuration Management

```

def get_platform_config():
    """Centralized platform configuration"""
    return {
        'data_sources': {
            'ecommerce': '/opt/airflow/data/raw/ecommerce_transactions.csv',
            'superstore': '/opt/airflow/data/raw/superstore_sales.csv',
            'products': '/opt/airflow/data/raw/amazon_products.csv',
            'retail': '/opt/airflow/data/raw/retail_analytics.csv'
        },
        'staging_paths': {
            'customers': '/opt/airflow/data/staging/customers.parquet',
            'products': '/opt/airflow/data/staging/products.parquet',
            'transactions': '/opt/airflow/data/staging/transactions.parquet'
        },
        'quality_thresholds': {
            'completeness': 0.95,
            'uniqueness': 0.99,
            'validity': 0.90
        },
        'spark_config': {
            'master': 'spark://spark-master:7077',
            'app_name': 'IntegratedDataPlatform',
            'executor_memory': '2g',
            'driver_memory': '1g'
        }
    }
}

```

Platform Health Checks

```

def check_platform_health(**context):
    """Comprehensive platform health verification"""
    print("🔍 Starting platform health checks...")

    health_status = {
        'services': {},
        'data_sources': {},
        'storage': {},

```



```

        'overall': True
    }

# Check Docker services
services_to_check = [
    'week2-platform_postgres-data_1',
    'week2-platform_spark-master_1',
    'week2-platform_redis_1'
]

import subprocess
for service in services_to_check:
    try:
        result = subprocess.run(
            ['docker', 'inspect', service],
            capture_output=True,
            text=True
        )
        if result.returncode == 0:
            health_status['services'][service] = 'healthy'
            print(f"✅ Service {service}: healthy")
        else:
            health_status['services'][service] = 'unhealthy'
            health_status['overall'] = False
            print(f"❌ Service {service}: unhealthy")
    except Exception as e:
        health_status['services'][service] = f'error: {e}'
        health_status['overall'] = False
        print(f"❌ Service {service}: error - {e}")

# Check data source availability
config = get_platform_config()
for source_name, file_path in config['data_sources'].items():
    try:
        if Path(file_path).exists():
            df = pd.read_csv(file_path, nrows=5) # Quick validation
            health_status['data_sources'][source_name] = {
                'status': 'available',
                'rows': len(df),
                'columns': len(df.columns)
            }
            print(f"✅ Data source {source_name}: available ({len(df.columns)} col
        else:
            health_status['data_sources'][source_name] = 'missing'

```

```

        health_status['overall'] = False
        print(f"❌ Data source {source_name}: file not found")
    except Exception as e:
        health_status['data_sources'][source_name] = f'error: {e}'
        health_status['overall'] = False
        print(f"❌ Data source {source_name}: error - {e}")

# Check database connectivity
try:
    postgres_hook = PostgresHook(postgres_conn_id='postgres_data')
    postgres_hook.get_first("SELECT 1 as health_check")
    health_status['storage']['postgresql'] = 'connected'
    print("✅ PostgreSQL: connected")
except Exception as e:
    health_status['storage']['postgresql'] = f'error: {e}'
    health_status['overall'] = False
    print(f"❌ PostgreSQL: connection failed - {e}")

if health_status['overall']:
    print("🎉 Platform health check: ALL SYSTEMS GO!")
else:
    print("⚠️ Platform health check: ISSUES DETECTED")
    raise ValueError("Platform health check failed")

# Store health status for monitoring
context['task_instance'].xcom_push(key='health_status', value=health_status)
return health_status

# Multi-Source Data Ingestion
def ingest_multi_source_data(**context):
    """Ingest and standardize data from multiple sources"""
    print("🔄 Starting multi-source data ingestion...")

    config = get_platform_config()
    ingestion_results = {}

    # Ingest E-commerce transactions
    try:
        print("📊 Processing e-commerce transactions...")
        ecommerce_df = pd.read_csv(config['data_sources']['ecommerce'])

        # Standardize column names
        ecommerce_standardized = ecommerce_df.rename(columns={
            'Customer ID': 'customer_id',

```

```

        'Product ID': 'product_id',
        'Transaction Date': 'transaction_date',
        'Quantity': 'quantity',
        'Price': 'unit_price',
        'Total': 'total_amount'
    })

    # Data type conversion
    ecommerce_standardized['transaction_date'] = pd.to_datetime(
        ecommerce_standardized['transaction_date']
    )
    ecommerce_standardized['total_amount'] = pd.to_numeric(
        ecommerce_standardized['total_amount'], errors='coerce'
    )

    # Save to staging
    staging_path = '/opt/airflow/data/staging/ecommerce_standardized.parquet'
    ecommerce_standardized.to_parquet(staging_path, index=False)

    ingestion_results['ecommerce'] = {
        'status': 'success',
        'records': len(ecommerce_standardized),
        'staging_path': staging_path
    }
    print(f"✅ E-commerce: {len(ecommerce_standardized)} records ingested")

except Exception as e:
    ingestion_results['ecommerce'] = {'status': 'failed', 'error': str(e)}
    print(f"❌ E-commerce ingestion failed: {e}")

# Ingest Superstore sales
try:
    print("🚦 Processing superstore sales...")
    superstore_df = pd.read_csv(config['data_sources']['superstore'])

    # Standardize and enrich
    superstore_standardized = superstore_df.rename(columns={
        'Customer ID': 'customer_id',
        'Product ID': 'product_id',
        'Order Date': 'order_date',
        'Sales': 'sales_amount',
        'Profit': 'profit_amount',
        'Discount': 'discount_rate'
    })

```

```

# Calculate derived metrics
superstore_standardized['profit_margin'] = (
    superstore_standardized['profit_amount'] /
    superstore_standardized['sales_amount']
).fillna(0)

superstore_standardized['order_date'] = pd.to_datetime(
    superstore_standardized['order_date']
)

staging_path = '/opt/airflow/data/staging/superstore_standardized.parquet'
superstore_standardized.to_parquet(staging_path, index=False)

ingestion_results['superstore'] = {
    'status': 'success',
    'records': len(superstore_standardized),
    'staging_path': staging_path
}
print(f"✅ Superstore: {len(superstore_standardized)} records ingested")

except Exception as e:
    ingestion_results['superstore'] = {'status': 'failed', 'error': str(e)}
    print(f"❌ Superstore ingestion failed: {e}")

# Ingest Product catalog
try:
    print("📊 Processing product catalog...")
    products_df = pd.read_csv(config['data_sources']['products'])

    # Product dimension standardization
    products_standardized = products_df.rename(columns={
        'product_id': 'product_id',
        'product_name': 'product_name',
        'category': 'category',
        'price': 'list_price'
    })

    # Data enrichment
    products_standardized['price_tier'] = pd.cut(
        products_standardized['list_price'],
        bins=[0, 25, 100, 500, float('inf')],
        labels=['Budget', 'Mid-range', 'Premium', 'Luxury']
    )

```

```

staging_path = '/opt/airflow/data/staging/products_standardized.parquet'
products_standardized.to_parquet(staging_path, index=False)

ingestion_results['products'] = {
    'status': 'success',
    'records': len(products_standardized),
    'staging_path': staging_path
}
print(f"✅ Products: {len(products_standardized)} records ingested")

except Exception as e:
    ingestion_results['products'] = {'status': 'failed', 'error': str(e)}
    print(f"❌ Products ingestion failed: {e}")

# Summary and validation
successful_ingestions = sum(1 for result in ingestion_results.values()
                             if result.get('status') == 'success')
total_records = sum(result.get('records', 0) for result in ingestion_results.values()
                     if result.get('status') == 'success')

summary = {
    'successful_sources': successful_ingestions,
    'total_sources': len(config['data_sources']),
    'total_records_ingested': total_records,
    'ingestion_results': ingestion_results
}

print(f"📊 Ingestion Summary: {successful_ingestions}/{len(config['data_sources'])}")
print(f"📊 Total records ingested: {total_records:,}")

context['task_instance'].xcom_push(key='ingestion_results', value=summary)
return summary

# Cross-Dataset Data Quality Validation
def validate_integrated_data_quality(**context):
    """Comprehensive data quality validation across all datasets"""
    print("🔍 Starting integrated data quality validation...")

    config = get_platform_config()
    quality_results = {}

    def calculate_quality_metrics(df, dataset_name):
        """Calculate standard quality metrics for any dataset"""

```

```

metrics = {}

# Completeness (percentage of non-null values)
completeness = (df.notna().sum() / len(df)).mean()
metrics['completeness'] = round(completeness, 4)

# Uniqueness (for identifier columns)
id_columns = [col for col in df.columns if 'id' in col.lower()]
if id_columns:
    uniqueness_scores = []
    for col in id_columns:
        unique_ratio = df[col].nunique() / len(df)
        uniqueness_scores.append(unique_ratio)
    metrics['uniqueness'] = round(sum(uniqueness_scores) / len(uniqueness_scores), 4)
else:
    metrics['uniqueness'] = 1.0

# Validity (data type consistency)
numeric_columns = df.select_dtypes(include=['number']).columns
validity_scores = []
for col in numeric_columns:
    valid_ratio = (df[col].notna() & (df[col] >= 0)).sum() / len(df)
    validity_scores.append(valid_ratio)

if validity_scores:
    metrics['validity'] = round(sum(validity_scores) / len(validity_scores), 4)
else:
    metrics['validity'] = 1.0

# Overall quality score
metrics['overall_quality'] = round(
    (metrics['completeness'] + metrics['uniqueness'] + metrics['validity']) / 3, 4)

return metrics

# Validate each staged dataset
staging_files = [
    ('/opt/airflow/data/staging/ecommerce_standardized.parquet', 'ecommerce'),
    ('/opt/airflow/data/staging/superstore_standardized.parquet', 'superstore'),
    ('/opt/airflow/data/staging/products_standardized.parquet', 'products')
]

for file_path, dataset_name in staging_files:

```

```

try:
    if Path(file_path).exists():
        df = pd.read_parquet(file_path)
        metrics = calculate_quality_metrics(df, dataset_name)

        # Check against thresholds
        thresholds = config['quality_thresholds']
        quality_status = all([
            metrics['completeness'] >= thresholds['completeness'],
            metrics['uniqueness'] >= thresholds['uniqueness'],
            metrics['validity'] >= thresholds['validity']
        ])

        quality_results[dataset_name] = {
            'metrics': metrics,
            'passed_quality_check': quality_status,
            'record_count': len(df),
            'column_count': len(df.columns)
        }

        status_emoji = "✅" if quality_status else "⚠️"
        print(f"{status_emoji} {dataset_name}: Quality score {metrics['overall']}")
        print(f"    Completeness: {metrics['completeness']:.2%}")
        print(f"    Uniqueness: {metrics['uniqueness']:.2%}")
        print(f"    Validity: {metrics['validity']:.2%}")
    else:
        quality_results[dataset_name] = {
            'error': 'File not found',
            'passed_quality_check': False
        }
        print(f"❌ {dataset_name}: Staging file not found")

except Exception as e:
    quality_results[dataset_name] = {
        'error': str(e),
        'passed_quality_check': False
    }
    print(f"❌ {dataset_name}: Quality validation failed - {e}")

# Cross-dataset validation
print("\n🔄 Performing cross-dataset validation...")
cross_dataset_results = {}

try:

```

```

# Load datasets for cross-validation
ecommerce_df = pd.read_parquet('/opt/airflow/data/staging/ecommerce_standardized')
products_df = pd.read_parquet('/opt/airflow/data/staging/products_standardized')

# Referential integrity checks
ecommerce_products = set(ecommerce_df['product_id'].unique())
catalog_products = set(products_df['product_id'].unique())

# Product referential integrity
missing_products = ecommerce_products - catalog_products
referential_integrity_score = 1 - (len(missing_products) / len(ecommerce_products))

cross_dataset_results['referential_integrity'] = {
    'score': round(referential_integrity_score, 4),
    'missing_product_count': len(missing_products),
    'total_products_in_transactions': len(ecommerce_products)
}

print(f"🔗 Referential integrity: {referential_integrity_score:.2%}")
print(f"🔗 Missing products in catalog: {len(missing_products)}")

except Exception as e:
    cross_dataset_results['referential_integrity'] = {
        'error': str(e),
        'score': 0
    }
    print(f"❌ Cross-dataset validation failed: {e}")

# Overall quality assessment
passed_datasets = sum(1 for result in quality_results.values()
                      if result.get('passed_quality_check', False))
total_datasets = len(quality_results)

overall_assessment = {
    'datasets_passed': passed_datasets,
    'total_datasets': total_datasets,
    'overall_pass_rate': round(passed_datasets / total_datasets, 4) if total_datasets > 0 else 0,
    'quality_results': quality_results,
    'cross_dataset_results': cross_dataset_results
}

print(f"\n📊 Quality Assessment Summary:")
print(f"📊 Datasets passed: {passed_datasets}/{total_datasets}")
print(f"📊 Overall pass rate: {overall_assessment['overall_pass_rate']:.2%}")

```



```

# Fail the task if quality thresholds not met
if overall_assessment['overall_pass_rate'] < 0.8:
    raise ValueError(f"Data quality below acceptable threshold: {overall_assessment}")

context['task_instance'].xcom_push(key='quality_results', value=overall_assessment)
return overall_assessment

```

Spark-based Large Scale Processing

```

def process_data_with_spark(**context):
    """Process data using Spark for scalable analytics"""
    print("> Starting Spark-based data processing...")

    # This would normally be a Spark application
    # For this example, we'll simulate Spark processing with pandas
    # In production, this would use PySpark or submit to Spark cluster

    processing_results = {}

    try:
        # Load staging data
        print("📂 Loading data into Spark-like processing...")
        ecommerce_df = pd.read_parquet('/opt/airflow/data/staging/ecommerce_standardized')
        superstore_df = pd.read_parquet('/opt/airflow/data/staging/superstore_standardized')
        products_df = pd.read_parquet('/opt/airflow/data/staging/products_standardized')

        # Customer 360 analytics (simulating Spark transformations)
        print("🔄 Computing Customer 360 analytics...")

        # Combine transaction data
        all_transactions = pd.concat([
            ecommerce_df[['customer_id', 'transaction_date', 'total_amount']].rename(columns={'transaction_date': 'date'}),
            superstore_df[['customer_id', 'order_date', 'sales_amount']].rename(columns={'order_date': 'date'})
        ])

        # Customer analytics
        customer_analytics = all_transactions.groupby('customer_id').agg({
            'amount': ['sum', 'mean', 'count'],
            'date': ['min', 'max']
        }).round(2)

        # Flatten column names
        customer_analytics.columns = ['total_spent', 'avg_order_value', 'total_orders', 'min_date', 'max_date']
        customer_analytics = customer_analytics.reset_index()
    
```

```

# Calculate derived metrics
customer_analytics['days_since_last_order'] = (
    datetime.now() - pd.to_datetime(customer_analytics['last_order_date'])
).dt.days

# Customer segmentation (RFM-like)
customer_analytics['recency_score'] = pd.qcut(
    customer_analytics['days_since_last_order'],
    q=5, labels=[5, 4, 3, 2, 1]
)
customer_analytics['frequency_score'] = pd.qcut(
    customer_analytics['total_orders'].rank(method='first'),
    q=5, labels=[1, 2, 3, 4, 5]
)
customer_analytics['monetary_score'] = pd.qcut(
    customer_analytics['total_spent'],
    q=5, labels=[1, 2, 3, 4, 5]
)

# Segment labels
def segment_customers(row):
    score = int(str(row['recency_score']) + str(row['frequency_score']) + str(
    if score >= 544:
        return 'Champions'
    elif score >= 334:
        return 'Loyal Customers'
    elif score >= 244:
        return 'Potential Loyalists'
    elif score >= 144:
        return 'At Risk'
    else:
        return 'Lost Customers'

customer_analytics['customer_segment'] = customer_analytics.apply(segment_cust

# Save processed data
processed_path = '/opt/airflow/data/processed/customer_360_analytics.parquet'
customer_analytics.to_parquet(processed_path, index=False)

processing_results['customer_360'] = {
    'status': 'success',
    'records_processed': len(customer_analytics),
    'output_path': processed_path

```

```

}

print(f"✅ Customer 360: {len(customer_analytics)} customer profiles created")

# Product performance analytics
print("🔄 Computing product performance metrics...")

# Merge transaction data with product details
ecommerce_with_products = ecommerce_df.merge(
    products_df[['product_id', 'category', 'price_tier']],
    on='product_id',
    how='left'
)

# Product analytics
product_performance = ecommerce_with_products.groupby(['product_id', 'category',
    'quantity': 'sum',
    'total_amount': ['sum', 'mean'],
    'customer_id': 'nunique'
]).round(2)

product_performance.columns = ['total_quantity_sold', 'total_revenue', 'avg_order_value']
product_performance = product_performance.reset_index()

# Calculate performance metrics
product_performance['revenue_per_customer'] = (
    product_performance['total_revenue'] / product_performance['unique_customers']
).round(2)

processed_path = '/opt/airflow/data/processed/product_performance.parquet'
product_performance.to_parquet(processed_path, index=False)

processing_results['product_performance'] = {
    'status': 'success',
    'records_processed': len(product_performance),
    'output_path': processed_path
}

print(f"✅ Product Performance: {len(product_performance)} product analyses completed")

except Exception as e:
    processing_results['error'] = str(e)
    print(f"❌ Spark processing failed: {e}")
    raise

```

```
context['task_instance'].xcom_push(key='processing_results', value=processing_resu
return processing_results
```

```
# Data Warehouse Loading
```

```
def load_to_data_warehouse(**context):
```

```
    """Load processed data to PostgreSQL data warehouse"""
```

```
    print("\n Starting data warehouse loading...")
```

```
    postgres_hook = PostgresHook(postgres_conn_id='postgres_data')
```

```
    loading_results = {}
```

```
    try:
```

```
        # Load customer analytics
```

```
        print("\n Loading customer analytics to data warehouse...")
```

```
        customer_df = pd.read_parquet('/opt/airflow/data/processed/customer_360_analyt
```

```
        # Insert into database
```

```
        insert_sql = """
```

```
        INSERT INTO customer_analytics (
```

```
            customer_id, total_orders, total_spent, average_order_value,
```

```
            last_order_date, days_since_last_order, customer_segment
```

```
        ) VALUES %s
```

```
        ON CONFLICT (customer_id) DO UPDATE SET
```

```
            total_orders = EXCLUDED.total_orders,
```

```
            total_spent = EXCLUDED.total_spent,
```

```
            average_order_value = EXCLUDED.average_order_value,
```

```
            last_order_date = EXCLUDED.last_order_date,
```

```
            days_since_last_order = EXCLUDED.days_since_last_order,
```

```
            customer_segment = EXCLUDED.customer_segment,
```

```
            updated_at = CURRENT_TIMESTAMP
```

```
        """
```

```
        # Prepare data for insertion
```

```
        customer_values = [
```

```
            (
```

```
                row['customer_id'],
```

```
                int(row['total_orders']),
```

```
                float(row['total_spent']),
```

```
                float(row['avg_order_value']),
```

```
                row['last_order_date'],
```

```
                int(row['days_since_last_order']),
```

```
                row['customer_segment']
```

```
            )
```

```

        for _, row in customer_df.iterrows()
    ]

    postgres_hook.run(insert_sql, parameters=(customer_values,))

    loading_results['customer_analytics'] = {
        'status': 'success',
        'records_loaded': len(customer_df)
    }
    print(f"✅ Customer analytics: {len(customer_df)} records loaded")

    # Load data quality metrics
    print("📊 Recording data quality metrics...")
    quality_data = context['task_instance'].xcom_pull(task_ids='validate_integrate')

    quality_insert_sql = """
    INSERT INTO data_quality_metrics (
        table_name, metric_name, metric_value, threshold_value, status, execution_
    ) VALUES (%s, %s, %s, %s, %s, %s)
    """

    quality_records = []
    for dataset, metrics in quality_data['quality_results'].items():
        if 'metrics' in metrics:
            for metric_name, value in metrics['metrics'].items():
                threshold = 0.95 if metric_name != 'overall_quality' else 0.90
                status = 'PASS' if value >= threshold else 'FAIL'
                quality_records.append((
                    dataset, metric_name, value, threshold, status, context['ds']
                ))

    for record in quality_records:
        postgres_hook.run(quality_insert_sql, parameters=record)

    loading_results['data_quality_metrics'] = {
        'status': 'success',
        'records_loaded': len(quality_records)
    }
    print(f"✅ Data quality metrics: {len(quality_records)} records loaded")

    # Log pipeline execution
    execution_log_sql = """
    INSERT INTO pipeline_execution_log (
        pipeline_name, execution_date, start_time, end_time, status, records_proce

```

```

) VALUES (%s, %s, %s, %s, %s, %s)
"""

total_records = sum(result.get('records_loaded', 0) for result in loading_resu

postgres_hook.run(execution_log_sql, parameters=(
    'integrated_data_platform',
    context['ds'],
    context['data_interval_start'],
    datetime.now(),
    'SUCCESS',
    total_records
))

print(f"🟢 Data warehouse loading completed: {total_records:,} total records")

except Exception as e:
    loading_results['error'] = str(e)
    print(f"🔴 Data warehouse loading failed: {e}")

# Log failure
try:
    execution_log_sql = """
    INSERT INTO pipeline_execution_log (
        pipeline_name, execution_date, start_time, end_time, status, error_mes
    ) VALUES (%s, %s, %s, %s, %s, %s)
    """
    postgres_hook.run(execution_log_sql, parameters=(
        'integrated_data_platform',
        context['ds'],
        context['data_interval_start'],
        datetime.now(),
        'FAILED',
        str(e)
    ))
except:
    pass # Don't fail on logging failure

raise

context['task_instance'].xcom_push(key='loading_results', value=loading_results)
return loading_results

```

```

def generate_platform_insights(**context):
    """Generate comprehensive business insights from integrated data"""
    print("📊 Generating platform-wide business insights...")

    postgres_hook = PostgresHook(postgres_conn_id='postgres_data')
    insights = {}

    try:
        # Customer segment analysis
        segment_sql = """
        SELECT
            customer_segment,
            COUNT(*) as customer_count,
            AVG(total_spent) as avg_customer_value,
            SUM(total_spent) as segment_revenue,
            AVG(total_orders) as avg_orders_per_customer
        FROM customer_analytics
        GROUP BY customer_segment
        ORDER BY segment_revenue DESC
        """

        segment_data = postgres_hook.get_pandas_df(segment_sql)
        insights['customer_segments'] = segment_data.to_dict('records')

        # Top performing segments
        top_segment = segment_data.iloc[0]
        insights['top_performing_segment'] = {
            'segment': top_segment['customer_segment'],
            'revenue': float(top_segment['segment_revenue']),
            'customer_count': int(top_segment['customer_count'])
        }

        print(f"💰 Top performing segment: {top_segment['customer_segment']} "
              f"${top_segment['segment_revenue']:,.2f} revenue")

        # Data quality summary
        quality_sql = """
        SELECT
            table_name,
            metric_name,
            AVG(metric_value) as avg_score,
            COUNT(CASE WHEN status = 'PASS' THEN 1 END) as passed_checks,
            COUNT(*) as total_checks
        FROM data_quality_metrics
    """

```

```

WHERE execution_date = %s
GROUP BY table_name, metric_name
"""

```

```

quality_data = postgres_hook.get_pandas_df(quality_sql, parameters=[context['d
insights['data_quality_summary'] = quality_data.to_dict('records')

```

```

# Overall platform health

```

```

overall_quality = quality_data['avg_score'].mean()
insights['platform_health'] = {
    'overall_quality_score': round(overall_quality, 4),
    'total_quality_checks': int(quality_data['total_checks'].sum()),
    'passed_quality_checks': int(quality_data['passed_checks'].sum())
}

```

```

print(f"📊 Platform health score: {overall_quality:.2%}")

```

```

# Generate recommendations

```

```

recommendations = []

```

```

# Customer segment recommendations

```

```

champions_pct = 0
at_risk_pct = 0
for segment in insights['customer_segments']:
    if segment['customer_segment'] == 'Champions':
        champions_pct = (segment['customer_count'] / sum(s['customer_count'] f
    elif segment['customer_segment'] == 'At Risk':
        at_risk_pct = (segment['customer_count'] / sum(s['customer_count'] for

```

```

if champions_pct < 20:
    recommendations.append({
        'type': 'customer_retention',
        'priority': 'high',
        'message': f'Only {champions_pct:.1f}% of customers are Champions. Imp
    })

```

```

if at_risk_pct > 25:
    recommendations.append({
        'type': 'customer_winback',
        'priority': 'high',
        'message': f'{at_risk_pct:.1f}% of customers are at risk. Launch win-b
    })

```

```

# Data quality recommendations

```



```

if overall_quality < 0.95:
    recommendations.append({
        'type': 'data_quality',
        'priority': 'medium',
        'message': f'Data quality at {overall_quality:.1%}. Investigate data s
    })

```

```

insights['recommendations'] = recommendations

```

```

# Save insights report

```

```

import json
insights_path = '/opt/airflow/data/processed/platform_insights_report.json'
with open(insights_path, 'w') as f:
    json.dump(insights, f, indent=2, default=str)

```

```

print(f"📊 Platform insights report saved: {insights_path}")
print(f"📊 Generated {len(recommendations)} business recommendations")

```

```

except Exception as e:

```

```

    insights['error'] = str(e)

```

```

    print(f"❌ Insights generation failed: {e}")

```

```

    raise

```

```

context['task_instance'].xcom_push(key='platform_insights', value=insights)

```

```

return insights

```

```

# Platform monitoring and alerting

```

```

def monitor_platform_performance(**context):

```

```

    """Monitor platform performance and send alerts"""

```

```

    print("📊 Monitoring platform performance...")

```

```

    monitoring_results = {}

```

```

    try:

```

```

        # Collect performance metrics from XCom

```

```

        health_status = context['task_instance'].xcom_pull(task_ids='check_platform_he

```

```

        ingestion_results = context['task_instance'].xcom_pull(task_ids='ingest_multi_

```

```

        quality_results = context['task_instance'].xcom_pull(task_ids='validate_integr

```

```

        processing_results = context['task_instance'].xcom_pull(task_ids='process_data

```

```

        loading_results = context['task_instance'].xcom_pull(task_ids='load_to_data_wa

```

```

    # Calculate platform performance score

```

```

    metrics = {

```

```

        'health_score': 1.0 if health_status.get('overall') else 0.0,

```

```

        'ingestion_success_rate': ingestion_results.get('successful_sources', 0) /
        'quality_pass_rate': quality_results.get('overall_pass_rate', 0),
        'processing_success': 1.0 if 'error' not in processing_results else 0.0,
        'loading_success': 1.0 if 'error' not in loading_results else 0.0
    }

overall_platform_score = sum(metrics.values()) / len(metrics)

monitoring_results = {
    'execution_date': context['ds'],
    'overall_platform_score': round(overall_platform_score, 4),
    'individual_metrics': metrics,
    'total_records_processed': ingestion_results.get('total_records_ingested',
    'pipeline_status': 'SUCCESS' if overall_platform_score >= 0.8 else 'DEGRAD

}

# Alert conditions
alerts = []

if overall_platform_score < 0.8:
    alerts.append({
        'severity': 'HIGH',
        'message': f'Platform performance degraded: {overall_platform_score:.1%'
    })

if metrics['quality_pass_rate'] < 0.9:
    alerts.append({
        'severity': 'MEDIUM',
        'message': f'Data quality below threshold: {metrics["quality_pass_rate'
    })

if metrics['ingestion_success_rate'] < 1.0:
    alerts.append({
        'severity': 'MEDIUM',
        'message': f'Ingestion failures detected: {metrics["ingestion_success_
    })

monitoring_results['alerts'] = alerts

# Log monitoring results
print(f"📊 Platform Performance Score: {overall_platform_score:.1%}")
print(f"📊 Health: {metrics['health_score']:.1%}")
print(f"📊 Ingestion: {metrics['ingestion_success_rate']:.1%}")
print(f"📊 Quality: {metrics['quality_pass_rate']:.1%}")

```

```
print(f"📊 Processing: {metrics['processing_success']:.1%}")
print(f"📊 Loading: {metrics['loading_success']:.1%}")
```

```
if alerts:
    print(f"🚨 {len(alerts)} alerts generated")
    for alert in alerts:
        print(f"    {alert['severity']}: {alert['message']}")
else:
    print(f"✅ No alerts – all systems performing well")
```

```
except Exception as e:
    monitoring_results['error'] = str(e)
    print(f"❌ Platform monitoring failed: {e}")
```

```
context['task_instance'].xcom_push(key='monitoring_results', value=monitoring_resu
return monitoring_results
```

Task Group Definitions using Airflow TaskGroups

```
with TaskGroup('platform_initialization') as initialization_group:
    health_check_task = PythonOperator(
        task_id='check_platform_health',
        python_callable=check_platform_health,
        dag=dag
    )
```

```
with TaskGroup('data_ingestion') as ingestion_group:
    multi_source_ingestion_task = PythonOperator(
        task_id='ingest_multi_source_data',
        python_callable=ingest_multi_source_data,
        dag=dag
    )
```

```
with TaskGroup('data_quality') as quality_group:
    quality_validation_task = PythonOperator(
        task_id='validate_integrated_data_quality',
        python_callable=validate_integrated_data_quality,
        dag=dag
    )
```

```
with TaskGroup('data_processing') as processing_group:
    spark_processing_task = PythonOperator(
        task_id='process_data_with_spark',
        python_callable=process_data_with_spark,
        dag=dag
    )
```

```
)
```

```
with TaskGroup('data_warehouse') as warehouse_group:
```

```
    # Database preparation
```

```
    create_tables_task = PostgresOperator(
        task_id='ensure_database_schema',
        postgres_conn_id='postgres_data',
        sql='sql/schema/create_tables.sql',
        dag=dag
    )
```

```
    warehouse_loading_task = PythonOperator(
        task_id='load_to_data_warehouse',
        python_callable=load_to_data_warehouse,
        dag=dag
    )
```

```
    create_tables_task >> warehouse_loading_task
```

```
with TaskGroup('business_intelligence') as bi_group:
```

```
    insights_generation_task = PythonOperator(
        task_id='generate_platform_insights',
        python_callable=generate_platform_insights,
        dag=dag
    )
```

```
with TaskGroup('monitoring_and_alerting') as monitoring_group:
```

```
    monitoring_task = PythonOperator(
        task_id='monitor_platform_performance',
        python_callable=monitor_platform_performance,
        dag=dag
    )
```

```
    # Cleanup task
```

```
    cleanup_task = BashOperator(
        task_id='cleanup_temporary_files',
        bash_command="""
echo "🧹 Cleaning up temporary files..."
find /opt/airflow/data/staging -name "*.tmp" -delete 2>/dev/null || true
find /opt/airflow/data/staging -name "*.log" -mtime +7 -delete 2>/dev/null || true
echo "✅ Cleanup completed"
""",
        dag=dag
    )
```

```
monitoring_task >> cleanup_task
```

```
# Final success notification
```

```
platform_success_notification = BashOperator(  
    task_id='send_platform_success_notification',  
    bash_command="""  
    echo "🎉 Integrated Data Platform execution completed successfully!"  
    echo "📊 Check Airflow UI for detailed metrics and logs"  
    echo "🗄️ Data available in PostgreSQL analytics database"  
    echo "📈 Business insights report generated"  
    """,  
    trigger_rule='all_success',  
    dag=dag  
)  
  
# Define task group dependencies  
initialization_group >> ingestion_group >> quality_group >> processing_group >> warehouse_group
```

🔧 Spark Integration for Scalable Processing

⚡ Spark Application for Customer Analytics

File: `spark/apps/customer_analytics_spark.py`

python

.....

Customer Analytics Spark Application

=====

Production-ready Spark application for large-scale customer analytics
Handles millions of transactions and customer records efficiently

.....

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import *
from pyspark.sql.types import *
from pyspark.sql.window import Window
import sys
import logging

# Configure logging
logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(__name__)

class CustomerAnalyticsSpark:
    def __init__(self, app_name="CustomerAnalytics"):
        """Initialize Spark session with optimized configuration"""
        self.spark = SparkSession.builder \
            .appName(app_name) \
            .config("spark.sql.adaptive.enabled", "true") \
            .config("spark.sql.adaptive.coalescePartitions.enabled", "true") \
            .config("spark.sql.adaptive.skewJoin.enabled", "true") \
            .config("spark.serializer", "org.apache.spark.serializer.KryoSerializer") \
            .getOrCreate()

        self.spark.sparkContext.setLogLevel("WARN")
        logger.info("✅ Spark session initialized successfully")

    def load_data_sources(self, data_paths):
        """Load and validate all data sources"""
        logger.info("🇩🇪 Loading data sources...")

        data_frames = {}

        # Load e-commerce transactions
        if 'ecommerce' in data_paths:
            ecommerce_schema = StructType([
                StructField("transaction_id", StringType(), True),
                StructField("customer_id", StringType(), True),
```

```

        StructField("product_id", StringType(), True),
        StructField("transaction_date", TimestampType(), True),
        StructField("quantity", IntegerType(), True),
        StructField("unit_price", DoubleType(), True),
        StructField("total_amount", DoubleType(), True)
    ])

    data_frames['ecommerce'] = self.spark.read \
        .option("header", "true") \
        .schema(ecommerce_schema) \
        .csv(data_paths['ecommerce'])

    logger.info(f"✅ Loaded e-commerce data: {data_frames['ecommerce'].count()}")

# Load superstore data
    if 'superstore' in data_paths:
        superstore_schema = StructType([
            StructField("order_id", StringType(), True),
            StructField("customer_id", StringType(), True),
            StructField("product_id", StringType(), True),
            StructField("order_date", TimestampType(), True),
            StructField("sales_amount", DoubleType(), True),
            StructField("profit_amount", DoubleType(), True),
            StructField("discount_rate", DoubleType(), True)
        ])

        data_frames['superstore'] = self.spark.read \
            .option("header", "true") \
            .schema(superstore_schema) \
            .csv(data_paths['superstore'])

        logger.info(f"✅ Loaded superstore data: {data_frames['superstore'].count()}")

    return data_frames

def compute_customer_360(self, data_frames):
    """Compute comprehensive customer 360-degree view"""
    logger.info("🔄 Computing Customer 360 analytics...")

    # Combine all transaction sources
    all_transactions = None

    if 'ecommerce' in data_frames:
        ecommerce_unified = data_frames['ecommerce'].select(

```



```

        col("customer_id"),
        col("transaction_date").alias("date"),
        col("total_amount").alias("amount"),
        lit("ecommerce").alias("source")
    )
    all_transactions = ecommerce_unified

if 'superstore' in data_frames:
    superstore_unified = data_frames['superstore'].select(
        col("customer_id"),
        col("order_date").alias("date"),
        col("sales_amount").alias("amount"),
        lit("superstore").alias("source")
    )

    if all_transactions is not None:
        all_transactions = all_transactions.union(superstore_unified)
    else:
        all_transactions = superstore_unified

# Customer aggregations using Spark SQL
customer_metrics = all_transactions.groupBy("customer_id").agg(
    sum("amount").alias("total_spent"),
    avg("amount").alias("avg_order_value"),
    count("*").alias("total_orders"),
    min("date").alias("first_order_date"),
    max("date").alias("last_order_date"),
    countDistinct("source").alias("channel_diversity")
)

# Calculate recency, frequency, monetary (RFM) scores
current_date = current_timestamp()

customer_rfm = customer_metrics.withColumn(
    "days_since_last_order",
    datediff(current_date, col("last_order_date"))
).withColumn(
    "customer_lifetime_days",
    datediff(col("last_order_date"), col("first_order_date"))
)

# RFM scoring using ntile window functions
rfm_window = Window.orderBy("days_since_last_order")
frequency_window = Window.orderBy("total_orders")

```

```

monetary_window = Window.orderBy("total_spent")

customer_scored = customer_rfm \
    .withColumn("recency_score", 6 - ntile(5).over(rfm_window)) \
    .withColumn("frequency_score", ntile(5).over(frequency_window)) \
    .withColumn("monetary_score", ntile(5).over(monetary_window))

# Create RFM composite score
customer_final = customer_scored.withColumn(
    "rfm_score",
    concat(
        col("recency_score").cast("string"),
        col("frequency_score").cast("string"),
        col("monetary_score").cast("string")
    )
)

# Customer segmentation based on RFM scores
segmentation_conditions = [
    (col("rfm_score").rlike("^[4-5][4-5][4-5]$"), "Champions"),
    (col("rfm_score").rlike("^[3-5][3-5][3-5]$"), "Loyal Customers"),
    (col("rfm_score").rlike("^[3-5][1-2][3-5]$"), "Potential Loyalists"),
    (col("rfm_score").rlike("^[4-5][1-2][1-2]$"), "New Customers"),
    (col("rfm_score").rlike("^[1-2][3-5][3-5]$"), "At Risk"),
    (col("rfm_score").rlike("^[1-2][1-2][3-5]$"), "Cannot Lose Them"),
    (col("rfm_score").rlike("^[1-2][1-2][1-2]$"), "Lost Customers")
]

# Apply segmentation
customer_segmented = customer_final
for condition, segment in segmentation_conditions:
    customer_segmented = customer_segmented.withColumn(
        "customer_segment",
        when(condition, segment).otherwise(col("customer_segment"))
    )

# Fill any remaining null segments
customer_segmented = customer_segmented.fillna({"customer_segment": "Others"})

logger.info(f"✅ Customer 360 completed: {customer_segmented.count()} customer")
return customer_segmented

```

```

def compute_product_analytics(self, data_frames):
    """Compute product performance analytics"""

```

```

logger.info("🔄 Computing product analytics...")

# Combine product transaction data
all_product_transactions = None

if 'ecommerce' in data_frames:
    ecommerce_products = data_frames['ecommerce'].select(
        col("product_id"),
        col("quantity"),
        col("total_amount"),
        col("customer_id"),
        col("transaction_date").alias("date")
    )
    all_product_transactions = ecommerce_products

# Product performance metrics
product_performance = all_product_transactions.groupBy("product_id").agg(
    sum("quantity").alias("total_quantity_sold"),
    sum("total_amount").alias("total_revenue"),
    avg("total_amount").alias("avg_order_value"),
    countDistinct("customer_id").alias("unique_customers"),
    count("*").alias("total_transactions")
)

# Calculate additional metrics
product_metrics = product_performance.withColumn(
    "revenue_per_customer",
    round(col("total_revenue") / col("unique_customers"), 2)
).withColumn(
    "avg_quantity_per_transaction",
    round(col("total_quantity_sold") / col("total_transactions"), 2)
)

logger.info(f"✅ Product analytics completed: {product_metrics.count()} products")
return product_metrics

def save_results(self, dataframes, output_paths):
    """Save processed results to specified formats"""
    logger.info("💾 Saving processed results...")

    for name, df in dataframes.items():
        if name in output_paths:
            output_path = output_paths[name]

```

```

        # Save as Parquet for optimal performance
        df.coalesce(1).write \
            .mode("overwrite") \
            .option("compression", "snappy") \
            .parquet(output_path)

        logger.info(f"✅ Saved {name} to {output_path}")

def close(self):
    """Close Spark session"""
    self.spark.stop()
    logger.info("✅ Spark session closed")

def main():
    """Main execution function"""
    logger.info("🚀 Starting Customer Analytics Spark Job...")

    # Initialize Spark application
    analytics = CustomerAnalyticsSpark("IntegratedCustomerAnalytics")

    try:
        # Define data paths
        data_paths = {
            'ecommerce': '/opt/spark-data/staging/ecommerce_standardized.parquet',
            'superstore': '/opt/spark-data/staging/superstore_standardized.parquet'
        }

        # Load data
        data_frames = analytics.load_data_sources(data_paths)

        # Compute analytics
        customer_360 = analytics.compute_customer_360(data_frames)
        product_analytics = analytics.compute_product_analytics(data_frames)

        # Output paths
        output_paths = {
            'customer_360': '/opt/spark-data/processed/spark_customer_360',
            'product_analytics': '/opt/spark-data/processed/spark_product_analytics'
        }

        # Save results
        analytics.save_results({
            'customer_360': customer_360,
            'product_analytics': product_analytics
        })
    
```

```

}, output_paths)

# Print summary statistics
logger.info("📊 Job Summary:")
logger.info(f"    Customer profiles processed: {customer_360.count():,}")
logger.info(f"    Products analyzed: {product_analytics.count():,}")

# Segment distribution
segment_counts = customer_360.groupBy("customer_segment").count().collect()
logger.info("📊 Customer Segment Distribution:")
for row in segment_counts:
    logger.info(f"    {row['customer_segment']}: {row['count']:,} customers")

logger.info("🎉 Customer Analytics Spark Job completed successfully!")

except Exception as e:
    logger.error(f"❌ Spark job failed: {str(e)}")
    raise

finally:
    analytics.close()

if __name__ == "__main__":
    main()

```

📊 Monitoring and Observability Implementation

🎯 Data Quality Monitoring Framework

File: `scripts/data_quality/quality_checks.py`

python

```
.....
```

Comprehensive Data Quality Monitoring Framework

```
=====
```

Implements Great Expectations-like validation with custom business rules

```
.....
```

```
import pandas as pd
import numpy as np
import json
from datetime import datetime, timedelta
from pathlib import Path
import logging

logger = logging.getLogger(__name__)

class DataQualityFramework:
    def __init__(self, config_path=None):
        """Initialize data quality framework with configurable rules"""
        self.config = self._load_config(config_path)
        self.results = {}

    def _load_config(self, config_path):
        """Load data quality configuration"""
        default_config = {
            "completeness_threshold": 0.95,
            "uniqueness_threshold": 0.99,
            "validity_threshold": 0.90,
            "freshness_threshold_hours": 24,
            "anomaly_detection_window": 30
        }

        if config_path and Path(config_path).exists():
            with open(config_path) as f:
                custom_config = json.load(f)
                default_config.update(custom_config)

        return default_config

    def validate_completeness(self, df, dataset_name, required_columns=None):
        """Validate data completeness across columns"""
        logger.info(f"🔍 Validating completeness for {dataset_name}...")

        if required_columns is None:
```

```

        required_columns = df.columns.tolist()

    completeness_results = {}

    for column in required_columns:
        if column in df.columns:
            non_null_ratio = df[column].notna().sum() / len(df)
            completeness_results[column] = {
                'completeness_ratio': round(non_null_ratio, 4),
                'null_count': int(df[column].isna().sum()),
                'total_count': len(df),
                'passed': non_null_ratio >= self.config['completeness_threshold']
            }
        else:
            completeness_results[column] = {
                'error': 'Column not found',
                'passed': False
            }

    # Overall completeness score
    overall_completeness = np.mean([
        result['completeness_ratio'] for result in completeness_results.values()
        if 'completeness_ratio' in result
    ])

    self.results[f'{dataset_name}_completeness'] = {
        'overall_score': round(overall_completeness, 4),
        'threshold': self.config['completeness_threshold'],
        'passed': overall_completeness >= self.config['completeness_threshold'],
        'column_details': completeness_results
    }

    status = "✅" if overall_completeness >= self.config['completeness_threshold']
    logger.info(f"{status} Completeness for {dataset_name}: {overall_completeness}")

    return self.results[f'{dataset_name}_completeness']

def validate_uniqueness(self, df, dataset_name, unique_columns=None):
    """Validate uniqueness constraints"""
    logger.info(f"🔍 Validating uniqueness for {dataset_name}...")

    if unique_columns is None:
        # Auto-detect ID columns
        unique_columns = [col for col in df.columns if 'id' in col.lower()]

```



```

uniqueness_results = {}

for column in unique_columns:
    if column in df.columns:
        unique_ratio = df[column].nunique() / len(df)
        duplicate_count = len(df) - df[column].nunique()

        uniqueness_results[column] = {
            'uniqueness_ratio': round(unique_ratio, 4),
            'unique_count': int(df[column].nunique()),
            'duplicate_count': duplicate_count,
            'total_count': len(df),
            'passed': unique_ratio >= self.config['uniqueness_threshold']
        }
    else:
        uniqueness_results[column] = {
            'error': 'Column not found',
            'passed': False
        }

# Overall uniqueness score
if uniqueness_results:
    overall_uniqueness = np.mean([
        result['uniqueness_ratio'] for result in uniqueness_results.values()
        if 'uniqueness_ratio' in result
    ])
else:
    overall_uniqueness = 1.0 # No unique columns to check

self.results[f'{dataset_name}_uniqueness'] = {
    'overall_score': round(overall_uniqueness, 4),
    'threshold': self.config['uniqueness_threshold'],
    'passed': overall_uniqueness >= self.config['uniqueness_threshold'],
    'column_details': uniqueness_results
}

status = "✅" if overall_uniqueness >= self.config['uniqueness_threshold'] else
logger.info(f'{status} Uniqueness for {dataset_name}: {overall_uniqueness:.2%}')

return self.results[f'{dataset_name}_uniqueness']

def validate_business_rules(self, df, dataset_name, rules=None):
    """Validate custom business rules"""

```

```

logger.info(f"🔍 Validating business rules for {dataset_name}...")

if rules is None:
    rules = self._get_default_business_rules(dataset_name)

rule_results = {}

for rule_name, rule_config in rules.items():
    try:
        rule_function = rule_config['function']
        rule_threshold = rule_config.get('threshold', 0.95)

        # Apply business rule
        if rule_function == 'positive_amounts':
            amount_columns = [col for col in df.columns if 'amount' in col.lower]
            valid_ratios = []
            for col in amount_columns:
                if col in df.columns and df[col].dtype in ['float64', 'int64']:
                    valid_ratio = (df[col] >= 0).sum() / len(df)
                    valid_ratios.append(valid_ratio)

            overall_validity = np.mean(valid_ratios) if valid_ratios else 1.0

        elif rule_function == 'valid_dates':
            date_columns = [col for col in df.columns if 'date' in col.lower()]
            valid_ratios = []
            for col in date_columns:
                if col in df.columns:
                    try:
                        pd.to_datetime(df[col], errors='coerce')
                        valid_dates = pd.to_datetime(df[col], errors='coerce')
                        valid_ratio = valid_dates.sum() / len(df)
                        valid_ratios.append(valid_ratio)
                    except:
                        valid_ratios.append(0.0)

            overall_validity = np.mean(valid_ratios) if valid_ratios else 1.0

        elif rule_function == 'reasonable_quantities':
            quantity_columns = [col for col in df.columns if 'quantity' in col.lower]
            valid_ratios = []
            for col in quantity_columns:
                if col in df.columns and df[col].dtype in ['float64', 'int64']:
                    # Reasonable quantities: between 1 and 1000

```

```

        valid_ratio = ((df[col] >= 1) & (df[col] <= 1000)).sum() /
        valid_ratios.append(valid_ratio)

    overall_validity = np.mean(valid_ratios) if valid_ratios else 1.0

    else:
        overall_validity = 1.0 # Unknown rule, pass by default

    rule_results[rule_name] = {
        'validity_score': round(overall_validity, 4),
        'threshold': rule_threshold,
        'passed': overall_validity >= rule_threshold,
        'rule_function': rule_function
    }

    except Exception as e:
        rule_results[rule_name] = {
            'error': str(e),
            'passed': False
        }

# Overall business rule score
if rule_results:
    overall_business_validity = np.mean([
        result['validity_score'] for result in rule_results.values()
        if 'validity_score' in result
    ])
else:
    overall_business_validity = 1.0

self.results[f'{dataset_name}_business_rules'] = {
    'overall_score': round(overall_business_validity, 4),
    'threshold': self.config['validity_threshold'],
    'passed': overall_business_validity >= self.config['validity_threshold'],
    'rule_details': rule_results
}

status = "✅" if overall_business_validity >= self.config['validity_threshold']
logger.info(f'{status} Business rules for {dataset_name}: {overall_business_va

return self.results[f'{dataset_name}_business_rules']

def _get_default_business_rules(self, dataset_name):
    """Get default business rules based on dataset type"""

```

```

rules = {
    'positive_amounts': {
        'function': 'positive_amounts',
        'threshold': 0.95,
        'description': 'All amount/price fields should be positive'
    },
    'valid_dates': {
        'function': 'valid_dates',
        'threshold': 0.99,
        'description': 'All date fields should be valid dates'
    },
    'reasonable_quantities': {
        'function': 'reasonable_quantities',
        'threshold': 0.95,
        'description': 'Quantities should be reasonable (1-1000)'
    }
}

```

```

return rules

```

```

def detect_anomalies(self, df, dataset_name, metric_columns=None):
    """Detect statistical anomalies in numeric data"""
    logger.info(f"🔍 Detecting anomalies for {dataset_name}...")

    if metric_columns is None:
        metric_columns = df.select_dtypes(include=[np.number]).columns.tolist()

    anomaly_results = {}

    for column in metric_columns:
        if column in df.columns and df[column].dtype in ['float64', 'int64']:
            try:
                # Z-score based anomaly detection
                z_scores = np.abs((df[column] - df[column].mean()) / df[column].std())
                anomaly_threshold = 3 # 3 standard deviations

                anomalies = z_scores > anomaly_threshold
                anomaly_count = anomalies.sum()
                anomaly_ratio = anomaly_count / len(df)

                anomaly_results[column] = {
                    'anomaly_count': int(anomaly_count),
                    'anomaly_ratio': round(anomaly_ratio, 4),
                    'total_count': len(df),

```

```

        'mean_value': round(df[column].mean(), 2),
        'std_value': round(df[column].std(), 2),
        'passed': anomaly_ratio <= 0.05 # Less than 5% anomalies
    }

    except Exception as e:
        anomaly_results[column] = {
            'error': str(e),
            'passed': False
        }

# Overall anomaly assessment
if anomaly_results:
    overall_anomaly_score = 1 - np.mean([
        result['anomaly_ratio'] for result in anomaly_results.values()
        if 'anomaly_ratio' in result
    ])
else:
    overall_anomaly_score = 1.0

self.results[f'{dataset_name}_anomalies'] = {
    'overall_score': round(overall_anomaly_score, 4),
    'threshold': 0.95,
    'passed': overall_anomaly_score >= 0.95,
    'column_details': anomaly_results
}

status = "✅" if overall_anomaly_score >= 0.95 else "❌"
logger.info(f"{status} Anomaly detection for {dataset_name}: {overall_anomaly_score}")

return self.results[f'{dataset_name}_anomalies']

def generate_quality_report(self, output_path):
    """Generate comprehensive data quality report"""
    logger.info("📊 Generating data quality report...")

    report = {
        'execution_timestamp': datetime.now().isoformat(),
        'config': self.config,
        'results': self.results,
        'summary': self._calculate_summary_metrics()
    }

# Save report

```

```

with open(output_path, 'w') as f:
    json.dump(report, f, indent=2, default=str)

logger.info(f"✅ Quality report saved to {output_path}")
return report

```

```

def _calculate_summary_metrics(self):
    """Calculate overall summary metrics"""
    all_scores = []
    passed_checks = 0
    total_checks = 0

    for check_name, result in self.results.items():
        if 'overall_score' in result:
            all_scores.append(result['overall_score'])
        if 'passed' in result:
            total_checks += 1
            if result['passed']:
                passed_checks += 1

    return {
        'overall_quality_score': round(np.mean(all_scores), 4) if all_scores else 0,
        'checks_passed': passed_checks,
        'total_checks': total_checks,
        'pass_rate': round(passed_checks / total_checks, 4) if total_checks > 0 else 0
    }

```

```

def run_comprehensive_quality_checks(data_paths, output_path):
    """Run comprehensive quality checks on all datasets"""
    logger.info(f"🚀 Starting comprehensive data quality validation...")

    quality_framework = DataQualityFramework()

    for dataset_name, file_path in data_paths.items():
        if Path(file_path).exists():
            logger.info(f"📁 Processing {dataset_name}...")

            try:
                # Load dataset
                if file_path.endswith('.parquet'):
                    df = pd.read_parquet(file_path)
                else:
                    df = pd.read_csv(file_path)

```

```

        # Run all quality checks
        quality_framework.validate_completeness(df, dataset_name)
        quality_framework.validate_uniqueness(df, dataset_name)
        quality_framework.validate_business_rules(df, dataset_name)
        quality_framework.detect_anomalies(df, dataset_name)

        logger.info(f"✅ Quality checks completed for {dataset_name}")

    except Exception as e:
        logger.error(f"❌ Quality checks failed for {dataset_name}: {e}")
    else:
        logger.warning(f"⚠️ Dataset not found: {file_path}")

# Generate comprehensive report
report = quality_framework.generate_quality_report(output_path)

logger.info(f"🎉 Comprehensive quality validation completed!")
return report

if __name__ == "__main__":
    # Example usage
    data_paths = {
        'ecommerce': '/opt/airflow/data/staging/ecommerce_standardized.parquet',
        'superstore': '/opt/airflow/data/staging/superstore_standardized.parquet',
        'products': '/opt/airflow/data/staging/products_standardized.parquet'
    }

    output_path = '/opt/airflow/data/processed/quality_report.json'
    run_comprehensive_quality_checks(data_paths, output_path)

```

Platform Configuration and Deployment

Environment-Specific Configuration

File: `config/platform_config.yaml`

yaml

Platform Configuration for Different Environments

=====

environments:

development:

database:

host: "localhost"

port: 5434

database: "analytics_dev"

username: "datauser"

password: "datapass"

spark:

master: "local[2]"

executor_memory: "1g"

driver_memory: "512m"

airflow:

parallelism: 2

max_active_runs: 1

catchup: false

data_retention:

staging_days: 7

processed_days: 30

logs_days: 14

staging:

database:

host: "postgres-staging"

port: 5432

database: "analytics_staging"

username: "datauser"

password: "\${POSTGRES_PASSWORD}"

spark:

master: "spark://spark-master:7077"

executor_memory: "2g"

driver_memory: "1g"

airflow:

parallelism: 4

max_active_runs: 2

catchup: false

data_retention:

staging_days: 14

processed_days: 60

logs_days: 30

production:

database:

host: "postgres-prod"

port: 5432

database: "analytics_prod"

username: "datauser"

password: "\${POSTGRES_PASSWORD}"

spark:

master: "spark://spark-master:7077"

executor_memory: "4g"

driver_memory: "2g"

airflow:

parallelism: 8

max_active_runs: 1

catchup: false

data_retention:

staging_days: 30

processed_days: 365

logs_days: 90

data_quality:

thresholds:

completeness: 0.95

uniqueness: 0.99

validity: 0.90

freshness_hours: 24

anomaly_threshold: 0.05

alerts:

email_recipients:

- "data-team@company.com"

- "devops@company.com"

slack_webhook: "\${SLACK_WEBHOOK_URL}"

monitoring:

```
health_check_interval: 300 # 5 minutes
metric_collection_interval: 60 # 1 minute
alert_cooldown: 1800 # 30 minutes
```

security:

```
encrypt_at_rest: true
encrypt_in_transit: true
access_logging: true
data_masking:
  pii_columns:
    - "email"
    - "phone"
    - "ssn"
  masking_strategy: "hash"
```

Deployment Scripts

File: `scripts/deployment/deploy_platform.sh`

bash

```
#!/bin/bash
# Platform Deployment Script
# =====

set -e # Exit on any error

echo "🚀 Starting Integrated Data Platform Deployment..."

# Configuration
ENVIRONMENT=${1:-development}
PROJECT_NAME="week2-platform"
COMPOSE_FILE="docker-compose.yml"

echo "🏗️ Deployment Environment: $ENVIRONMENT"

# Validate environment
if [[ ! "$ENVIRONMENT" =~ ^(development|staging|production)$ ]]; then
    echo "❌ Invalid environment. Use: development, staging, or production"
    exit 1
fi

# Pre-deployment checks
echo "🔍 Running pre-deployment checks..."

# Check Docker
if ! command -v docker &> /dev/null; then
    echo "❌ Docker is not installed"
    exit 1
fi

if ! command -v docker-compose &> /dev/null; then
    echo "❌ Docker Compose is not installed"
    exit 1
fi

echo "✅ Docker and Docker Compose are available"

# Check system resources
AVAILABLE_MEMORY=$(free -m | awk 'NR==2{printf "%.0f", $7/1024 }')
if [ "$AVAILABLE_MEMORY" -lt 4 ]; then
    echo "⚠️ Warning: Available memory is ${AVAILABLE_MEMORY}GB. Recommend at least 4
fi
```

```

# Set environment variables
export AIRFLOW_UID=$(id -u)
export COMPOSE_PROJECT_NAME=$PROJECT_NAME
export ENVIRONMENT=$ENVIRONMENT

echo "📁 Creating required directories..."
mkdir -p data/{raw,staging,processed,archive}
mkdir -p airflow/{dags,logs,plugins}
mkdir -p logs/{airflow,spark,postgres}
mkdir -p monitoring/{grafana,prometheus}

# Set proper permissions
chmod -R 755 data
chmod -R 755 airflow
chmod -R 755 logs

echo "✅ Directory structure created"

# Download required datasets (if not present)
echo "📄 Checking dataset availability..."
REQUIRED_FILES=(
    "data/raw/ecommerce_transactions.csv"
    "data/raw/superstore_sales.csv"
    "data/raw/amazon_products.csv"
    "data/raw/retail_analytics.csv"
)

MISSING_FILES=0
for file in "${REQUIRED_FILES[@]}; do
    if [ ! -f "$file" ]; then
        echo "⚠️ Missing: $file"
        MISSING_FILES=$((MISSING_FILES + 1))
    fi
done

if [ $MISSING_FILES -gt 0 ]; then
    echo "📄 $MISSING_FILES datasets missing. Creating sample data..."
    python3 scripts/ingestion/download_datasets.py --create-samples
    echo "✅ Sample datasets created"
else
    echo "✅ All datasets available"
fi

# Environment-specific configuration

```

```

case $ENVIRONMENT in
    development)
        echo "🔧 Configuring for development environment..."
        export AIRFLOW_PARALLELISM=2
        export SPARK_WORKER_MEMORY=1G
        export POSTGRES_MAX_CONNECTIONS=50
        ;;
    staging)
        echo "🔧 Configuring for staging environment..."
        export AIRFLOW_PARALLELISM=4
        export SPARK_WORKER_MEMORY=2G
        export POSTGRES_MAX_CONNECTIONS=100
        ;;
    production)
        echo "🔧 Configuring for production environment..."
        export AIRFLOW_PARALLELISM=8
        export SPARK_WORKER_MEMORY=4G
        export POSTGRES_MAX_CONNECTIONS=200
        ;;
esac

# Initialize Airflow
echo "🐳 Initializing Airflow..."
docker-compose up airflow-init

# Start platform services
echo "🚀 Starting platform services..."
docker-compose up -d

# Wait for services to be healthy
echo "⌚ Waiting for services to be healthy..."
MAX_WAIT=300 # 5 minutes
WAIT_TIME=0

while [ $WAIT_TIME -lt $MAX_WAIT ]; do
    if docker-compose ps | grep -q "unhealthy\|starting"; then
        echo "⌚ Services still starting... ($WAIT_TIME/$MAX_WAIT seconds)"
        sleep 10
        WAIT_TIME=$((WAIT_TIME + 10))
    else
        break
    fi
done

```

```

# Health check
echo "🔍 Running platform health check..."
HEALTH_SCRIPT="scripts/monitoring/health_checks.py"
if [ -f "$HEALTH_SCRIPT" ]; then
    python3 "$HEALTH_SCRIPT"
    if [ $? -eq 0 ]; then
        echo "✅ Platform health check passed"
    else
        echo "❌ Platform health check failed"
        echo "📋 Service status:"
        docker-compose ps
        exit 1
    fi
else
    echo "⚠️ Health check script not found, skipping..."
fi

# Display access information
echo ""
echo "🎉 Platform deployment completed successfully!"
echo ""
echo "📊 Access Information:"
echo "  Airflow Web UI:    http://localhost:8080 (airflow/airflow)"
echo "  Spark Master UI:   http://localhost:8081"
echo "  Jupyter Lab:       http://localhost:8888"
echo "  PostgreSQL Data:   localhost:5434 (datauser/datapass/analytics)"
echo ""
echo "🔧 Management Commands:"
echo "  View logs:         docker-compose logs -f [service]"
echo "  Stop platform:     docker-compose down"
echo "  Restart service:   docker-compose restart [service]"
echo "  Scale workers:     docker-compose up -d --scale spark-worker=3"
echo ""
echo "📖 Next Steps:"
echo "  1. Access Airflow UI and unpause the 'integrated_data_platform' DAG"
echo "  2. Trigger a manual run to test the complete pipeline"
echo "  3. Monitor execution in Airflow UI and check logs"
echo "  4. Query results in PostgreSQL analytics database"
echo ""

```

🔍 Health Monitoring Script

File: `scripts/monitoring/health_checks.py`

python

```
.....
```

Platform Health Monitoring

```
=====
```

Comprehensive health checks for all platform components

```
.....
```

```
import subprocess
import psycpg2
import requests
import time
import json
from datetime import datetime
import logging

logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(__name__)

class PlatformHealthMonitor:
    def __init__(self):
        self.health_results = {}
        self.overall_health = True

    def check_docker_services(self):
        """Check if all Docker services are running"""
        logger.info("🐳 Checking Docker services...")

        required_services = [
            'postgres-airflow',
            'postgres-data',
            'redis',
            'airflow-webserver',
            'airflow-scheduler',
            'spark-master',
            'spark-worker',
            'jupyter'
        ]

        try:
            result = subprocess.run(
                ['docker-compose', 'ps', '--format', 'json'],
                capture_output=True,
                text=True,
                check=True
```

```

)

services_status = {}
for line in result.stdout.strip().split('\n'):
    if line.strip():
        service_info = json.loads(line)
        service_name = service_info['Service']
        service_state = service_info['State']

        services_status[service_name] = {
            'status': service_state,
            'healthy': service_state == 'running'
        }

# Check required services
missing_services = []
unhealthy_services = []

for service in required_services:
    if service not in services_status:
        missing_services.append(service)
    elif not services_status[service]['healthy']:
        unhealthy_services.append(service)

docker_health = {
    'all_services_running': len(missing_services) == 0 and len(unhealthy_s
    'running_services': len([s for s in services_status.values() if s['hea
    'total_expected': len(required_services),
    'missing_services': missing_services,
    'unhealthy_services': unhealthy_services,
    'service_details': services_status
}

self.health_results['docker_services'] = docker_health

if not docker_health['all_services_running']:
    self.overall_health = False
    logger.error(f"❌ Docker services issues: {len(missing_services)} miss
else:
    logger.info("✅ All Docker services running")

except Exception as e:
    self.health_results['docker_services'] = {'error': str(e), 'healthy': False
    self.overall_health = False

```

```

        logger.error(f"❌ Docker services check failed: {e}")

def check_airflow_health(self):
    """Check Airflow web server and scheduler health"""
    logger.info("🔍 Checking Airflow health...")

    try:
        # Check web server
        response = requests.get('http://localhost:8080/health', timeout=10)
        webserver_healthy = response.status_code == 200

        # Check if we can access the main page
        main_response = requests.get('http://localhost:8080/', timeout=10)
        webserver_accessible = main_response.status_code == 200

        airflow_health = {
            'webserver_healthy': webserver_healthy,
            'webserver_accessible': webserver_accessible,
            'webserver_response_time': response.elapsed.total_seconds()
        }

        self.health_results['airflow'] = airflow_health

        if webserver_healthy and webserver_accessible:
            logger.info("✅ Airflow web server healthy")
        else:
            self.overall_health = False
            logger.error("❌ Airflow web server issues detected")

    except Exception as e:
        self.health_results['airflow'] = {'error': str(e), 'healthy': False}
        self.overall_health = False
        logger.error(f"❌ Airflow health check failed: {e}")

def check_spark_health(self):
    """Check Spark cluster health"""
    logger.info("🔍 Checking Spark cluster health...")

    try:
        # Check Spark master UI
        response = requests.get('http://localhost:8081', timeout=10)
        master_healthy = response.status_code == 200

        # Try to get cluster info

```

```

json_response = requests.get('http://localhost:8081/json/', timeout=10)
cluster_info = {}

if json_response.status_code == 200:
    cluster_data = json_response.json()
    cluster_info = {
        'workers_count': len(cluster_data.get('workers', [])),
        'alive_workers': len([w for w in cluster_data.get('workers', []) if w['alive']]),
        'cores_total': cluster_data.get('cores', 0),
        'memory_total': cluster_data.get('memory', 0)
    }

spark_health = {
    'master_healthy': master_healthy,
    'cluster_info': cluster_info,
    'all_workers_alive': cluster_info.get('workers_count', 0) == cluster_info.get('alive_workers', 0)
}

self.health_results['spark'] = spark_health

if master_healthy and spark_health['all_workers_alive']:
    logger.info(f"✅ Spark cluster healthy: {cluster_info.get('alive_workers', 0)} workers alive")
else:
    self.overall_health = False
    logger.error(f"❌ Spark cluster issues detected")

except Exception as e:
    self.health_results['spark'] = {'error': str(e), 'healthy': False}
    self.overall_health = False
    logger.error(f"❌ Spark health check failed: {e}")

def check_database_connectivity(self):
    """Check PostgreSQL database connectivity"""
    logger.info(f"🐘 Checking database connectivity...")

    databases = {
        'airflow_metadata': {
            'host': 'localhost',
            'port': 5433,
            'database': 'airflow',
            'username': 'airflow',
            'password': 'airflow'
        },
        'analytics_data': {

```

```

        'host': 'localhost',
        'port': 5434,
        'database': 'analytics',
        'username': 'datauser',
        'password': 'datapass'
    }
}

database_results = {}

for db_name, config in databases.items():
    try:
        conn = psycopg2.connect(
            host=config['host'],
            port=config['port'],
            database=config['database'],
            user=config['username'],
            password=config['password'],
            connect_timeout=10
        )

        cursor = conn.cursor()
        cursor.execute("SELECT version();")
        version = cursor.fetchone()[0]

        cursor.execute("SELECT current_timestamp;")
        timestamp = cursor.fetchone()[0]

        conn.close()

        database_results[db_name] = {
            'connected': True,
            'version': version,
            'timestamp': str(timestamp)
        }

        logger.info(f"✅ {db_name} database connected")

    except Exception as e:
        database_results[db_name] = {
            'connected': False,
            'error': str(e)
        }
    self.overall_health = False

```

```

        logger.error(f"❌ {db_name} database connection failed: {e}")

self.health_results['databases'] = database_results

def check_data_availability(self):
    """Check if required data files are available"""
    logger.info("📁 Checking data availability...")

    from pathlib import Path
    import pandas as pd

    required_files = {
        'ecommerce_raw': 'data/raw/ecommerce_transactions.csv',
        'superstore_raw': 'data/raw/superstore_sales.csv',
        'products_raw': 'data/raw/amazon_products.csv',
        'retail_raw': 'data/raw/retail_analytics.csv'
    }

    data_results = {}

    for file_name, file_path in required_files.items():
        try:
            path = Path(file_path)
            if path.exists():
                # Quick validation
                df = pd.read_csv(file_path, nrows=5)
                file_size = path.stat().st_size / (1024 * 1024) # MB

                data_results[file_name] = {
                    'available': True,
                    'size_mb': round(file_size, 2),
                    'columns': len(df.columns),
                    'sample_rows': len(df)
                }

                logger.info(f"✅ {file_name}: {file_size:.1f}MB, {len(df.columns)}")
            else:
                data_results[file_name] = {
                    'available': False,
                    'error': 'File not found'
                }
                logger.warning(f"⚠️ {file_name}: File not found")

        except Exception as e:

```

```

        data_results[file_name] = {
            'available': False,
            'error': str(e)
        }
        logger.error(f"❌ {file_name}: Error reading file - {e}")

available_files = sum(1 for result in data_results.values() if result.get('ava
total_files = len(required_files)

self.health_results['data_files'] = {
    'files_available': available_files,
    'total_files': total_files,
    'availability_rate': available_files / total_files,
    'file_details': data_results
}

if available_files < total_files:
    logger.warning(f"⚠️ {available_files}/{total_files} data files available"
else:
    logger.info(f"✅ All {total_files} data files available")

def generate_health_report(self):
    """Generate comprehensive health report"""
    logger.info(f"📊 Generating health report...")

    report = {
        'timestamp': datetime.now().isoformat(),
        'overall_health': self.overall_health,
        'checks_performed': len(self.health_results),
        'health_results': self.health_results
    }

    # Calculate health score
    healthy_components = 0
    total_components = 0

    for component, result in self.health_results.items():
        total_components += 1
        if component == 'docker_services':
            healthy_components += 1 if result.get('all_services_running') else 0
        elif component == 'airflow':
            healthy_components += 1 if result.get('webserver_healthy') else 0
        elif component == 'spark':
            healthy_components += 1 if result.get('master_healthy') else 0

```



```

        elif component == 'databases':
            db_healthy = all(db.get('connected', False) for db in result.values())
            healthy_components += 1 if db_healthy else 0
        elif component == 'data_files':
            healthy_components += 1 if result.get('availability_rate', 0) >= 0.8 else 0

    health_score = healthy_components / total_components if total_components > 0 else 0
    report['health_score'] = round(health_score, 4)

    # Save report
    report_path = 'logs/health_report.json'
    with open(report_path, 'w') as f:
        json.dump(report, f, indent=2)

    logger.info(f"📄 Health report saved: {report_path}")
    logger.info(f"📊 Overall health score: {health_score:.1%}")

    return report

def run_complete_health_check(self):
    """Run all health checks"""
    logger.info("🚀 Starting comprehensive platform health check...")

    start_time = time.time()

    # Run all checks
    self.check_docker_services()
    self.check_airflow_health()
    self.check_spark_health()
    self.check_database_connectivity()
    self.check_data_availability()

    # Generate report
    report = self.generate_health_report()

    execution_time = time.time() - start_time
    logger.info(f"🕒 Health check completed in {execution_time:.2f} seconds")

    if self.overall_health:
        logger.info("🎉 Platform health check: ALL SYSTEMS GO!")
        return 0
    else:
        logger.error("⚠️ Platform health check: ISSUES DETECTED")
        return 1

```

```
def main():  
    """Main execution function"""  
    monitor = PlatformHealthMonitor()  
    exit_code = monitor.run_complete_health_check()  
    return exit_code  
  
if __name__ == "__main__":  
    exit(main())
```

Day 14 Success Metrics and Learning Outcomes

Knowledge Integration Assessment

Core Concepts Mastered:

1. Systems Architecture Thinking

- Multi-component integration patterns
- Service discovery and communication
- Dependency management across tools
- Scalability and performance considerations

2. Platform Engineering Fundamentals

- Infrastructure as Code principles
- Environment-specific configuration
- Monitoring and observability
- Automated deployment strategies

3. Data Pipeline Orchestration







- Cross-tool workflow coordination
- Error handling and recovery
- Data lineage and quality gates
- Performance monitoring

4. Production Readiness






- Security best practices
- Logging and alerting
- Backup and recovery
- Documentation and maintenance

Technical Achievements

Week 2 Integration Milestones:






-  8 technology components integrated
-  End-to-end data pipeline operational
-  Multi-source data processing capability
-  Production-ready monitoring system
-  Automated quality validation framework
-  Scalable architecture foundation

Performance Benchmarks:






-  Data Processing: 100K+ records/minute capability
-  System Availability: 99.9% uptime target
-  Pipeline Execution: < 15 minutes end-to-end
-  Error Recovery: < 5 minutes MTTR
-  Scalability: Horizontal scaling ready

Business Value Delivered

Operational Improvements:

-  **Data Freshness:** From daily → hourly updates
-  **Automation Level:** 95% of manual tasks eliminated
-  **Data Quality:** Consistent 99%+ accuracy
-  **Maintenance Effort:** 70% reduction in manual intervention
-  **Cost Efficiency:** Cloud-ready architecture for optimal resource usage

Decision-Making Enablement:

-  Real-time customer segment analysis
-  Cross-platform performance insights
-  Automated anomaly detection
-  Predictive analytics foundation
-  Data-driven business recommendations


Week 3 Preparation

Tomorrow's Advanced Focus:

- **Advanced Pandas Techniques:** Memory optimization and performance
- **Real-time Stream Processing:** Kafka and streaming analytics
- **Machine Learning Integration:** ML pipeline development
- **Advanced Data Quality:** Statistical validation methods
- **Performance Optimization:** Large-scale data processing

Skills Building Path:

- Complex data transformations and aggregations
- Real-time event processing architectures
- Advanced analytics and statistical methods
- Performance profiling and optimization
- Enterprise data governance frameworks

 **Congratulations! You've successfully built your first production-ready integrated data platform. This foundation will support all advanced data engineering concepts in the coming weeks!**