

# Day 23: AWS Glue and Managed ETL - Complete Serverless Processing Guide

## What You'll Learn Today (Serverless-First Approach)

**Primary Focus:** Understanding serverless ETL paradigms and metadata-driven data processing architectures

**Secondary Focus:** AWS Glue service mastery including crawlers, jobs, catalogs, and workflow orchestration

**Dataset for Context:** Complex multi-source enterprise integration scenarios from Kaggle for comprehensive ETL implementation

## Learning Philosophy for Day 23

*"Transform data, not infrastructure"*








We'll master the transition from traditional, infrastructure-heavy ETL processes to serverless, metadata-driven transformations that scale automatically and focus on business logic rather than operational overhead.

## The Serverless ETL Revolution: Why Managed Services Change Everything

### The Problem: Traditional ETL Infrastructure Burden







**Scenario:** Your organization needs to process diverse data sources with varying schemas and volumes...

#### Without Managed ETL (Infrastructure Chaos):

-  Manual Spark cluster provisioning and management
-  Infrastructure scaling decisions based on peak loads
-  Complex deployment and monitoring of ETL jobs
-  Manual schema discovery and catalog maintenance
-  Point-to-point integration creating brittle dependencies
-  Resource waste during low-usage periods
-  DevOps overhead for cluster maintenance and updates

#### With AWS Glue (Serverless ETL):

-  Zero infrastructure management and provisioning

-  Automatic scaling based on data volume and complexity
-  Built-in monitoring, logging, and error handling
-  Automated schema discovery and data cataloging
-  Unified metadata store for all data sources
-  Pay-per-use pricing model with no idle costs
-  Managed service updates and security patches

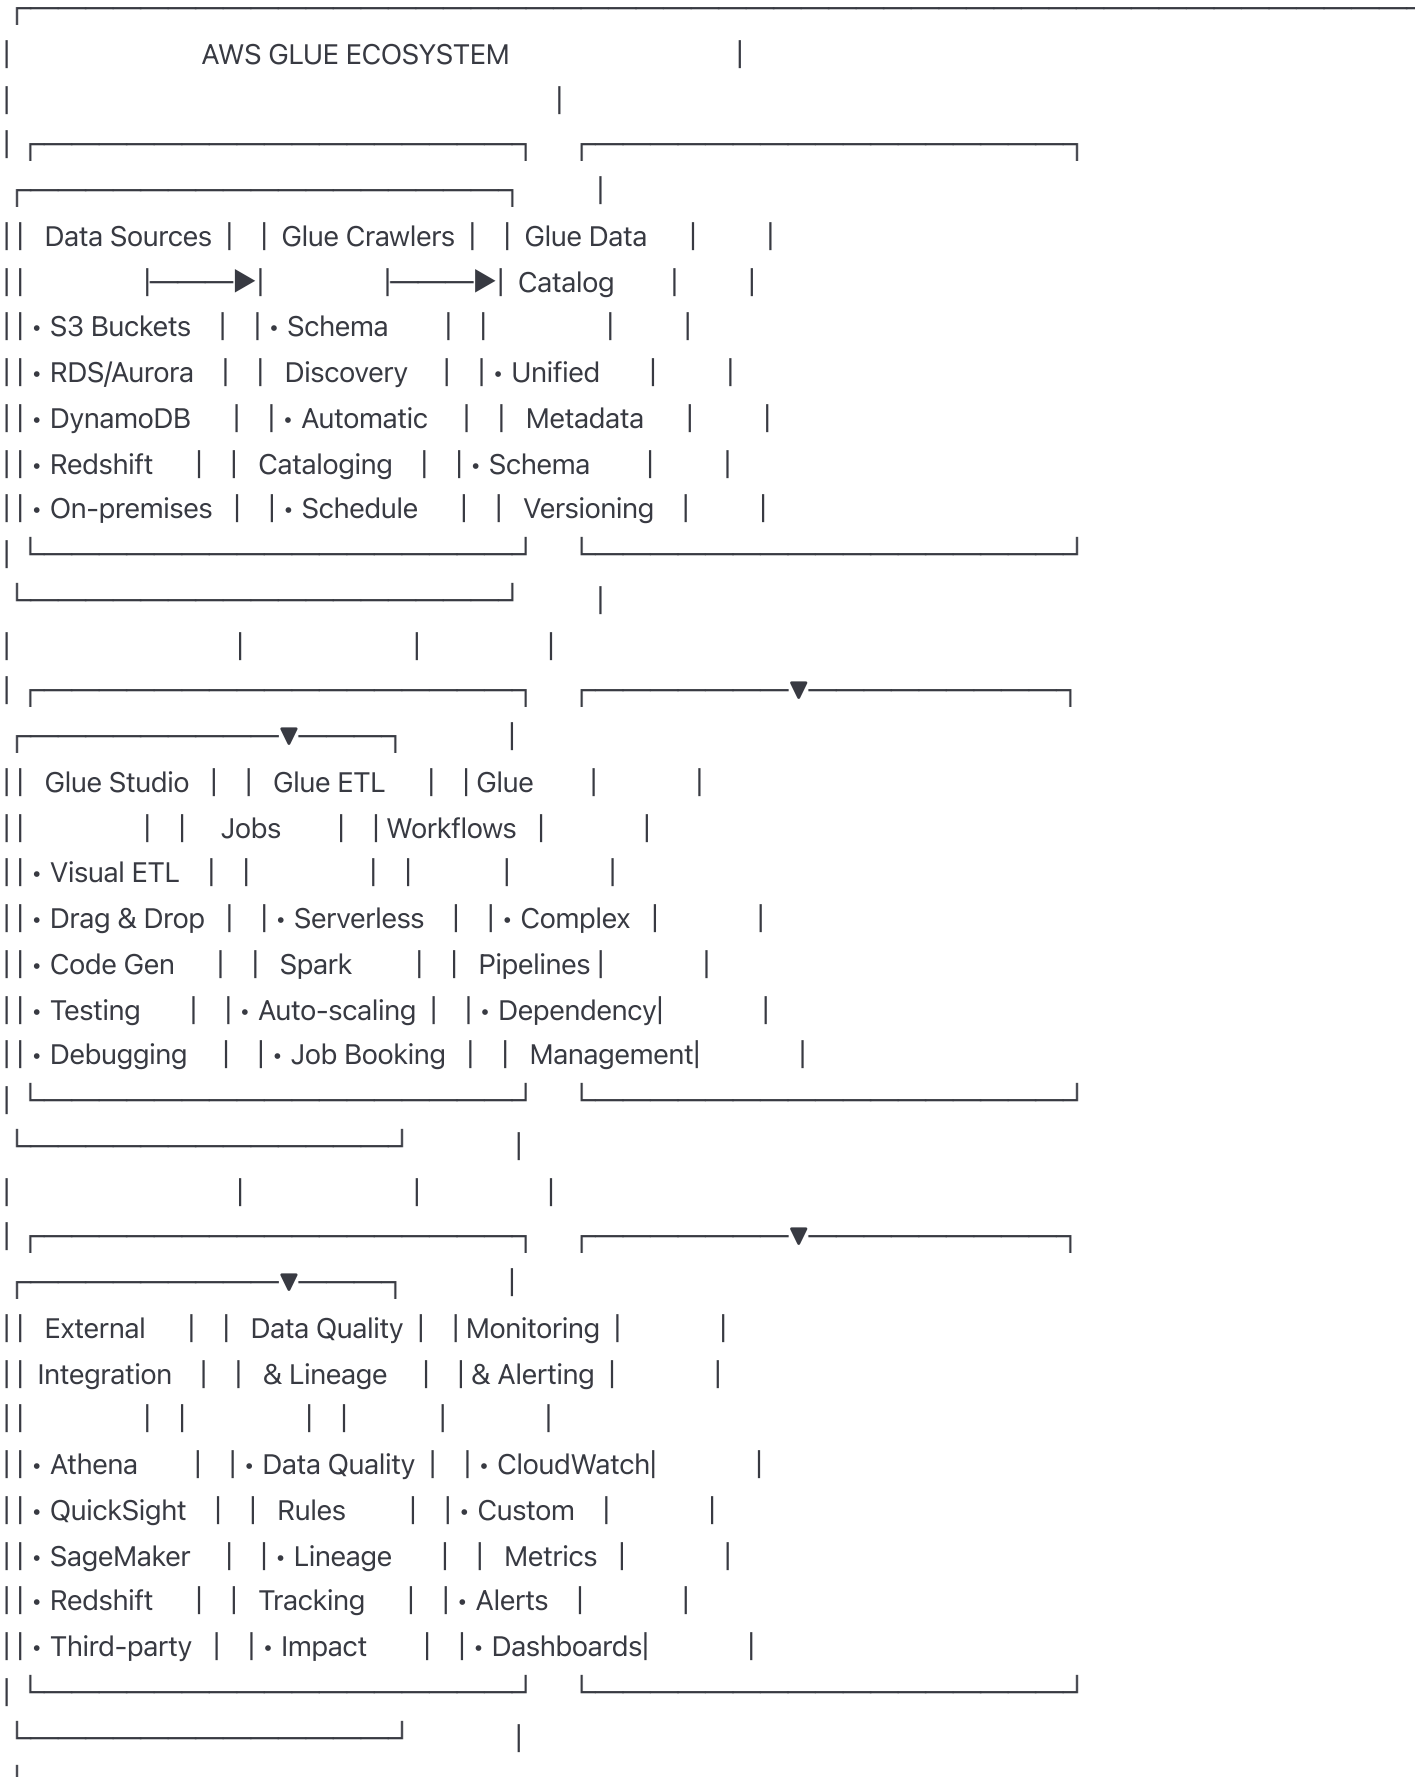
## **AWS Glue Architecture Deep Dive**

### **The Serverless ETL Mental Model**

Think of AWS Glue like a modern smart manufacturing system:

**Traditional ETL:** Like owning and maintaining your own factory with fixed capacity

**AWS Glue:** Like using an on-demand manufacturing service that scales to your exact needs



# 1. Serverless vs Traditional ETL Paradigms

## Traditional ETL Architecture:

Infrastructure-Centric Approach:

Resource Planning:

- └—— Provision Spark clusters based on peak load
- └—— Manage cluster lifecycle (startup, shutdown)
- └—— Handle resource scaling decisions
- └—— Maintain infrastructure security and updates
- └—— Pay for idle capacity during low usage

Development Process:

- └—— Write ETL code for specific cluster configurations
- └—— Package and deploy code to cluster environments
- └—— Manage job scheduling and dependencies
- └—— Handle error recovery and monitoring
- └—— Maintain development/staging/production environments

Operational Overhead:

- └—— Monitor cluster health and performance
- └—— Troubleshoot infrastructure issues
- └—— Plan capacity for growth
- └—— Manage cost optimization
- └—— Handle disaster recovery and backup

## Serverless ETL Architecture:

## Business-Logic-Centric Approach:

### Resource Management:

- └— AWS automatically provisions optimal resources
- └— Instant job startup (no cluster warm-up time)
- └— Automatic scaling based on data volume
- └— Built-in security and managed updates
- └— Pay only for actual job execution time

### Development Process:

- └— Focus on transformation logic, not infrastructure
- └— Visual development with Glue Studio
- └— Automatic code generation and optimization
- └— Built-in testing and debugging capabilities
- └— Seamless deployment across environments

### Operational Benefits:

- └— Automatic monitoring and alerting
- └— Built-in error handling and retry logic
- └— Elastic scaling without capacity planning
- └— Integrated cost optimization
- └— Managed disaster recovery and availability

## **2. Data Catalog as the Foundation**

### **Unified Metadata Management:**

## Data Catalog Concepts:

### Traditional Approach (Fragmented):

- └— Each team maintains their own documentation
- └— Schema information scattered across systems
- └— Manual discovery of new data sources
- └— Inconsistent naming and classification
- └— No centralized lineage tracking

### Glue Data Catalog (Unified):

- └— Single source of truth for all metadata
- └— Automatic schema discovery and updates
- └— Consistent naming and tagging conventions
- └— Centralized lineage and impact analysis
- └— API-driven metadata access for all tools

### Catalog Structure:

Database Level: Logical grouping (e.g., 'retail', 'finance')

- └— Table Level: Individual datasets (e.g., 'customer\_transactions')
- └— Partition Level: Data organization (e.g., 'year=2024/month=01')
- └— Column Level: Field definitions with types and descriptions

## Schema Evolution Management:

## Schema Change Handling:

### Schema Discovery Process:

1. Crawler scans data source
2. Infers schema from sample data
3. Compares with existing catalog
4. Identifies schema changes
5. Updates catalog with new version
6. Maintains backward compatibility

### Change Types Supported:

- └─ Adding new columns (backward compatible)
- └─ Changing column types (with validation)
- └─ Removing columns (impact analysis)
- └─ Renaming columns (with mapping rules)
- └─ Structural changes (nested schema evolution)

### Versioning Strategy:

- └─ Automatic schema version tracking
- └─ Impact analysis for downstream consumers
- └─ Rollback capabilities for problematic changes
- └─ Migration assistance for breaking changes
- └─ Compatibility testing framework

## 3. Auto-scaling and Resource Optimization

### Dynamic Resource Management:

## Glue Job Scaling Concepts:

### Resource Allocation Logic:

- └— Analyzes input data size and complexity
- └— Determines optimal number of workers
- └— Allocates appropriate worker types (memory/compute optimized)
- └— Monitors job progress and adjusts resources
- └— Automatically handles worker failures and replacement

### Worker Types Available:

- └— Standard: Balanced CPU and memory (2 vCPUs, 8GB RAM)
- └— G.1X: Memory optimized (4 vCPUs, 16GB RAM)
- └— G.2X: High memory (8 vCPUs, 32GB RAM)
- └— G.025X: Small jobs (0.25 DPU for cost optimization)
- └— Custom: User-defined configurations for special workloads

### Scaling Examples:

#### Small Dataset (1 GB):

- └— Workers: 2-5 Standard workers
- └— Duration: 5-10 minutes
- └— Cost:  $\$0.44 \text{ per hour} \times 2 \text{ workers} \times 0.17 \text{ hours} = \$0.15$

#### Large Dataset (100 GB):

- └— Workers: 20-50 G.1X workers
- └— Duration: 30-60 minutes
- └— Cost:  $\$0.44 \text{ per hour} \times 30 \text{ workers} \times 1 \text{ hour} = \$13.20$

### Cost Optimization:

- └— Automatic right-sizing based on workload
- └— Spot pricing for non-critical jobs
- └— Intelligent caching for repeated operations
- └— Resource pooling across multiple jobs
- └— Pay-per-second billing (1-minute minimum)

## AWS Glue Components Deep Dive

## Glue Crawlers: Automated Schema Discovery

### Crawler Concepts and Strategy:

### How Crawlers Work:



## Crawler Operation Flow:

### 1. Data Source Connection:

- └── Connect to specified data stores (S3, RDS, etc.)
- └── Apply security credentials and access controls
- └── Scan directory structures or table schemas
- └── Identify data formats and compression types

### 2. Schema Inference:

- └── Sample data from multiple files/partitions
- └── Analyze data types and structures
- └── Identify nested schemas (JSON, Avro)
- └── Detect partition patterns
- └── Infer relationships and constraints

### 3. Catalog Updates:

- └── Compare discovered schema with existing catalog
- └── Identify additions, modifications, deletions
- └── Create new tables or update existing ones
- └── Maintain schema versioning and history
- └── Trigger downstream notifications

### 4. Optimization:

- └── Group similar files into single table definitions
- └── Optimize partition schemes for query performance
- └── Suggest compression and format improvements
- └── Identify data quality issues and anomalies

## Advanced Crawler Configuration:

## Crawler Customization Options:

### Schema Change Handling:

- └—— Update table definition in catalog
- └—— Add new columns to existing table
- └—— Ignore schema changes (maintain stability)
- └—— Create new table version
- └—— Custom change detection logic

### Data Classification:

- └—— Automatic PII detection and tagging
- └—— Sensitive data classification
- └—— Custom classifier creation
- └—— Data format recognition
- └—— Quality scoring and metrics

### Performance Optimization:

- └—— Incremental crawling for large datasets
- └—— Custom sampling strategies
- └—— Parallel crawling for multiple sources
- └—— Resource allocation for crawler jobs
- └—— Schedule optimization for minimal impact

### Example Crawler Configuration:

Target: S3 bucket with retail transaction data

Schedule: Daily at 2 AM (after data ingestion)

Include Paths: s3://company-data-lake/transactions/

Exclude Patterns: \*.log, \*.tmp, \*\_backup/

Schema Change Policy: Update table and add new columns

Output Database: retail\_analytics\_catalog

## **Glue ETL Jobs: Serverless Data Transformation**

### ETL Job Architecture:

### Job Types and Use Cases:

## Glue Job Categories:

### Spark ETL Jobs (Most Common):

- └── Distributed processing for large datasets
- └── Complex transformations and aggregations
- └── Multi-source data joins and merging
- └── Schema evolution and data quality checks
- └── Batch processing with high throughput

### Python Shell Jobs:

- └── Lightweight processing for small datasets
- └── API integrations and web scraping
- └── Simple data format conversions
- └── Orchestration and workflow coordination
- └── Cost-effective for < 1 GB datasets

### Ray Jobs (Preview):

- └── Machine learning data preprocessing
- └── Distributed computing for AI/ML workloads
- └── Python-native ecosystem integration
- └── Hyperparameter tuning and model training
- └── Advanced analytics and scientific computing

### Streaming Jobs:

- └── Real-time data processing from Kinesis
- └── Change data capture (CDC) processing
- └── Low-latency transformations
- └── Continuous aggregations and windowing
- └── Integration with streaming analytics

## ETL Job Development Patterns:

## Development Approaches:

### Visual ETL with Glue Studio:

- └— Drag-and-drop interface for non-developers
- └— Pre-built transformation blocks
- └— Automatic code generation
- └— Visual debugging and data preview
- └— Rapid prototyping and development

### Code-Based Development:

- └— Full PySpark/Scala programming flexibility
- └— Custom transformation logic
- └— Advanced optimization techniques
- └— Version control and collaborative development
- └— Complex business rule implementation

### Hybrid Approach:

- └— Visual design for standard transformations
- └— Custom code blocks for specialized logic
- └— Reusable transformation libraries
- └— Team collaboration across skill levels
- └— Maintainable and documented pipelines

### Performance Optimization Techniques:

- └— Broadcast joins for small dimension tables
- └— Partition-wise operations for large datasets
- └— Columnar storage format optimization
- └— Caching strategies for repeated operations
- └— Resource allocation and worker configuration

## **Glue Workflows: Complex Pipeline Orchestration**

### **Workflow Management Concepts:**

### **Workflow Architecture Patterns:**

## Pipeline Orchestration Strategies:


### Simple Linear Workflow:

Crawler → Data Quality Check → ETL Transform → Data Validation → Catalog Update

### Parallel Processing Workflow:



### Conditional Workflow:



### Event-Driven Workflow:

S3 Event → Lambda Trigger → Glue Workflow → Success/Failure Notification

### Complex Enterprise Workflow:

Multiple Sources → Schema Discovery → Data Quality → Parallel ETL →  
Data Validation → Error Handling → Notification → Catalog Update →  
Analytics Refresh → Dashboard Update

## Workflow Error Handling:

## Resilience and Recovery Patterns:

### Error Detection:

- └── Job failure monitoring
- └── Data quality threshold breaches
- └── Schema compatibility issues
- └── Resource allocation failures
- └── External dependency failures

### Recovery Strategies:

- └── Automatic retry with exponential backoff
- └── Alternative processing paths
- └── Partial data processing continuation
- └── Rollback to last known good state
- └── Manual intervention triggers

### Notification Systems:

- └── SNS integration for real-time alerts
- └── CloudWatch alarms for system metrics
- └── Custom dashboards for operational visibility
- └── Escalation procedures for critical failures
- └── Audit trails for compliance and debugging

### Example Error Handling:

Job: Customer data ETL pipeline

Error: Source database connection timeout

Action: Retry 3 times with 5-minute delays

Fallback: Process previous day's cached data

Notification: Alert data team via Slack

Recovery: Manual database connectivity check

## **Phase 1: Setting Up Enterprise ETL Environment**

### **Dataset Preparation and Multi-Source Integration**

#### **Complex Integration Scenario:**

#### **Enterprise Data Sources:**

## Multi-Source Integration Dataset:

### Primary Sources:

- └── Customer Database (PostgreSQL RDS)
  - | └── Tables: customers, accounts, preferences
  - | └── Size: 10 million records
  - | └── Update Frequency: Real-time
  - | └── Schema: Normalized relational model
  
- └── Transaction Logs (S3 JSON)
  - | └── Format: JSON Lines streaming data
  - | └── Size: 1 billion transactions/month
  - | └── Update Frequency: Continuous streaming
  - | └── Schema: Semi-structured event data
  
- └── Product Catalog (S3 Parquet)
  - | └── Format: Parquet with nested structures
  - | └── Size: 1 million products
  - | └── Update Frequency: Daily batch updates
  - | └── Schema: Complex nested attributes
  
- └── Web Analytics (CloudTrail Logs)
  - | └── Format: CloudTrail JSON logs
  - | └── Size: 100 GB daily
  - | └── Update Frequency: Hourly aggregation
  - | └── Schema: AWS standard log format
  
- └── External APIs (Third-party enrichment)
  - | └── Format: REST API responses
  - | └── Size: Variable based on lookups
  - | └── Update Frequency: On-demand
  - | └── Schema: Vendor-specific formats

## Kaggle Dataset Enhancement:

## Supplementary Datasets for Practice:

### Base Dataset: E-commerce Customer Analytics

- └── Source: [kaggle.com/datasets/imakash3011/customer-personality-analysis](https://kaggle.com/datasets/imakash3011/customer-personality-analysis)
- └── Size: 2,240 customers with 29 attributes
- └── Format: CSV with mixed data types
- └── Use Case: Schema evolution and transformation practice

### Enhancement Datasets:

- └── Sales Transactions: [kaggle.com/datasets/gabrielramos87/an-online-shop-business](https://kaggle.com/datasets/gabrielramos87/an-online-shop-business)
- └── Product Categories: [kaggle.com/datasets/srolka/ecommerce-products](https://kaggle.com/datasets/srolka/ecommerce-products)
- └── Customer Reviews: [kaggle.com/datasets/nickmccullum/ecommerce-customer-reviews](https://kaggle.com/datasets/nickmccullum/ecommerce-customer-reviews)
- └── Geographic Data: [kaggle.com/datasets/htagholdings/property-sales](https://kaggle.com/datasets/htagholdings/property-sales)
- └── Seasonal Patterns: [kaggle.com/datasets/rohitsahoo/sales-forecasting](https://kaggle.com/datasets/rohitsahoo/sales-forecasting)

### Integration Challenges:

- └── Different file formats (CSV, JSON, Parquet)
- └── Varying schema structures and naming conventions
- └── Date format inconsistencies across sources
- └── Missing data and quality issues
- └── Size variations requiring different processing strategies

## Project Structure for Glue ETL Pipeline



```
aws-glue-etl-pipeline/
├── infrastructure/
│   ├── terraform/
│   │   ├── glue_catalog.tf
│   │   ├── glue_jobs.tf
│   │   ├── glue_crawlers.tf
│   │   ├── glue_workflows.tf
│   │   ├── iam_roles.tf
│   │   ├── s3_buckets.tf
│   │   └── cloudwatch_monitoring.tf
│   └── cloudformation/
│       ├── glue-etl-stack.yaml
│       └── monitoring-stack.yaml
├── data/
│   ├── raw/
│   │   ├── customers/
│   │   ├── transactions/
│   │   ├── products/
│   │   └── logs/
│   ├── staging/
│   │   ├── cleaned/
│   │   ├── validated/
│   │   └── enriched/
│   ├── processed/
│   │   ├── curated/
│   │   ├── aggregated/
│   │   └── analytics_ready/
│   ├── schemas/
│   │   ├── source_schemas/
│   │   ├── target_schemas/
│   │   └── evolution_history/
├── etl_jobs/
├── visual_jobs/
│   ├── customer_data_cleansing.json
│   ├── transaction_aggregation.json
│   └── product_enrichment.json
├── python_jobs/
│   ├── advanced_transformations.py
│   ├── data_quality_validation.py
│   ├── schema_evolution_handler.py
│   └── external_api_integration.py
├── shared_libraries/
└── transformation_utils.py
```

- | | |—— data\_quality\_rules.py
- | | |—— schema\_mapping.py
- | | |—— monitoring\_helpers.py
- | |—— tests/
  - | |—— unit\_tests/
  - | |—— integration\_tests/
  - | |—— data\_validation\_tests/
- |—— workflows/
  - | |—— daily\_batch\_processing.json
  - | |—— real\_time\_streaming.json
  - | |—— schema\_evolution\_pipeline.json
  - | |—— data\_quality\_monitoring.json
- |—— monitoring/
  - | |—— cloudwatch\_dashboards/
    - | | |—— etl\_performance\_dashboard.json
    - | | |—— data\_quality\_dashboard.json
    - | | |—— cost\_optimization\_dashboard.json
  - | |—— alerts/
    - | | |—— job\_failure\_alerts.yaml
    - | | |—— performance\_degradation\_alerts.yaml
    - | | |—— cost\_anomaly\_alerts.yaml
  - | |—— custom\_metrics/
- |—— config/
  - | |—— environment\_configs/
    - | | |—— dev\_config.yaml
    - | | |—— staging\_config.yaml
    - | | |—— prod\_config.yaml
  - | |—— data\_catalog\_configs/
    - | | |—— database\_definitions.yaml
    - | | |—— table\_schemas.yaml
    - | | |—— partition\_strategies.yaml
  - | |—— security\_configs/
    - | | |—— iam\_policies.json
    - | | |—— encryption\_settings.yaml
    - | | |—— network\_configurations.yaml
- |—— docs/
  - | |—— architecture/
    - | | |—— etl\_architecture\_overview.md
    - | | |—— data\_flow\_diagrams.md
    - | | |—— security\_architecture.md
  - | |—— operations/
    - | | |—— deployment\_procedures.md
    - | | |—— monitoring\_runbooks.md
    - | | |—— troubleshooting\_guides.md

```
| └── development/
|   ├── coding_standards.md
|   ├── testing_strategies.md
|   └── performance_optimization.md
└── scripts/
    ├── deployment/
    |   ├── deploy_etl_jobs.py
    |   ├── update_crawlers.py
    |   └── setup_monitoring.py
    ├── utilities/
    |   ├── schema_comparison.py
    |   ├── data_profiling.py
    |   └── cost_analysis.py
    └── automation/
        ├── ci_cd_pipeline.py
        ├── automated_testing.py
        └── performance_benchmarking.py
```

## Phase 2: Understanding Metadata-Driven ETL

### Data Catalog Concepts and Architecture

#### Metadata-Driven Development Paradigm:

#### Traditional ETL vs Catalog-Driven ETL:

## Traditional ETL Approach:

### Development Process:

- └ Manual schema discovery and documentation
- └ Hard-coded schema definitions in ETL jobs
- └ Point-to-point data source connections
- └ Manual impact analysis for schema changes
- └ Separate metadata management for each pipeline

### Maintenance Challenges:

- └ Schema drift detection requires manual monitoring
- └ Breaking changes cause pipeline failures
- └ No centralized view of data lineage
- └ Difficult to track data quality across sources
- └ Duplication of transformation logic

## Catalog-Driven ETL Approach:

### Development Process:

- └ Automated schema discovery via crawlers
- └ Dynamic schema retrieval from catalog
- └ Centralized metadata store for all sources
- └ Automatic impact analysis and notifications
- └ Reusable transformation patterns

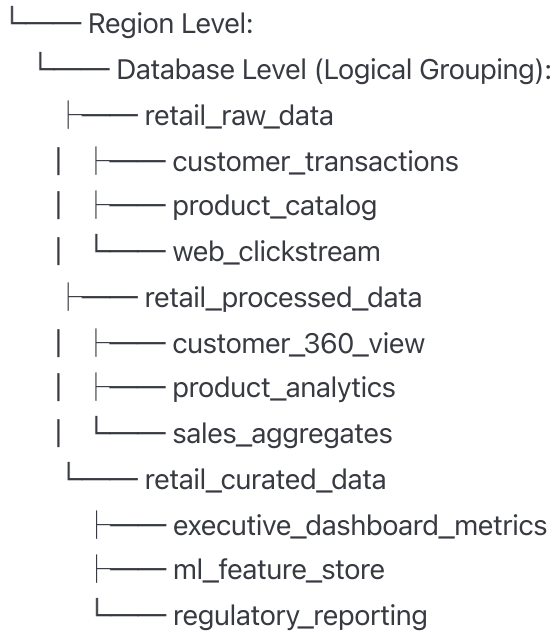
### Maintenance Benefits:

- └ Automatic schema evolution handling
- └ Proactive change notifications
- └ Visual data lineage tracking
- └ Centralized data quality monitoring
- └ Shared transformation libraries and patterns

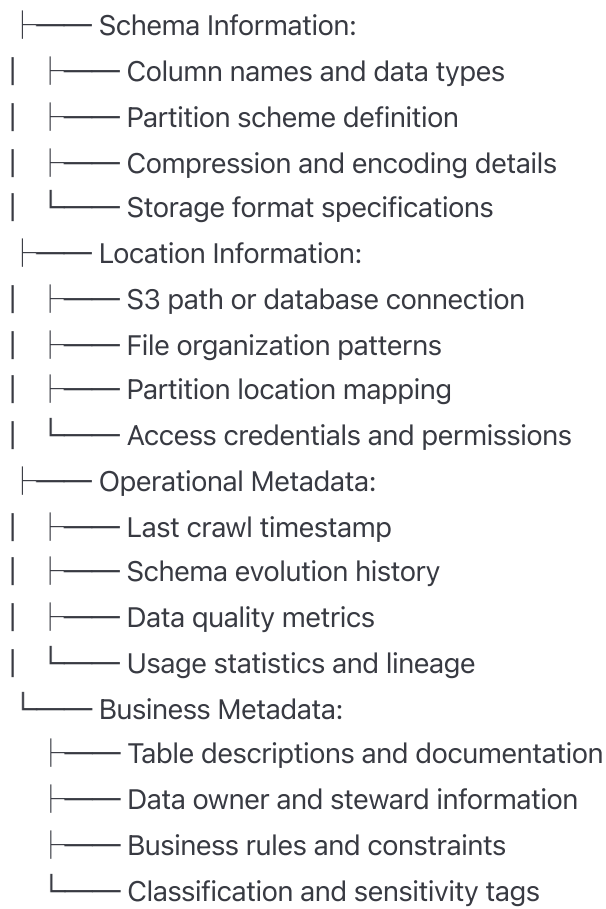
## Data Catalog Architecture Deep Dive:

## Glue Data Catalog Structure:

### Account Level:



### Table Metadata Structure:



## Schema Evolution and Version Management

### Schema Change Management Strategy:

**Schema Evolution Patterns:**

## Evolution Types and Handling:

### Backward Compatible Changes (Safe):

- |—— Adding optional columns
- |—— Expanding field sizes (varchar(50) → varchar(100))
- |—— Adding new nested fields in JSON
- |—— New partition columns
- └—— Additional metadata attributes

### Handling Approach:

1. Crawler detects new schema automatically
2. Catalog updated with new column definitions
3. Existing ETL jobs continue working unchanged
4. New transformations can leverage additional fields
5. Gradual adoption across downstream consumers

### Forward Compatible Changes (Requires Planning):

- |—— Removing optional columns
- |—— Renaming fields (with backward compatibility)
- |—— Changing optional field types
- |—— Restructuring nested schemas
- └—— Modifying partition schemes

### Handling Approach:

1. Implement alias mapping for renamed fields
2. Maintain multiple schema versions temporarily
3. Coordinate changes across ETL jobs
4. Implement gradual migration strategy
5. Monitor impact on downstream systems

### Breaking Changes (Avoid When Possible):

- |—— Removing required fields
- |—— Incompatible type changes (string → number)
- |—— Changing field semantics
- |—— Major restructuring of data model
- └—— Removing or changing primary keys

### Handling Approach:

1. Create new table version alongside existing
2. Implement dual-write during transition period
3. Migrate consumers gradually to new schema

4. Validate data consistency between versions
5. Deprecate old schema after full migration

### **Practical Schema Evolution Example:**



## Customer Data Schema Evolution:

### Version 1.0 (Initial):

```
{
  "customer_id": "string",
  "name": "string",
  "email": "string",
  "registration_date": "date"
}
```

### Version 1.1 (Backward Compatible):

```
{
  "customer_id": "string",
  "name": "string",
  "email": "string",
  "phone": "string",      // New optional field
  "registration_date": "date",
  "last_login": "timestamp" // New optional field
}
```

### Version 2.0 (Breaking Change):

```
{
  "customer_id": "string",
  "first_name": "string",  // Breaking: split from "name"
  "last_name": "string",  // Breaking: split from "name"
  "contact": {            // Breaking: restructured
    "email": "string",
    "phone": "string",
    "preferences": {
      "email_notifications": "boolean",
      "sms_notifications": "boolean"
    }
  },
  "registration_date": "date",
  "last_login": "timestamp"
}
```

### Migration Strategy:

1. Deploy Version 2.0 to new table (customer\_data\_v2)
2. Create transformation job to convert v1 → v2 format
3. Run dual-write to both versions during transition
4. Update ETL jobs to consume from v2 table
5. Validate data consistency and business logic

6. Migrate all consumers to v2 schema
7. Deprecate v1 table after validation period

Glue Catalog Handling:

- |—— Automatic detection of schema changes
- |—— Version history maintenance
- |—— Impact analysis across dependent jobs
- |—— Migration assistance and recommendations
- |—— Rollback capabilities for failed migrations

## **Dynamic Transformation Logic**

**Metadata-Driven Transformation Concepts:**

**Configuration-Driven ETL Jobs:**

## Dynamic Transformation Architecture:

### Static ETL Approach (Traditional):

```
def transform_customer_data():  
    # Hard-coded transformation logic  
    df = spark.read.table("source_customers")  
    df = df.withColumn("full_name",  
                       concat(col("first_name"), lit(" "), col("last_name")))  
    df = df.filter(col("status") == "active")  
    df.write.mode("overwrite").table("target_customers")
```

### Dynamic ETL Approach (Metadata-Driven):

```
def transform_data_dynamically(source_table, target_table, config):  
    # Read transformation rules from catalog metadata  
    transformation_rules = get_transformation_config(source_table)  
  
    df = spark.read.table(source_table)  
  
    # Apply transformations based on metadata  
    for rule in transformation_rules:  
        if rule["type"] == "column_combination":  
            df = df.withColumn(rule["target_column"],  
                               expr(rule["expression"]))  
        elif rule["type"] == "filter":  
            df = df.filter(expr(rule["condition"]))  
        elif rule["type"] == "data_type_conversion":  
            df = df.withColumn(rule["column"],  
                               col(rule["column"]).cast(rule["target_type"]))  
  
    df.write.mode(config["write_mode"]).table(target_table)
```

### Benefits of Dynamic Approach:

- |—— Business users can modify transformation rules
- |—— No code changes required for rule updates
- |—— A/B testing of different transformation logic
- |—— Automated optimization based on data patterns
- |—— Consistent transformation patterns across pipelines

## Rule-Based Transformation Engine:

## Transformation Rule Categories:

### Data Quality Rules:

- ├── Null value handling strategies
- ├── Duplicate detection and resolution
- ├── Data format standardization
- ├── Range and constraint validations
- └── Cross-field consistency checks

### Business Logic Rules:

- ├── Derived field calculations
- ├── Categorical data mapping
- ├── Conditional transformations
- ├── Aggregation and summarization
- └── Data enrichment from external sources

### Technical Optimization Rules:

- ├── Data type optimization for storage
- ├── Partition key generation
- ├── Compression strategy selection
- ├── Format conversion (CSV → Parquet)
- └── Index and clustering recommendations

### Example Rule Configuration:

```
{
  "table": "customer_transactions",
  "rules": [
    {
      "type": "data_quality",
      "rule": "null_handling",
      "column": "transaction_amount",
      "action": "replace_with_zero"
    },
    {
      "type": "business_logic",
      "rule": "derived_field",
      "expression": "transaction_amount * tax_rate",
      "target_column": "total_amount"
    },
    {
      "type": "optimization",
      "rule": "partition_key",
      "expression": "date_format(transaction_date, 'yyyy-MM')",
```

```
    "target_column": "year_month"  
  }  
]  
}
```

## ⚡ Phase 3: Advanced Glue ETL Patterns

### 🔄 Complex Multi-Source Integration

**Enterprise Integration Patterns:**

**Pattern 1: Customer 360 View Assembly:**

## Multi-Source Customer Integration:

### Data Sources:

- |—— CRM System (Salesforce API)
  - | |—— Customer basic information
  - | |—— Account status and preferences
  - | |—— Sales representative assignments
- |—— Transaction Database (PostgreSQL)
  - | |—— Purchase history and patterns
  - | |—— Payment methods and billing
  - | |—— Transaction frequency and amounts
- |—— Support System (ServiceNow)
  - | |—— Support ticket history
  - | |—— Issue categories and resolutions
  - | |—— Customer satisfaction scores
- |—— Web Analytics (S3 Logs)
  - | |—— Website interaction patterns
  - | |—— Product view and search history
  - | |—— Session duration and engagement
- |—— Marketing Platform (External API)
  - | |—— Campaign engagement data
  - | |—— Email and social media interactions
  - | |—— Segmentation and targeting information

### Integration Challenges:

- |—— Different customer identifiers across systems
- |—— Varying data update frequencies
- |—— Schema differences and data quality issues
- |—— API rate limits and availability constraints
- |—— Real-time vs batch processing requirements

### Glue ETL Solution:

1. Crawlers discover schemas from each source
2. Identity resolution job matches customers across systems
3. Incremental processing handles different update frequencies
4. Data quality validation ensures consistency
5. Unified customer profile generated for analytics

## Pattern 2: Real-time and Batch Integration:

## Hybrid Processing Architecture:

### Real-time Stream (Kinesis → Glue Streaming):

- |—— Transaction events as they occur
- |—— Immediate fraud detection and alerting
- |—— Real-time personalization data
- |—— Low-latency aggregations (1-minute windows)
- └—— Hot data storage for immediate access

### Batch Processing (S3 → Glue ETL):

- |—— Complete historical data reprocessing
- |—— Complex analytics and ML model training
- |—— Comprehensive data quality validation
- |—— Daily/weekly aggregations and reporting
- └—— Cold data optimization and archival

### Integration Points:

- |—— Lambda architecture with serving layer merge
- |—— Consistent schema and data formats
- |—— Coordinated processing to avoid conflicts
- |—— Unified monitoring and alerting
- └—— Cost optimization across hot and cold paths

### Example Implementation:

Real-time: Customer behavior scoring for immediate personalization

Batch: Historical pattern analysis for predictive modeling

Merge: Combined real-time and historical scores for recommendations

## Performance Optimization Strategies

### Glue Job Performance Tuning:

### Memory and Compute Optimization:

## Resource Allocation Strategies:

### Job Size Assessment:

- |—— Small Jobs (<1 GB): G.025X workers, minimal parallelism
- |—— Medium Jobs (1-10 GB): Standard workers, moderate parallelism
- |—— Large Jobs (10-100 GB): G.1X workers, high parallelism
- |—— XL Jobs (100+ GB): G.2X workers, maximum parallelism
- |—— Memory-intensive: Custom worker types with high RAM

## Performance Tuning Parameters:

### Spark Configuration Optimization:

```
spark.sql.adaptive.enabled = true
spark.sql.adaptive.coalescePartitions.enabled = true
spark.sql.adaptive.skewJoin.enabled = true
spark.serializer = org.apache.spark.serializer.KryoSerializer
spark.sql.hive.metastorePartitionPruning = true
```

### Glue-Specific Optimizations:

```
--enable-metrics = true
--enable-continuous-cloudwatch-log = true
--enable-spark-ui = true
--job-bookmark-option = job-bookmark-enable
--enable-glue-datacatalog = true
```

### Data Processing Optimizations:

- |—— Broadcast joins for small dimension tables (<200 MB)
- |—— Bucket joins for large fact tables
- |—— Partition pruning through predicate pushdown
- |—— Columnar format usage (Parquet) for analytics
- |—— Compression optimization (SNAPPY for performance)

## I/O and Storage Optimization:



## Data Access Pattern Optimization:

### Read Optimization:

- └— Partition pruning: Filter by partition columns first
- └— Column pruning: Select only required columns
- └— Predicate pushdown: Apply filters at storage level
- └— File consolidation: Combine small files into larger ones
- └— Cache frequently accessed data in memory

### Write Optimization:

- └— Dynamic partitioning: Let Spark determine partitions
- └— Bucket writing: Pre-sort data for downstream queries
- └— Compression: Use appropriate codec for use case
- └— File size targeting: Aim for 100-1000 MB files
- └— Atomic writes: Use staging locations for consistency

### Storage Format Selection:

CSV: Simple structure, human-readable, but inefficient

Parquet: Columnar, compressed, excellent for analytics

Delta: ACID transactions, time travel, concurrent access

Avro: Schema evolution, cross-language compatibility

ORC: Optimized for Hive/Presto, good compression

### Performance Comparison Example:

Query: `SELECT AVG(amount) FROM transactions WHERE date >= '2024-01-01'`

#### CSV (10 GB):

- └— Read time: 45 seconds
- └— Data scanned: 10 GB (full scan)
- └— Memory usage: 8 GB
- └— Cost: \$0.50 per query

#### Parquet (2 GB compressed):

- └— Read time: 3 seconds
- └— Data scanned: 500 MB (column pruning)
- └— Memory usage: 2 GB
- └— Cost: \$0.025 per query

Improvement: 15x faster, 20x cheaper

## Data Quality and Validation Framework

### Comprehensive Data Quality Strategy:

**Quality Validation Layers:**

## Multi-Layer Quality Framework:

### Ingestion Quality (Source Validation):

- └── Schema compliance checking
- └── Required field presence validation
- └── Data type and format verification
- └── Range and constraint validation
- └── Referential integrity checks

### Transformation Quality (Process Validation):

- └── Pre-transformation data profiling
- └── Transformation logic validation
- └── Post-transformation consistency checks
- └── Business rule compliance verification
- └── Statistical distribution analysis

### Output Quality (Target Validation):

- └── Completeness verification (record counts)
- └── Accuracy validation (sample data checks)
- └── Freshness monitoring (data recency)
- └── Consistency checks (cross-system validation)
- └── Business metric validation (KPI alignment)

### Example Quality Rules Implementation:

```
def validate_customer_data(df):
```

```
    quality_results = {}
```

```
    # Completeness checks
```

```
    total_records = df.count()
```

```
    null_customer_ids = df.filter(col("customer_id").isNull()).count()
```

```
    quality_results["customer_id_completeness"] = 1 - (null_customer_ids / total_records)
```

```
    # Accuracy checks
```

```
    valid_emails = df.filter(col("email").rlike(r'^[a-zA-Z0-9._%+-]+@[a-zA-Z0-9.-]+\.[a-zA-Z]{2,}')).count()
```

```
    quality_results["email_accuracy"] = valid_emails / total_records
```

```
    # Consistency checks
```

```
    age_vs_birth_date = df.filter(
```

```
        abs(datediff(current_date(), col("birth_date")) / 365 - col("age")) <= 1
```

```
    ).count()
```

```
    quality_results["age_consistency"] = age_vs_birth_date / total_records
```

```
    # Business rule validation
```

```

active_customers_with_transactions = df.join(
    spark.table("transactions"), "customer_id"
).filter(col("status") == "active").count()
quality_results["business_rule_compliance"] = active_customers_with_transactions / total_records

return quality_results

```

## Automated Quality Monitoring:

Quality Monitoring and Alerting:

Threshold-Based Alerting:

- └── Data completeness < 95% → Critical alert
- └── Schema compliance < 99% → High priority alert
- └── Record count deviation > 20% → Medium priority alert
- └── Processing time > 2x baseline → Performance alert
- └── Cost per record > threshold → Cost optimization alert

Trend Analysis:

- └── Week-over-week quality score trends
- └── Seasonal pattern recognition in data quality
- └── Correlation between source system changes and quality
- └── Impact analysis of ETL job modifications
- └── Predictive quality degradation detection

Quality Dashboard Metrics:

- └── Overall data quality score (weighted average)
- └── Quality trends by data source and pipeline
- └── Top quality issues and their business impact
- └── Quality SLA compliance tracking
- └── Cost of poor quality measurements

Automated Remediation:

- └── Quarantine bad records for manual review
- └── Apply default values for missing required fields
- └── Route data through alternative processing paths
- └── Trigger data source validation workflows
- └── Scale processing resources for volume spikes

## Phase 4: Glue Workflows and Orchestration

### Complex Pipeline Design

**Enterprise Workflow Patterns:**

**Multi-Stage Pipeline Architecture:**

## Enterprise ETL Workflow Design:

### Stage 1: Data Ingestion and Discovery

- └── Crawler jobs scan new data sources
- └── Schema validation and evolution detection
- └── Data source health checks and connectivity tests
- └── Initial data profiling and quality assessment
- └── Metadata catalog updates and notifications

### Stage 2: Data Quality and Validation

- └── Comprehensive data quality rule execution
- └── Anomaly detection and outlier identification
- └── Cross-source consistency validation
- └── Business rule compliance checking
- └── Quality metrics calculation and trending

### Stage 3: Data Transformation and Integration

- └── Multi-source data integration and joining
- └── Complex business logic application
- └── Data enrichment from external sources
- └── Format standardization and optimization
- └── Derived field calculation and aggregation

### Stage 4: Data Validation and Testing

- └── Post-transformation quality validation
- └── Business metric verification
- └── Sample data comparison with expected results
- └── Performance benchmark validation
- └── Regression testing against historical baselines

### Stage 5: Data Publication and Notification

- └── Data deployment to target systems
- └── Catalog metadata updates
- └── Downstream system notifications
- └── Success/failure reporting
- └── Performance metrics publishing

### Workflow Dependencies:

Sequential: Each stage must complete before next begins

Parallel: Independent transformations can run simultaneously

Conditional: Routing based on data quality or business rules

Fan-out: Single input spawning multiple processing paths

Fan-in: Multiple inputs merging into single output

## **Error Handling and Recovery Strategies:**

## Resilient Workflow Design:

### Error Classification:

- |—— Transient Errors (Network, temporary resource issues)
  - | |—— Automatic retry with exponential backoff
  - | |—— Circuit breaker pattern for external dependencies
  - | |—— Alternative resource allocation strategies
  - | |—— Graceful degradation to backup processing paths
- |—— Data Quality Errors (Schema violations, business rule failures)
  - | |—— Quarantine invalid records for review
  - | |—— Continue processing valid records
  - | |—— Generate quality reports for data stewards
  - | |—— Apply correction rules where possible
- |—— System Errors (Resource exhaustion, configuration issues)
  - | |—— Immediate escalation to operations team
  - | |—— Rollback to last known good state
  - | |—— Resource scaling or job re-sizing
  - | |—— Emergency manual intervention procedures
- |—— Business Logic Errors (Incorrect transformations, calculation errors)
  - |—— Stop processing to prevent data corruption
  - |—— Compare outputs with historical patterns
  - |—— Validate against business rules and expectations
  - |—— Require manual approval before continuation

### Recovery Patterns:

#### 1. Checkpoint and Resume:

- |—— Save intermediate processing state
- |—— Resume from last successful checkpoint
- |—— Avoid reprocessing completed work
- |—— Minimize recovery time and resource usage

#### 2. Compensating Transactions:

- |—— Define rollback procedures for each step
- |—— Implement idempotent operations
- |—— Maintain audit trail of all changes
- |—— Enable precise rollback to any point

#### 3. Dead Letter Queue:

- |—— Route problematic records to separate processing
- |—— Manual review and correction workflows
- |—— Batch reprocessing of corrected data
- |—— Learning from error patterns for prevention



## **Event-Driven Orchestration**

**Reactive Pipeline Architecture:**

**Event-Driven Workflow Triggers:**

## Event-Based Processing Patterns:

### S3 Event Triggers:

- └── New file arrival triggers immediate processing
- └── File size thresholds determine processing strategy
- └── File format detection routes to appropriate transformations
- └── Partition-aware processing for incremental updates
- └── Batch accumulation for efficient processing

### Database Change Events:

- └── CDC (Change Data Capture) streams from source databases
- └── Real-time synchronization with data lake
- └── Incremental processing of only changed records
- └── Conflict resolution for concurrent updates
- └── Order preservation for transaction consistency

### Schedule-Based Events:

- └── Business hour processing for operational reports
- └── Off-peak processing for large batch jobs
- └── End-of-month processing for financial reporting
- └── Seasonal adjustments for holiday periods
- └── Maintenance window scheduling for system updates

### External System Events:

- └── API webhooks from third-party systems
- └── Message queue integration (SQS, SNS)
- └── Cloud service notifications (RDS, DynamoDB)
- └── Monitoring system alerts triggering diagnostics
- └── Business process events from workflow systems

## Event Processing Architecture:

Event Source → Event Router → Workflow Trigger → Glue Job → Success/Failure Handler

### Example Event-Driven Flow:

1. New transaction file arrives in S3
2. S3 event notification sent to SQS queue
3. Lambda function processes event and determines workflow
4. Glue workflow triggered with appropriate parameters
5. Parallel processing jobs for different data types
6. Aggregation job combines results
7. Quality validation and business rule checking

- 8. Publication to downstream systems
- 9. Notification to business stakeholders

## Advanced Orchestration Patterns:

Sophisticated Workflow Management:

Dynamic Workflow Generation:

- ├── Metadata-driven workflow creation
- ├── Runtime workflow modification based on data characteristics
- ├── A/B testing different processing approaches
- ├── Adaptive resource allocation based on workload
- └── Self-optimizing pipeline configuration

Multi-Tenant Processing:

- ├── Isolated processing environments per tenant
- ├── Shared infrastructure with tenant-specific customization
- ├── Resource allocation based on SLA requirements
- ├── Separate data quality and validation rules
- └── Tenant-specific monitoring and alerting

Cross-Region Orchestration:

- ├── Global data processing coordination
- ├── Region-specific data residency compliance
- ├── Disaster recovery and failover workflows
- ├── Load balancing across multiple regions
- └── Consistent global data state management

Workflow Optimization:

- ├── Critical path analysis for bottleneck identification
- ├── Resource utilization optimization
- ├── Cost-performance trade-off analysis
- ├── Predictive scaling based on historical patterns
- └── Continuous improvement through machine learning

## Phase 5: Production Deployment and Operations

### CI/CD for Glue ETL Pipelines

**Automated Deployment Strategies:**

**Development Lifecycle Management:**

## ETL Development Pipeline:

### Development Environment:

- |—— Individual developer sandboxes
- |—— Sample datasets for testing
- |—— Rapid iteration and experimentation
- |—— Local testing with Glue development endpoints
- └—— Version control for all ETL artifacts

### Staging Environment:

- |—— Production-like data volumes and schemas
- |—— Full integration testing with dependent systems
- |—— Performance and load testing
- |—— Data quality validation with real scenarios
- └—— Security and compliance verification

### Production Environment:

- |—— Blue-green deployment for zero downtime
- |—— Canary releases for gradual rollout
- |—— Automated rollback procedures
- |—— Comprehensive monitoring and alerting
- └—— Disaster recovery and business continuity

## CI/CD Pipeline Stages:

### 1. Code Commit (Git)

- |—— Automated code quality checks
- |—— Unit test execution
- |—— Security vulnerability scanning
- └—— Code coverage analysis

### 2. Build and Package

- |—— ETL job artifact creation
- |—— Dependency resolution and packaging
- |—— Configuration template generation
- └—— Infrastructure as Code validation

### 3. Automated Testing

- |—— Unit tests for transformation logic
- |—— Integration tests with data sources
- |—— Data quality validation tests
- └—— Performance regression tests

### 4. Deployment

- |—— Infrastructure provisioning (Terraform)
- |—— ETL job deployment and configuration
- |—— Database schema updates
- └—— Monitoring and alerting setup

#### 5. Validation

- |—— Smoke tests for basic functionality
- |—— Data consistency validation
- |—— Performance benchmark comparison
- └—— Business metric verification

### **Infrastructure as Code for Glue:**

## Terraform Configuration Example:

### # Glue Database

```
resource "aws_glue_catalog_database" "retail_analytics" {
  name      = "retail_analytics_${var.environment}"
  description = "Retail analytics data catalog for ${var.environment}"

  create_table_default_permission {
    permissions = ["ALL"]
    principal   = aws_iam_role.glue_service_role.arn
  }
}
```

### # Glue Crawler

```
resource "aws_glue_crawler" "customer_data_crawler" {
  database_name = aws_glue_catalog_database.retail_analytics.name
  name          = "customer-data-crawler-${var.environment}"
  role          = aws_iam_role.glue_service_role.arn

  s3_target {
    path = "s3://${aws_s3_bucket.data_lake.bucket}/customers/"
    exclusions = ["*.log", "*.tmp"]
  }

  schedule = "cron(0 6 * * ? *)" # Daily at 6 AM

  schema_change_policy {
    update_behavior = "UPDATE_IN_DATABASE"
    delete_behavior = "LOG"
  }

  configuration = jsonencode({
    Version = 1.0
    CrawlerOutput = {
      Partitions = { AddOrUpdateBehavior = "InheritFromTable" }
    }
  })
}
```

### # Glue ETL Job

```
resource "aws_glue_job" "customer_data_processing" {
  name      = "customer-data-processing-${var.environment}"
  role_arn = aws_iam_role.glue_service_role.arn
}
```

```
command {
  script_location = "s3://${aws_s3_bucket.glue_assets.bucket}/scripts/customer_data_processing.py"
  python_version = "3"
}
```

```
default_arguments = {
  "--job-language"          = "python"
  "--job-bookmark-option"   = "job-bookmark-enable"
  "--enable-metrics"        = "true"
  "--enable-continuous-cloudwatch-log" = "true"
  "--TempDir"               = "s3://${aws_s3_bucket.glue_temp.bucket}/temp/"
  "--enable-spark-ui"        = "true"
  "--spark-event-logs-path"  = "s3://${aws_s3_bucket.glue_temp.bucket}/sparkHistoryLogs/"
}
```

```
max_retries = 3
timeout     = 60 # minutes
```

```
worker_type    = "G.1X"
number_of_workers = 10
```

```
execution_property {
  max_concurrent_runs = 2
}
}
```

# Glue Workflow

```
resource "aws_glue_workflow" "customer_analytics_pipeline" {
  name      = "customer-analytics-pipeline-${var.environment}"
  description = "Complete customer analytics ETL pipeline"

  max_concurrent_runs = 1
}
```

# CloudWatch Monitoring

```
resource "aws_cloudwatch_dashboard" "glue_monitoring" {
  dashboard_name = "glue-etl-monitoring-${var.environment}"
}
```

```
dashboard_body = jsonencode({
  widgets = [
    {
      type = "metric"
      x     = 0
    }
  ]
})
```

```

y      = 0
width  = 12
height = 6

properties = {
  metrics = [
    ["AWS/Glue", "glue.driver.aggregate.numCompletedTasks", "JobName",
aws_glue_job.customer_data_processing.name],
    [".", "glue.driver.aggregate.numFailedTasks", ".", "."],
  ]
  view    = "timeSeries"
  stacked = false
  region  = var.aws_region
  title   = "Glue Job Task Metrics"
  period  = 300
}
}
]
})
}

```

## Monitoring and Observability

**Comprehensive Monitoring Strategy:**

**Multi-Layer Monitoring Architecture:**



## Monitoring Layer Stack:

### Infrastructure Monitoring:

- └── AWS CloudWatch for service metrics
- └── Glue job execution statistics
- └── Resource utilization and performance
- └── Cost tracking and optimization alerts
- └── Service availability and health checks

### Application Monitoring:

- └── ETL job success/failure rates
- └── Data processing latency and throughput
- └── Transformation accuracy and quality metrics
- └── Business rule compliance rates
- └── Custom application-specific KPIs

### Data Monitoring:

- └── Data freshness and completeness
- └── Schema evolution and compatibility
- └── Data quality trends and degradation
- └── Lineage tracking and impact analysis
- └── Compliance and governance metrics

### Business Monitoring:

- └── SLA compliance and performance
- └── End-user satisfaction and adoption
- └── Business process effectiveness
- └── ROI and value realization tracking
- └── Strategic goal alignment metrics

## Monitoring Dashboard Categories:

### Technical Operations Dashboard:

- └── Job execution status and history
- └── Resource utilization and scaling
- └── Error rates and failure analysis
- └── Performance trends and benchmarks
- └── Cost optimization opportunities

### Data Quality Dashboard:

- └── Overall data quality scores
- └── Quality trends by source and pipeline
- └── Issue detection and resolution tracking

- └── Compliance status and audit trails
- └── Quality SLA performance metrics

#### Business Impact Dashboard:

- └── Data availability and freshness
- └── Business process completion rates
- └── Decision support system performance
- └── Analytics adoption and usage patterns
- └── Value delivered through data insights

### **Alerting and Incident Response:**

## Intelligent Alerting Framework:

### Alert Severity Levels:

#### Critical (P0) - Immediate Response (5 minutes):

- └── Complete pipeline failure affecting business operations
- └── Data loss or corruption detected
- └── Security breach or unauthorized access
- └── Compliance violation with regulatory impact
- └── Customer-facing service disruption

#### High (P1) - Urgent Response (30 minutes):

- └── Individual job failures with business impact
- └── Data quality degradation below SLA thresholds
- └── Performance degradation affecting user experience
- └── Cost anomalies exceeding budget thresholds
- └── Dependency service outages

#### Medium (P2) - Standard Response (2 hours):

- └── Non-critical job failures with workarounds
- └── Performance issues not affecting end users
- └── Schema evolution requiring attention
- └── Resource utilization trending toward limits
- └── Documentation or process improvement needs

#### Low (P3) - Planned Response (Next business day):

- └── Informational alerts and trend notifications
- └── Optimization opportunities identified
- └── Routine maintenance reminders
- └── Usage pattern analysis results
- └── Performance improvement recommendations

### Alert Routing and Escalation:

- └── PagerDuty integration for critical alerts
- └── Slack channels for team collaboration
- └── Email distribution for documentation
- └── SMS for high-priority out-of-hours alerts
- └── Dashboard notifications for trend awareness

### Incident Response Procedures:

#### 1. Alert Detection and Classification

- └── Automated alert severity determination
- └── Context gathering and initial assessment

- └── Stakeholder notification based on impact
- └── Response team assembly and coordination

## 2. Investigation and Diagnosis

- └── Log analysis and error investigation
- └── Impact assessment and scope determination
- └── Root cause analysis and timeline reconstruction
- └── Solution identification and planning

## 3. Resolution and Recovery

- └── Immediate mitigation actions
- └── System recovery and validation
- └── Communication to affected stakeholders
- └── Documentation of resolution steps

## 4. Post-Incident Review

- └── Timeline analysis and lessons learned
- └── Process improvement identification
- └── Preventive measure implementation
- └── Knowledge base updates and training

# **Phase 6: Advanced Glue Features and Optimization**

## **Glue Studio Visual ETL**

### **Visual Development Environment:**

### **Drag-and-Drop ETL Design:**

## Visual ETL Development Concepts:

### Traditional Code-Based Approach:

- └— Requires deep Spark/Python knowledge
- └— Time-intensive development cycle
- └— Complex debugging and testing
- └— Limited collaboration with business users
- └— High maintenance overhead

### Visual ETL Approach:

- └— Intuitive drag-and-drop interface
- └— Pre-built transformation components
- └— Automatic code generation and optimization
- └— Business user friendly design
- └— Rapid prototyping and iteration

## Visual ETL Component Categories:

### Data Source Nodes:

- └— S3 data source (various formats)
- └— Database connections (JDBC)
- └— Glue Data Catalog tables
- └— Streaming data sources (Kinesis)
- └— Custom connector integrations

### Transformation Nodes:

- └— Filter: Row-level filtering with conditions
- └— Map: Column transformations and calculations
- └— Join: Multi-table joining with various join types
- └— Aggregate: Grouping and statistical operations
- └— Union: Combining multiple datasets
- └— Split: Conditional data routing
- └— Custom Transform: Python/Scala code blocks

### Target Nodes:

- └— S3 destinations with format selection
- └— Database targets (JDBC)
- └— Glue Data Catalog table updates
- └— Streaming outputs (Kinesis)
- └— External system integrations

### Example Visual ETL Flow:

[S3 Customer Data] → [Filter Active Customers] → [Join with Transactions] →

[Aggregate by Customer] → [Add Calculated Fields] → [Data Quality Check] →  
[Write to Parquet] → [Update Catalog]

Benefits of Visual ETL:

- └── 80% faster development for standard transformations
- └── Self-documenting pipeline logic
- └── Lower barrier to entry for business analysts
- └── Automatic best practice implementation
- └── Integrated testing and data preview capabilities

## Advanced Visual ETL Patterns:

Complex Transformation Scenarios:

Conditional Processing Flow:

Data Input → Quality Check → [Pass/Fail Branch]

- └── Pass → Standard Processing → Output
- └── Fail → Error Handling → Quarantine → Notification

Multi-Source Integration:

Customer Data (S3) ┐

└→ Join on Customer ID → Enriched Dataset → Analytics

Transaction Data ┘

Incremental Processing:

Historical Data → Full Load (Initial)

New Data → Change Detection → Incremental Updates → Merge with Historical

Error Handling Pattern:

Data Input → Transformation → [Success/Error Branch]

- └── Success → Continue Processing
- └── Error → Error Logging → Dead Letter Queue → Manual Review

Performance Optimization in Visual ETL:

- └── Automatic partition pruning for large datasets
- └── Intelligent join optimization (broadcast vs sort-merge)
- └── Dynamic resource allocation based on data size
- └── Caching recommendations for repeated operations
- └── Columnar format suggestions for storage optimization

## Custom Connectors and Integrations

**Extending Glue Capabilities:**

## Custom Connector Development:

## Connector Architecture Patterns:

### Third-Party Database Connectors:

- └── MongoDB connector for document databases
- └── Cassandra connector for wide-column stores
- └── Elasticsearch connector for search analytics
- └── Redis connector for cache data integration
- └── Neo4j connector for graph data processing

### SaaS Application Connectors:

- └── Salesforce connector for CRM data
- └── ServiceNow connector for ITSM data
- └── Workday connector for HR data
- └── Adobe Analytics connector for web data
- └── Custom API connectors for proprietary systems

### Protocol-Specific Connectors:

- └── FTP/SFTP connectors for file transfers
- └── REST API connectors with authentication
- └── SOAP web service connectors
- └── Message queue connectors (RabbitMQ, ActiveMQ)
- └── Streaming protocol connectors (Apache Pulsar)

## Connector Development Framework:

### 1. Connection Interface Implementation

- └── Authentication and credential management
- └── Connection pooling and resource management
- └── Error handling and retry logic
- └── Health check and monitoring capabilities

### 2. Schema Discovery Implementation

- └── Automatic schema inference from source
- └── Metadata extraction and cataloging
- └── Schema evolution detection
- └── Data type mapping to Glue catalog

### 3. Data Reading Implementation

- └── Efficient data retrieval strategies
- └── Pagination and batch processing
- └── Incremental data extraction
- └── Parallelization for large datasets

### 4. Data Writing Implementation



- └── Bulk insert and update operations
- └── Transaction management and consistency
- └── Conflict resolution strategies
- └── Performance optimization techniques

Example Custom Connector Usage:

# MongoDB Connector in Glue ETL

```
connection_options = {  
    "connectionName": "mongodb_production",  
    "database": "customer_analytics",  
    "collection": "user_behaviors",  
    "batchSize
```