

Day 24: Amazon Redshift - Mastering Cloud Data Warehousing for Data Engineers

What You'll Learn Today (Concept-First Approach)

Primary Focus: Understanding cloud data warehousing and why columnar storage revolutionizes analytics

Secondary Focus: Hands-on implementation through Redshift console and performance optimization

Dataset for Context: Historical retail sales data from Kaggle for enterprise analytics

Learning Philosophy for Day 24

"Understand the warehouse before stocking the shelves"







We'll start with data warehousing concepts, explore columnar storage advantages, understand MPP architecture, and build production-ready analytics solutions.

The Data Warehousing Revolution: Why Redshift Matters

The Problem: Analytics Chaos in Traditional Databases

Scenario: You're running business intelligence queries on a traditional OLTP database...

Without Cloud Data Warehouse (Traditional Pain):

-  Monday 9 AM: Run monthly sales report
-  Monday 9:15 AM: Still waiting... database locked
-  Monday 9:30 AM: Query timeout after 15 minutes
-  Monday 9:45 AM: Restart query with smaller dataset
-  Monday 10:30 AM: Get partial results, business decisions delayed
-  Monday 11:00 AM: Operations team complains about database slowdown

Problems:

- OLTP databases not designed for analytics
- Row-based storage inefficient for aggregations
- Single-server limitations hit quickly
- Analytics queries interfere with operations
- No scalability for growing data volumes

- Complex maintenance and tuning required

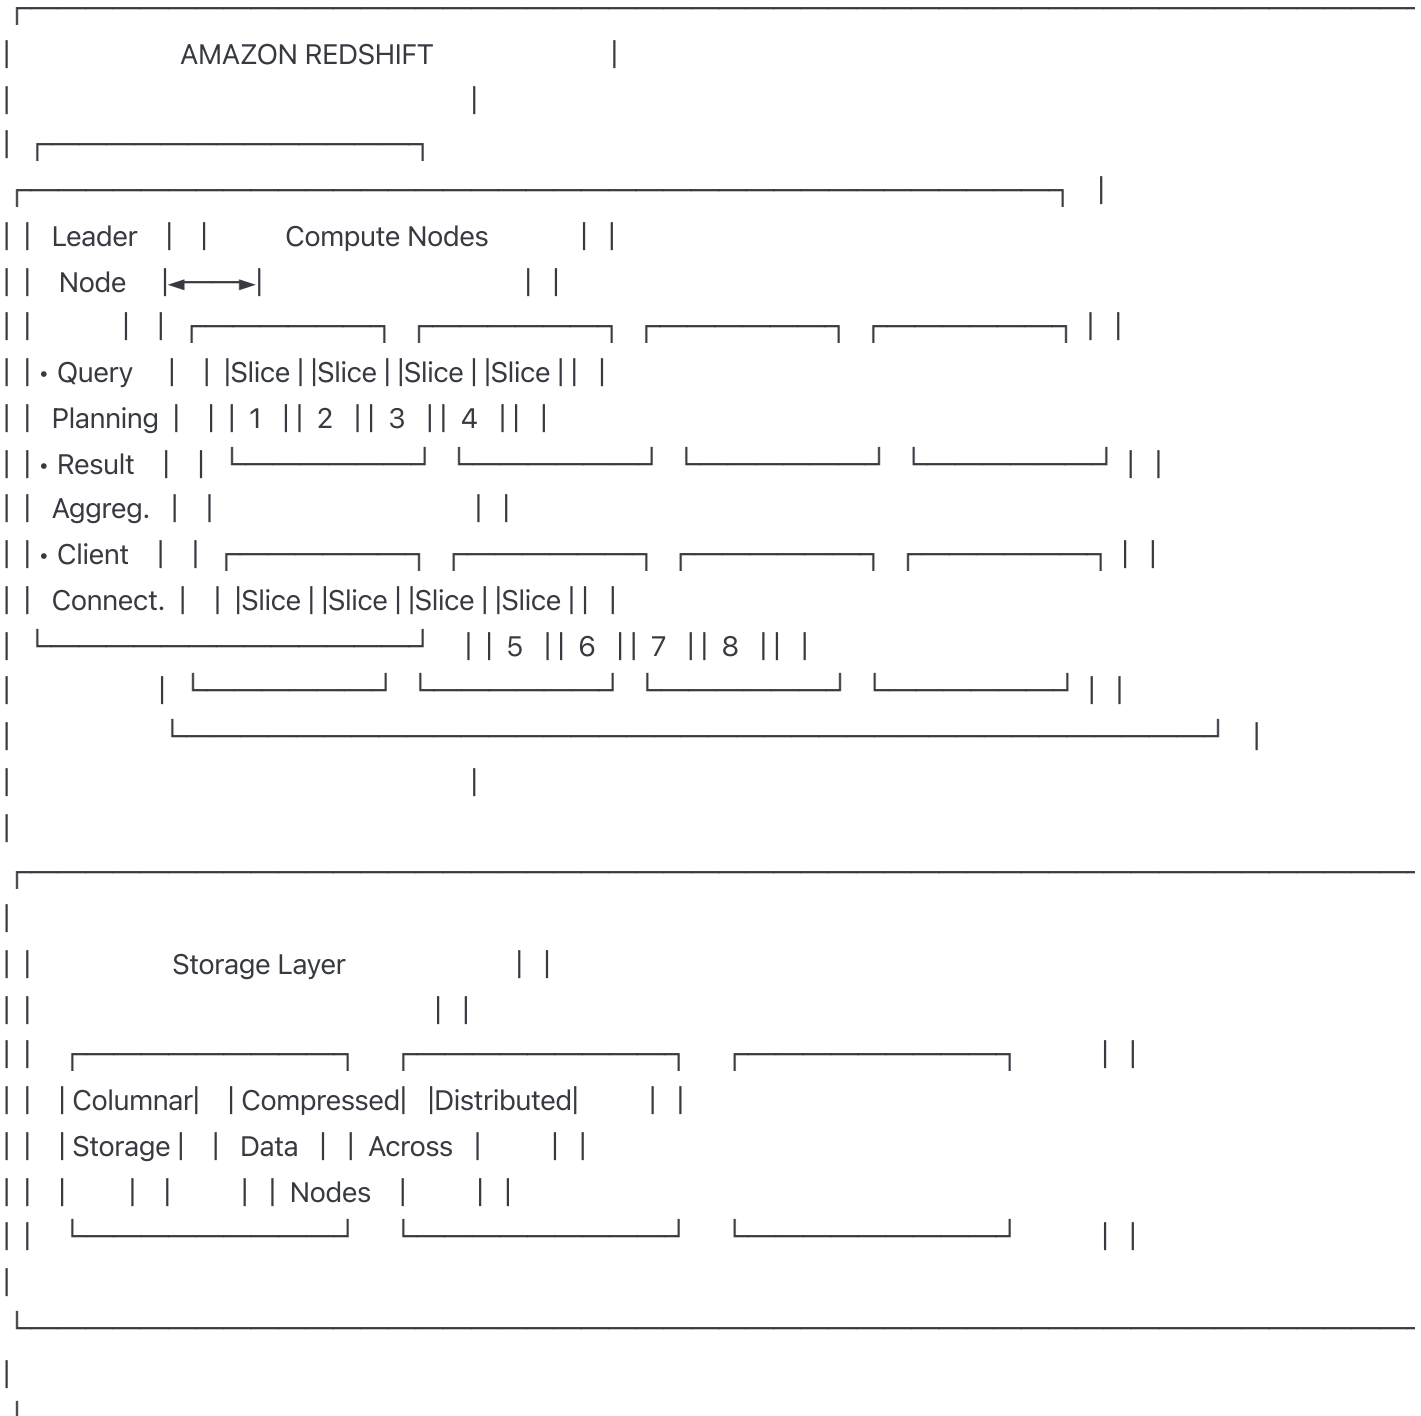
The Redshift Solution: Purpose-Built Analytics Engine

Think of Redshift like this:

- **Traditional Database:** Like a busy grocery store - optimized for many small transactions
- **Redshift:** Like a massive warehouse distribution center - optimized for bulk operations and analytics

Understanding Redshift Architecture (Visual Approach)

The Redshift Mental Model



Key Redshift Components

1. Leader Node

- **Query Coordinator:** Plans and coordinates query execution
- **Client Interface:** Handles all client connections
- **Result Aggregator:** Combines results from compute nodes
- **Metadata Manager:** Stores catalog information

2. Compute Nodes

- **Parallel Processing:** Execute queries in parallel
- **Data Storage:** Store actual table data
- **Local Processing:** Perform computations on local data
- **Network Communication:** Coordinate with other nodes

3. Node Slices

- **Parallel Units:** Each node divided into slices
- **CPU Allocation:** Each slice gets dedicated CPU and memory
- **Data Partitions:** Data distributed across slices
- **Concurrent Processing:** Multiple operations per node

Redshift Setup and Configuration (Concept-First Approach)

Understanding Cluster Architecture

Cluster Sizing Decision Framework:

Data Volume Decision Tree:

- ├── < 100 GB: Single dc2.large node
- ├── 100 GB - 1 TB: 2-4 dc2.large nodes
- ├── 1 TB - 10 TB: 4-8 dc2.8xlarge nodes
- ├── 10 TB - 100 TB: 8-32 ds2.8xlarge nodes
- └── > 100 TB: ra3.xlplus or ra3.4xlarge nodes

Cluster Configuration Concepts

Node Types Understanding:

1. Dense Compute (DC2):

- SSD storage, high CPU
- Best for: High-performance analytics
- Use when: < 10 TB data, complex queries

2. Dense Storage (DS2):

- HDD storage, cost-effective
- Best for: Large datasets, simple queries
- Use when: > 10 TB data, cost optimization

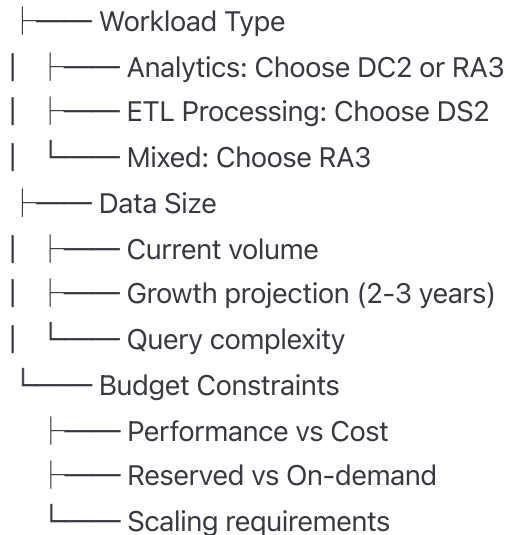
3. RA3:

- Managed storage, compute separation
- Best for: Variable workloads, scaling
- Use when: Unpredictable growth, flexibility needed

Initial Cluster Setup (AWS Console)

Step 1: Cluster Creation Decisions

Cluster Configuration Framework:



Understanding Columnar Storage (The Game Changer)

Row Storage vs Columnar Storage (Visual Learning)

Traditional Row Storage:

Customer Table (Row-based):

Record 1: [ID:1001, Name:"John", Age:25, City:"NYC", Salary:50000]

Record 2: [ID:1002, Name:"Jane", Age:30, City:"LA", Salary:60000]

Record 3: [ID:1003, Name:"Bob", Age:35, City:"NYC", Salary:70000]

Query: SELECT AVG(Salary) FROM Customer WHERE City='NYC'

Problem: Must read entire rows, even unused columns

Redshift Columnar Storage:

Customer Table (Column-based):

ID Column: [1001, 1002, 1003, ...]

Name Column: ["John", "Jane", "Bob", ...]

Age Column: [25, 30, 35, ...]

City Column: ["NYC", "LA", "NYC", ...]

Salary Column: [50000, 60000, 70000, ...]

Query: SELECT AVG(Salary) FROM Customer WHERE City='NYC'

Advantage: Only read City and Salary columns!

Columnar Advantages (Real Numbers)

Performance Impact:

- **I/O Reduction:** 80-90% less data read
- **Compression:** 75-85% storage savings
- **Cache Efficiency:** Better CPU cache utilization
- **Parallel Processing:** Column-level parallelism

Real Example:

Traditional Database: 1 TB table scan

Redshift Columnar: 100 GB data read (10x improvement)


With Compression: 25 GB actual I/O (40x improvement)

Data Distribution Strategies (Performance Foundation)

Understanding Distribution Styles

1. Even Distribution (DISTSTYLE EVEN)

Concept: Data spread equally across all nodes



Best for: Small tables, no joins, reference data

Example: Country codes, product categories

2. Key Distribution (DISTSTYLE KEY)

Concept: Data distributed by column value hash

Customer ID 1000-2499 → Node 1

Customer ID 2500-4999 → Node 2

Customer ID 5000-7499 → Node 3


Customer ID 7500-9999 → Node 4

Best for: Large tables with frequent joins

Example: Customer and Order tables joined on customer_id

3. All Distribution (DISTSTYLE ALL)

Concept: Complete table copy on every node



Best for: Small dimension tables (<2-3 million rows)

Example: Product dimension, Date dimension

Distribution Strategy Decision Framework

Distribution Decision Tree:

- └─ Table Size < 2M rows → DISTSTYLE ALL
- └─ Frequently Joined Tables → DISTSTYLE KEY (on join column)
- └─ No Clear Join Pattern → DISTSTYLE EVEN
- └─ Fact Table → DISTSTYLE KEY (on primary dimension)

Sort Keys: The Query Accelerator

Understanding Sort Key Concepts

Why Sort Keys Matter:

- **Zone Maps:** Redshift creates min/max statistics per 1MB block
- **Block Skipping:** Queries skip irrelevant blocks entirely
- **Compression:** Better compression on sorted data
- **Join Performance:** Merge joins become possible

Sort Key Types (Visual Understanding)

1. Compound Sort Keys

Multi-column sort: (date, region, product)

Data sorted like phone book: Smith, John then Smith, Mary

Block 1: [2024-01-01, East, Product A]

Block 2: [2024-01-01, East, Product B]

Block 3: [2024-01-01, West, Product A]

Block 4: [2024-01-02, East, Product A]

Query: WHERE date='2024-01-01' AND region='East'

Result: Skip blocks 3 and 4 (50% less I/O)

2. Interleaved Sort Keys

Multi-dimensional sort: Balanced across all columns

Better for: Queries filtering on different column combinations

Example: Sometimes filter by date, sometimes by region

Sort Key Selection Strategy

Sort Key Decision Framework:

- |—— Time-series data → Sort by timestamp
- |—— Geographic data → Sort by region/location
- |—— Customer analytics → Sort by customer_id, date
- |—— Sales data → Sort by date, store_id
- |—— Mixed queries → Consider interleaved sort

Hands-On Implementation: Building a Retail Analytics Warehouse

Dataset: Retail Sales Analytics

Source: Kaggle Retail Analytics Dataset **URL:** [kaggle.com/datasets/manjeetsingh/retaildataset](https://www.kaggle.com/datasets/manjeetsingh/retaildataset)

Files: sales_data.csv, products.csv, customers.csv, stores.csv

Business Context:

- Retail chain with 100+ stores
- 5 years of historical data

- Product catalog with 10K+ items
- Customer base of 500K+ users



Schema Design: Star Schema Implementation

sql

-- Fact Table: Sales (Largest table)

```
CREATE TABLE fact_sales (  
  sale_id BIGINT IDENTITY(1,1),  
  date_key DATE NOT NULL,  
  customer_key INT NOT NULL,  
  product_key INT NOT NULL,  
  store_key INT NOT NULL,  
  quantity INT NOT NULL,  
  unit_price DECIMAL(10,2) NOT NULL,  
  total_amount DECIMAL(12,2) NOT NULL,  
  discount_amount DECIMAL(10,2) DEFAULT 0  
)  
DISTSTYLE KEY  
DISTKEY(customer_key) -- Most common join  
SORTKEY(date_key, store_key); -- Time-series analysis
```

-- Dimension Table: Customers

```
CREATE TABLE dim_customer (  
  customer_key INT NOT NULL,  
  customer_id VARCHAR(50) NOT NULL,  
  first_name VARCHAR(100),  
  last_name VARCHAR(100),  
  email VARCHAR(200),  
  phone VARCHAR(20),  
  birth_date DATE,  
  gender VARCHAR(10),  
  city VARCHAR(100),  
  state VARCHAR(50),  
  country VARCHAR(50),  
  registration_date DATE  
)  
DISTSTYLE ALL -- Small dimension, replicate everywhere  
SORTKEY(customer_key);
```

-- Dimension Table: Products

```
CREATE TABLE dim_product (  
  product_key INT NOT NULL,  
  product_id VARCHAR(50) NOT NULL,  
  product_name VARCHAR(200) NOT NULL,  
  category VARCHAR(100),  
  subcategory VARCHAR(100),  
  brand VARCHAR(100),  
  supplier_name VARCHAR(100),
```

```
    unit_cost DECIMAL(10,2),
    retail_price DECIMAL(10,2),
    product_status VARCHAR(20)
)
DISTSTYLE ALL
SORTKEY(category, subcategory);
```

-- Dimension Table: Stores

```
CREATE TABLE dim_store (
    store_key INT NOT NULL,
    store_id VARCHAR(20) NOT NULL,
    store_name VARCHAR(100),
    address VARCHAR(200),
    city VARCHAR(100),
    state VARCHAR(50),
    postal_code VARCHAR(20),
    country VARCHAR(50),
    store_type VARCHAR(50),
    store_size_sqft INT,
    opening_date DATE
)
DISTSTYLE ALL
SORTKEY(state, city);
```

-- Dimension Table: Date (Analytics essential)

```
CREATE TABLE dim_date (
    date_key DATE NOT NULL,
    year INT NOT NULL,
    quarter INT NOT NULL,
    month INT NOT NULL,
    month_name VARCHAR(20),
    day_of_month INT NOT NULL,
    day_of_week INT NOT NULL,
    day_name VARCHAR(20),
    week_of_year INT,
    is_weekend BOOLEAN,
    is_holiday BOOLEAN,
    fiscal_year INT,
    fiscal_quarter INT
)
DISTSTYLE ALL
SORTKEY(date_key);
```

Data Loading Strategies (Concept-Driven)

Understanding COPY Command (Redshift's Strength)

Why COPY vs INSERT:

- **Parallel Loading:** COPY uses all nodes simultaneously
- **Compression Detection:** Automatic compression analysis
- **Error Handling:** Detailed error reporting and recovery
- **Performance:** 10-100x faster than individual INSERTs

Loading from S3 (Production Pattern)

```
sql

-- Step 1: Upload data to S3 (conceptually)
-- sales_data/
-- |----- year=2019/month=01/sales_part_001.csv
-- |----- year=2019/month=01/sales_part_002.csv
-- |----- year=2019/month=02/sales_part_001.csv
-- |----- ...

-- Step 2: COPY command with best practices
COPY fact_sales (
    date_key, customer_key, product_key, store_key,
    quantity, unit_price, total_amount, discount_amount
)
FROM 's3://retail-analytics-bucket/sales_data/'
IAM_ROLE 'arn:aws:iam::123456789:role/RedshiftS3AccessRole'
CSV
DELIMITER ','
IGNOREHEADER 1
REGION 'us-east-1'
COMPUPDATE ON      -- Analyze compression
STATUPDATE ON      -- Update table statistics
MAXERROR 1000      -- Allow some bad records
TRUNCATECOLUMNS   -- Handle data overflow gracefully
TIMEFORMAT 'YYYY-MM-DD HH:MI:SS';
```

Understanding Load Performance

COPY Performance Factors:

Parallel Loading Optimization:

- └── File Count = Number of Slices (optimal)
- └── File Size: 1-1000 MB per file
- └── Compression: GZIP recommended for S3
- └── Data Types: Choose appropriate sizes
- └── Error Handling: Balance between speed and accuracy

⚡ Performance Optimization (Deep Concepts)

🧠 Query Execution Understanding

Redshift Query Lifecycle:

1. Query Parsing → Validate syntax and permissions
2. Query Planning → Generate execution plan
3. Code Generation → Compile optimized code
4. Execution → Parallel processing across nodes
5. Result Assembly → Aggregate results on leader node

📊 Analyzing Query Performance

Using EXPLAIN Command:

```
sql

-- Understanding query execution plan
EXPLAIN (COSTS TRUE, BUFFERS TRUE)
SELECT
    p.category,
    s.state,
    DATE_TRUNC('month', f.date_key) as month,
    SUM(f.total_amount) as monthly_revenue,
    COUNT(*) as transaction_count
FROM fact_sales f
JOIN dim_product p ON f.product_key = p.product_key
JOIN dim_store s ON f.store_key = s.store_key
WHERE f.date_key >= '2023-01-01'
    AND p.category IN ('Electronics', 'Clothing')
GROUP BY 1, 2, 3
ORDER BY monthly_revenue DESC;
```

Reading Execution Plans:

- **DS_DIST_ALL_NONE:** Good - no data movement needed
- **DS_DIST_BOTH:** Warning - data redistribution required
- **DS_DIST_ALL_INNER:** Expensive - large table broadcast

Performance Tuning Strategies

1. Distribution Key Optimization

```
sql

-- Problem: Poor distribution causing data movement
SELECT query, slice, rows
FROM stv_slices
WHERE rows > AVG(rows) * 2; -- Find skewed distribution

-- Solution: Choose better distribution key
ALTER TABLE fact_sales
ALTER DISTSTYLE KEY
ALTER DISTKEY customer_key; -- If customer joins are common
```

2. Sort Key Maintenance

```
sql

-- Check sort key effectiveness
SELECT schema, "table", unsorted, vacuum_sort_benefit
FROM svv_table_info
WHERE unsorted > 20; -- Tables needing maintenance

-- Solution: Regular VACUUM operations
VACUUM SORT ONLY fact_sales;
VACUUM DELETE ONLY dim_customer;
```

3. Compression Optimization

sql

-- Analyze compression recommendations

ANALYZE COMPRESSION dim_product;

-- Apply compression encoding

ALTER TABLE dim_product

ALTER COLUMN product_name **TYPE VARCHAR**(200) **ENCODE** LZO,

ALTER COLUMN category **TYPE VARCHAR**(100) **ENCODE** BYTEDICT,

ALTER COLUMN retail_price **TYPE DECIMAL**(10,2) **ENCODE** DELTA32K;

Workload Management (WLM) Concepts

Understanding Query Queues

WLM Purpose:

- **Resource Allocation:** Control memory and concurrency
- **Query Prioritization:** Business-critical queries first
- **Performance Isolation:** Prevent resource conflicts

WLM Configuration Strategy

json

```
{
  "query_concurrency": 5,
  "query_queues": [
    {
      "name": "priority_queue",
      "query_concurrency": 2,
      "memory_percent": 40,
      "user_groups": ["executives", "analysts"]
    },
    {
      "name": "etl_queue",
      "query_concurrency": 1,
      "memory_percent": 35,
      "user_groups": ["etl_users"]
    },
    {
      "name": "general_queue",
      "query_concurrency": 3,
      "memory_percent": 25,
      "user_groups": ["general_users"]
    }
  ]
}
```

Query Monitoring Concepts

sql

-- Monitor query performance

SELECT

query,
userid,
query_start_time,
total_exec_time,
queue_time,
exec_time,
rows,
bytes

FROM stl_query_metrics

WHERE query_start_time >= CURRENT_DATE - 1

ORDER BY total_exec_time DESC

LIMIT 10;

Security and Access Control

Redshift Security Model

Multi-Layer Security:

Security Layers:

- |—— Network Level: VPC, Security Groups, CIDR blocks
- |—— Cluster Level: Encryption at rest and in transit
- |—— Database Level: User management, database isolation
- |—— Schema Level: Schema-specific permissions
- |—— Object Level: Table, view, function permissions

User and Permission Management

sql

-- Create user groups for different roles

CREATE GROUP analysts;

CREATE GROUP data_engineers;

CREATE GROUP executives;

-- Create users and assign to groups

CREATE USER john_analyst PASSWORD 'SecurePass123!';

ALTER GROUP analysts ADD USER john_analyst;

-- Grant schema-level permissions

GRANT USAGE ON SCHEMA retail_analytics TO GROUP analysts;

-- Grant table-level permissions

GRANT SELECT ON fact_sales TO GROUP analysts;

GRANT SELECT ON ALL TABLES IN SCHEMA retail_analytics TO GROUP executives;

-- Grant specific column access (for sensitive data)

CREATE VIEW customer_safe AS

SELECT customer_key, first_name, city, state

FROM dim_customer; *-- Exclude phone, email*

GRANT SELECT ON customer_safe TO GROUP analysts;



Integration with BI Tools



Connecting Business Intelligence Tools

Common BI Tool Connections:

- **Tableau:** Native Redshift connector
- **Power BI:** Amazon Redshift connector
- **Looker:** Database connection via JDBC
- **QuickSight:** Built-in AWS integration



Optimization for BI Workloads

sql

-- Create aggregated tables for faster BI queries

```
CREATE TABLE monthly_sales_summary AS
SELECT
  DATE_TRUNC('month', f.date_key) as month,
  p.category,
  s.state,
  SUM(f.total_amount) as total_revenue,
  SUM(f.quantity) as total_quantity,
  COUNT(*) as transaction_count,
  AVG(f.total_amount) as avg_transaction_value
FROM fact_sales f
JOIN dim_product p ON f.product_key = p.product_key
JOIN dim_store s ON f.store_key = s.store_key
GROUP BY 1, 2, 3;
```

-- Create materialized views for real-time dashboards

```
CREATE MATERIALIZED VIEW daily_kpis AS
SELECT
  f.date_key,
  SUM(f.total_amount) as daily_revenue,
  COUNT(DISTINCT f.customer_key) as unique_customers,
  COUNT(*) as transaction_count
FROM fact_sales f
WHERE f.date_key >= CURRENT_DATE - 30
GROUP BY f.date_key;
```

Cost Optimization Strategies

Understanding Redshift Pricing

Cost Components:

- **Compute:** Node hours (on-demand vs reserved)
- **Storage:** Managed storage for RA3 (separate billing)
- **Data Transfer:** Cross-region and internet egress
- **Backup:** Automated and manual snapshots

Cost Optimization Techniques

1. Reserved Instances

Reserved Instance Savings:

- |—— 1 Year: 20-40% savings
- |—— 3 Years: 40-60% savings
- |—— All Upfront: Maximum discount

2. Automatic Scaling

```
sql

-- Configure auto-pause for development clusters
ALTER CLUSTER dev-analytics
SET auto_pause = true,
   pause_timeout = 300; -- 5 minutes of inactivity
```

3. Storage Optimization

```
sql

-- Remove old data with lifecycle policies
DELETE FROM fact_sales
WHERE date_key < CURRENT_DATE - INTERVAL '5 years';

-- Compress historical data
VACUUM DELETE ONLY fact_sales;
ANALYZE fact_sales;
```

Monitoring and Maintenance

Cluster Health Monitoring

Key Metrics to Track:

- **CPU Utilization:** Target 60-80% average
- **Disk Space:** Keep below 80% capacity
- **Query Performance:** Monitor execution times
- **Concurrency:** Track queue wait times

Automated Maintenance Scripts

sql

```
-- Daily maintenance routine
-- 1. Update table statistics
ANALYZE fact_sales;
ANALYZE dim_customer;
ANALYZE dim_product;

-- 2. Vacuum tables (weekly)
VACUUM DELETE ONLY fact_sales;
VACUUM SORT ONLY fact_sales;

-- 3. Monitor query performance
SELECT
    schemaname,
    tablename,
    attname as column_name,
    n_distinct,
    correlation
FROM pg_stats
WHERE schemaname = 'retail_analytics'
    AND n_distinct < 100; -- Potential encoding opportunities
```

Backup and Disaster Recovery

Backup Strategy Concepts

Redshift Backup Types:

- **Automated Snapshots:** Daily, retention period configurable
- **Manual Snapshots:** On-demand, kept until deleted
- **Cross-Region Snapshots:** Disaster recovery protection

Recovery Planning

```
sql
```

```
-- Create manual snapshot before major changes
```

```
CREATE SNAPSHOT retail_analytics_snapshot_20241124  
FROM CLUSTER retail-analytics-prod;
```

```
-- Restore from snapshot (concept)
```

```
RESTORE CLUSTER retail-analytics-restored  
FROM SNAPSHOT retail_analytics_snapshot_20241124  
CLUSTER_SUBNET_GROUP default;
```



Advanced Redshift Features



Redshift Spectrum (Data Lake Integration)

Concept: Query S3 data directly without loading

```
sql
```

```
-- Create external schema for S3 data
```

```
CREATE EXTERNAL SCHEMA spectrum_schema  
FROM DATA CATALOG  
DATABASE 'retail_data_lake'  
IAM_ROLE 'arn:aws:iam::123456789:role/SpectrumRole';
```

```
-- Query S3 data alongside Redshift tables
```

```
SELECT  
    f.date_key,  
    s3.raw_events,  
    SUM(f.total_amount) as revenue  
FROM fact_sales f  
JOIN spectrum_schema.s3_events s3  
    ON f.date_key = s3.event_date  
GROUP BY 1, 2;
```



Redshift ML (Machine Learning Integration)

sql

-- Create ML model directly in Redshift

```
CREATE MODEL customer_lifetime_value_model
FROM (
  SELECT
    customer_key,
    recency_days,
    frequency_count,
    monetary_value,
    future_clv as target
  FROM customer_features
)
FUNCTION predict_clv
IAM_ROLE 'arn:aws:iam::123456789:role/RedshiftMLRole'
SETTINGS (
  S3_BUCKET 'ml-models-bucket',
  MAX_RUNTIME 7200
);
```



Performance Benchmarking



Establishing Baseline Performance

Key Performance Indicators:

sql

-- Query performance baseline

SELECT

DATE_TRUNC('hour', query_start_time) as hour,

COUNT(*) as query_count,

AVG(total_exec_time/1000000.0) as avg_execution_seconds,

PERCENTILE_CONT(0.95) WITHIN GROUP (ORDER BY total_exec_time/1000000.0) as p95_execution_seconds

FROM stl_query_metrics

WHERE query_start_time >= CURRENT_DATE - 7

GROUP BY 1

ORDER BY 1;

-- Storage utilization tracking

SELECT

schemaname,

tablename,

size_in_mb,

rows,

unsorted_rows,

vacuum_sort_benefit

FROM svv_table_info

WHERE schemaname = 'retail_analytics'

ORDER BY size_in_mb DESC;

Capacity Planning Framework

Growth Projection Model:

Current State Assessment:

- └── Data Volume: Current TB, growth rate %/month
- └── Query Volume: Queries/hour, complexity trends
- └── User Growth: Current users, projected growth
- └── Performance SLA: Response time targets

Future State Planning:

- └── 6 Month Projection: 2x data, 1.5x queries
- └── 12 Month Projection: 4x data, 2x queries
- └── 24 Month Projection: 8x data, 3x queries
- └── Scaling Strategy: Vertical vs horizontal scaling

ETL Integration Patterns

Redshift as Target for ETL Pipelines

Common ETL Patterns:

1. Batch Loading Pattern

ETL Pipeline Flow:

S3 Landing Zone → Data Validation → Transformation → Redshift COPY

- └── Staging Tables: Temporary data validation
- └── Transformation Logic: Business rules application
- └── Error Handling: Bad record isolation
- └── Incremental Loading: Delta processing

2. CDC (Change Data Capture) Pattern

```
sql

-- Incremental loading strategy
-- Step 1: Identify new/changed records
CREATE TEMP TABLE staging_sales AS
SELECT * FROM external_staging.sales_updates
WHERE last_modified > (
    SELECT MAX(last_updated)
    FROM fact_sales_control_table
);

-- Step 2: Upsert operation
BEGIN TRANSACTION;

-- Delete existing records that will be updated
DELETE FROM fact_sales
WHERE sale_id IN (SELECT sale_id FROM staging_sales);

-- Insert new and updated records
INSERT INTO fact_sales
SELECT * FROM staging_sales;

-- Update control table
UPDATE fact_sales_control_table
SET last_updated = CURRENT_TIMESTAMP;

COMMIT;
```

ELT vs ETL Decision Framework

When to use ELT with Redshift:

- Large data volumes (> 100 GB)
- Complex transformations benefit from MPP
- Source systems support bulk extraction
- Real-time transformation not required

When to use ETL before Redshift:

- Complex data quality rules
- Multiple source system integration
- Real-time streaming requirements
- Sensitive data masking/encryption needed



Real-World Use Cases and Patterns



Enterprise Analytics Scenarios

1. Customer 360 Analytics

sql

-- Customer behavior analysis combining multiple data sources

```
WITH customer_metrics AS (  
    SELECT  
        c.customer_key,  
        c.registration_date,  
        COUNT(DISTINCT f.date_key) as days_active,  
        SUM(f.total_amount) as lifetime_value,  
        COUNT(*) as total_transactions,  
        AVG(f.total_amount) as avg_transaction_value,  
        MAX(f.date_key) as last_purchase_date,  
        CURRENT_DATE - MAX(f.date_key) as days_since_last_purchase  
    FROM dim_customer c  
    LEFT JOIN fact_sales f ON c.customer_key = f.customer_key  
    GROUP BY 1, 2  
,  
customer_segments AS (  
    SELECT  
        *,  
        CASE  
            WHEN days_since_last_purchase <= 30 AND lifetime_value >= 1000 THEN 'VIP Active'  
            WHEN days_since_last_purchase <= 90 AND lifetime_value >= 500 THEN 'Active Loyal'  
            WHEN days_since_last_purchase <= 180 THEN 'Active Regular'  
            WHEN days_since_last_purchase <= 365 THEN 'At Risk'  
            ELSE 'Churned'  
        END as customer_segment  
    FROM customer_metrics  
)  
SELECT  
    customer_segment,  
    COUNT(*) as customer_count,  
    AVG(lifetime_value) as avg_lifetime_value,  
    AVG(days_since_last_purchase) as avg_recency,  
    SUM(lifetime_value) as total_segment_value  
FROM customer_segments  
GROUP BY 1  
ORDER BY total_segment_value DESC;
```

2. Supply Chain Analytics

sql

-- Inventory turnover and demand forecasting

```
WITH product_performance AS (  
    SELECT  
        p.product_key,  
        p.product_name,  
        p.category,  
        SUM(f.quantity) as total_sold,  
        SUM(f.total_amount) as total_revenue,  
        COUNT(DISTINCT f.date_key) as days_sold,  
        COUNT(DISTINCT f.store_key) as stores_sold_in,  
        AVG(f.quantity) as avg_daily_demand  
    FROM dim_product p  
    JOIN fact_sales f ON p.product_key = f.product_key  
    WHERE f.date_key >= CURRENT_DATE - 90 -- Last 3 months  
    GROUP BY 1, 2, 3  
,  
inventory_metrics AS (  
    SELECT  
        *,  
        total_sold / NULLIF(days_sold, 0) as velocity,  
        total_revenue / NULLIF(total_sold, 0) as avg_selling_price,  
        CASE  
            WHEN total_sold > 1000 AND days_sold > 60 THEN 'Fast Mover'  
            WHEN total_sold > 100 AND days_sold > 30 THEN 'Regular Mover'  
            WHEN total_sold > 10 THEN 'Slow Mover'  
            ELSE 'Dead Stock'  
        END as movement_category  
    FROM product_performance  
)  
SELECT  
    category,  
    movement_category,  
    COUNT(*) as product_count,  
    SUM(total_revenue) as category_revenue,  
    AVG(velocity) as avg_velocity,  
    AVG(avg_selling_price) as avg_price  
FROM inventory_metrics  
GROUP BY 1, 2  
ORDER BY category, category_revenue DESC;
```

sql

-- Financial performance dashboard with time intelligence

WITH monthly_financials AS (

SELECT

DATE_TRUNC('month', f.date_key) as month,
s.state,
p.category,
SUM(f.total_amount) as revenue,
SUM(f.quantity * p.unit_cost) as cost_of_goods,
SUM(f.total_amount) - SUM(f.quantity * p.unit_cost) as gross_profit,
COUNT(*) as transaction_count,
COUNT(DISTINCT f.customer_key) as unique_customers

FROM fact_sales f

JOIN dim_product p ON f.product_key = p.product_key

JOIN dim_store s ON f.store_key = s.store_key

WHERE f.date_key >= '2023-01-01'

GROUP BY 1, 2, 3

),

performance_metrics AS (

SELECT

*,
gross_profit / NULLIF(revenue, 0) * 100 as gross_margin_pct,
revenue / NULLIF(transaction_count, 0) as avg_transaction_value,
LAG(revenue, 1) OVER (PARTITION BY state, category ORDER BY month) as prev_month_revenue,
LAG(revenue, 12) OVER (PARTITION BY state, category ORDER BY month) as same_month_last_year

FROM monthly_financials

)

SELECT

month,
state,
category,
revenue,
gross_profit,
gross_margin_pct,

CASE

WHEN prev_month_revenue IS NOT NULL
THEN (revenue - prev_month_revenue) / prev_month_revenue * 100
ELSE NULL

END as month_over_month_growth_pct,

CASE

WHEN same_month_last_year IS NOT NULL
THEN (revenue - same_month_last_year) / same_month_last_year * 100
ELSE NULL

END as year_over_year_growth_pct

```
FROM performance_metrics
WHERE month >= '2024-01-01' -- Current year focus
ORDER BY month DESC, revenue DESC;
```

Troubleshooting Common Issues

Performance Problems and Solutions

1. Query Hanging or Slow Performance

```
sql

-- Diagnose long-running queries
SELECT
  query,
  userid,
  query_start_time,
  total_exec_time/1000000 as runtime_seconds,
  substr(query_text, 1, 100) as query_preview
FROM stl_query_metrics
WHERE total_exec_time > 300000000 -- > 5 minutes
ORDER BY total_exec_time DESC;

-- Check for blocking queries
SELECT
  blocking.query as blocking_query,
  blocked.query as blocked_query,
  blocking.userid as blocking_user,
  blocked.userid as blocked_user
FROM stv_locks blocking
JOIN stv_locks blocked ON blocking.table_id = blocked.table_id
WHERE blocking.query <> blocked.query;
```

Common Solutions:

- Add appropriate WHERE clauses to reduce data scanned
- Verify distribution and sort keys are optimal
- Check for data skew in distribution
- Consider query rewriting with CTEs or temp tables

2. Storage and Capacity Issues

sql

-- Monitor disk space usage

SELECT

node,
used_mb,
capacity_mb,
used_mb::FLOAT / capacity_mb * 100 as used_percentage

FROM stv_partitions

WHERE used_percentage > 80; *-- Alert threshold*

-- Identify largest tables

SELECT

schemaname,
tablename,
size_in_mb,
rows,
unsorted_pct

FROM svv_table_info

WHERE size_in_mb > 1000 *-- Focus on large tables*

ORDER BY size_in_mb DESC;

3. Connection and Concurrency Issues

```
sql
```

```
-- Monitor active connections
```

```
SELECT
```

```
    datname,  
    username,  
    application_name,  
    client_addr,  
    query_start,  
    state,  
    substr(query, 1, 50) as current_query
```

```
FROM pg_stat_activity
```

```
WHERE state = 'active';
```

```
-- Check WLM queue performance
```

```
SELECT
```

```
    service_class,  
    num_query_tasks,  
    num_executing_tasks,  
    num_queued_tasks,  
    queue_time_percentile
```

```
FROM stl_wlm_query
```

```
WHERE query_start_time >= CURRENT_DATE - 1;
```



Redshift vs Alternatives Comparison



Decision Framework: When to Choose Redshift

Redshift is Optimal for:

- Structured/semi-structured analytics workloads
- Complex SQL queries and joins
- Existing AWS ecosystem integration
- Cost-predictable, steady workloads
- Traditional BI tool integration

Consider Alternatives When:

- Primarily streaming/real-time analytics (→ Kinesis Analytics)
- Unstructured data processing (→ EMR/Spark)
- Serverless, sporadic queries (→ Athena)

- Multi-cloud requirements (→ Snowflake)
- Graph analytics (→ Neptune)

Performance Comparison Framework

Redshift Performance Characteristics:

- └── Query Latency: 100ms - 10 seconds (typical)
- └── Throughput: 1000+ concurrent users
- └── Data Volume: Petabytes supported
- └── Compression: 75-85% typical savings
- └── Scaling: Minutes (resize) to instant (concurrency scaling)

Cost Comparison (rough guidelines):

- └── Redshift: \$0.25/hour/TB (dc2.large)
- └── Snowflake: \$0.40/hour/TB (similar compute)
- └── BigQuery: \$5/TB scanned (on-demand)
- └── Athena: \$5/TB scanned

Production Best Practices

Security Hardening Checklist

Network Security:


- ☒ VPC deployment with private subnets
- ☒ Security groups restricting access to known IPs
- ☒ SSL/TLS encryption for all connections
- ☒ VPC endpoints for S3 access (avoid internet routing)

Data Protection:

- ☒ Encryption at rest using KMS keys
- ☒ Automated backup retention policies
- ☒ Cross-region backup for disaster recovery
- ☒ Column-level security for sensitive data

Access Control:

- ☒ IAM roles instead of embedded credentials
- ☒ Principle of least privilege for all users
- ☒ Regular access reviews and cleanup

-  Audit logging enabled and monitored

Operational Excellence

```
sql

-- Create monitoring views for operational oversight
CREATE VIEW ops_cluster_health AS
SELECT
    'CPU Utilization' as metric,
    AVG(cpu_percent) as current_value,
    80 as threshold,
    CASE WHEN AVG(cpu_percent) > 80 THEN 'CRITICAL' ELSE 'OK' END as status
FROM stl_query_metrics
WHERE query_start_time >= CURRENT_TIMESTAMP - INTERVAL '1 hour'

UNION ALL

SELECT
    'Disk Usage' as metric,
    MAX(used_mb::FLOAT / capacity_mb * 100) as current_value,
    85 as threshold,
    CASE WHEN MAX(used_mb::FLOAT / capacity_mb * 100) > 85 THEN 'CRITICAL' ELSE 'OK' END as status
FROM stv_partitions

UNION ALL

SELECT
    'Queue Wait Time' as metric,
    AVG(queue_time/1000000.0) as current_value,
    30 as threshold,
    CASE WHEN AVG(queue_time/1000000.0) > 30 THEN 'WARNING' ELSE 'OK' END as status
FROM stl_wlm_query
WHERE query_start_time >= CURRENT_TIMESTAMP - INTERVAL '1 hour';
```

Change Management Process

Deployment Pipeline:

Development → Staging → Production

- └—— Schema Changes: DDL review and testing
- └—— Query Changes: Performance testing
- └—— Data Changes: Backup and rollback plan
- └—— Configuration: Gradual rollout strategy

Advanced Analytics Patterns

Time Series Analysis in Redshift

sql

-- Advanced time series analytics with window functions

```
WITH daily_metrics AS (  
    SELECT  
        date_key,  
        SUM(total_amount) as daily_revenue,  
        COUNT(*) as daily_transactions,  
        COUNT(DISTINCT customer_key) as daily_unique_customers  
    FROM fact_sales  
    WHERE date_key >= '2023-01-01'  
    GROUP BY date_key  
)  
time_series_analysis AS (  
    SELECT  
        date_key,  
        daily_revenue,  
        -- Moving averages  
        AVG(daily_revenue) OVER (  
            ORDER BY date_key  
            ROWS BETWEEN 6 PRECEDING AND CURRENT ROW  
        ) as revenue_7day_ma,  
        AVG(daily_revenue) OVER (  
            ORDER BY date_key  
            ROWS BETWEEN 29 PRECEDING AND CURRENT ROW  
        ) as revenue_30day_ma,  
  
        -- Growth calculations  
        LAG(daily_revenue, 1) OVER (ORDER BY date_key) as prev_day_revenue,  
        LAG(daily_revenue, 7) OVER (ORDER BY date_key) as week_ago_revenue,  
        LAG(daily_revenue, 365) OVER (ORDER BY date_key) as year_ago_revenue,  
  
        -- Volatility measures  
        STDDEV(daily_revenue) OVER (  
            ORDER BY date_key  
            ROWS BETWEEN 29 PRECEDING AND CURRENT ROW  
        ) as revenue_volatility_30day  
    FROM daily_metrics  
)  
SELECT  
    date_key,  
    daily_revenue,  
    revenue_7day_ma,  
    revenue_30day_ma,  
    CASE
```

```
    WHEN prev_day_revenue IS NOT NULL
    THEN (daily_revenue - prev_day_revenue) / prev_day_revenue * 100
END as day_over_day_growth_pct,
CASE
    WHEN year_ago_revenue IS NOT NULL
    THEN (daily_revenue - year_ago_revenue) / year_ago_revenue * 100
END as year_over_year_growth_pct,
revenue_volatility_30day
FROM time_series_analysis
WHERE date_key >= '2024-01-01'
ORDER BY date_key;
```

Cohort Analysis for Customer Retention

sql

-- Customer cohort analysis for retention insights

```
WITH customer_cohorts AS (  
    SELECT  
        customer_key,  
        DATE_TRUNC('month', MIN(date_key)) as cohort_month,  
        DATE_TRUNC('month', date_key) as purchase_month,  
        DATEDIFF('month',  
            DATE_TRUNC('month', MIN(date_key)) OVER (PARTITION BY customer_key),  
            DATE_TRUNC('month', date_key)  
        ) as period_number  
    FROM fact_sales  
    GROUP BY customer_key, DATE_TRUNC('month', date_key)  
,  
cohort_sizes AS (  
    SELECT  
        cohort_month,  
        COUNT(DISTINCT customer_key) as cohort_size  
    FROM customer_cohorts  
    WHERE period_number = 0 -- Initial cohort size  
    GROUP BY cohort_month  
,  
retention_analysis AS (  
    SELECT  
        c.cohort_month,  
        c.period_number,  
        COUNT(DISTINCT c.customer_key) as active_customers,  
        s.cohort_size  
    FROM customer_cohorts c  
    JOIN cohort_sizes s ON c.cohort_month = s.cohort_month  
    GROUP BY c.cohort_month, c.period_number, s.cohort_size  
)  
SELECT  
    cohort_month,  
    period_number,  
    active_customers,  
    cohort_size,  
    active_customers::FLOAT / cohort_size * 100 as retention_rate_pct  
FROM retention_analysis  
WHERE cohort_month >= '2023-01-01'  
    AND period_number <= 12 -- Focus on first year  
ORDER BY cohort_month, period_number;
```

Learning Resources and Next Steps

Essential Documentation and Guides

Official AWS Resources:

- Amazon Redshift Documentation: <https://docs.aws.amazon.com/redshift/>
- Redshift Best Practices: <https://docs.aws.amazon.com/redshift/latest/dg/best-practices.html>
- Redshift Performance Tuning: https://docs.aws.amazon.com/redshift/latest/dg/c_optimizing-query-performance.html

Performance and Optimization:

- Query Performance Tuning: Focus on execution plans and WLM
- Distribution Key Selection: Understand data access patterns
- Sort Key Optimization: Align with query filter patterns
- Compression Analysis: Regular compression reviews

Hands-On Practice Recommendations

Progressive Skill Building:

1. **Week 1:** Basic cluster setup and data loading
2. **Week 2:** Query optimization and performance tuning
3. **Week 3:** Advanced analytics and window functions
4. **Week 4:** Integration with ETL tools and BI platforms

Real-World Project Ideas:

- Build customer analytics dashboard for e-commerce data
- Create financial reporting system with time-series analysis
- Implement real-time streaming analytics with Kinesis → Redshift
- Design multi-tenant analytics platform with proper security

Certification and Career Path

Relevant AWS Certifications:

- AWS Certified Data Analytics - Specialty
- AWS Certified Solutions Architect - Associate
- AWS Certified Database - Specialty

Career Progression:

- **Junior Data Engineer:** Focus on data loading and basic queries
- **Data Engineer:** Advanced performance tuning and architecture
- **Senior Data Engineer:** Design enterprise data platforms
- **Data Architect:** Strategic data platform and governance design

Tomorrow's Preview: Data Pipeline Monitoring

Building on today's Redshift mastery, tomorrow we'll explore:

- **Comprehensive monitoring strategies** for data pipelines
- **Alerting systems** that catch issues before they impact business
- **Observability patterns** for complex data architectures
- **Incident response** procedures for data engineering teams


The combination of powerful analytics engines like Redshift with robust monitoring creates truly reliable data platforms that businesses can depend on.

Key Takeaways for Day 24

Conceptual Mastery Achieved:  **Columnar Storage:** Understand why it revolutionizes analytics performance

 **MPP Architecture:** Grasp how parallel processing scales linearly

 **Distribution Strategies:** Know when and how to optimize data placement

 **Performance Tuning:** Apply systematic optimization approaches

 **Production Operations:** Implement enterprise-grade data warehousing

Mental Model Transformation:

- Row-based thinking → Column-optimized analytics mindset
- Single-server limitations → Distributed computing possibilities
- Manual optimization → Intelligent automated systems
- Isolated databases → Integrated analytics ecosystem

Real Business Impact Understanding:

- Query performance improvements: 10-100x faster analytics
- Cost optimization: 60-80% storage savings through compression
- Scalability: Linear performance scaling with cluster growth

- Integration: Seamless AWS ecosystem connectivity

Tomorrow, we'll ensure these powerful systems stay healthy and performant through comprehensive monitoring and observability! 🚀